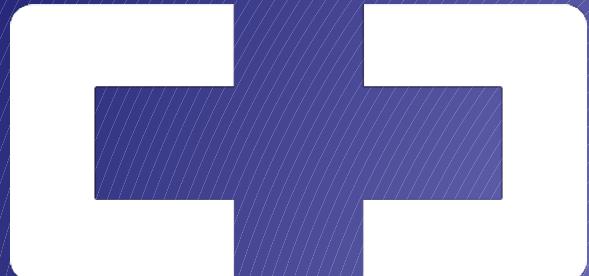


A Universal Framework for Sequential Decision Problems with Applications in Finance

Program in Digital Finance
University of Twente

February 4, 2024



Warren B Powell
Chief Innovation Officer, Optimal Dynamics
Professor Emeritus, Princeton University
Executive-in-Residence, Rutgers University

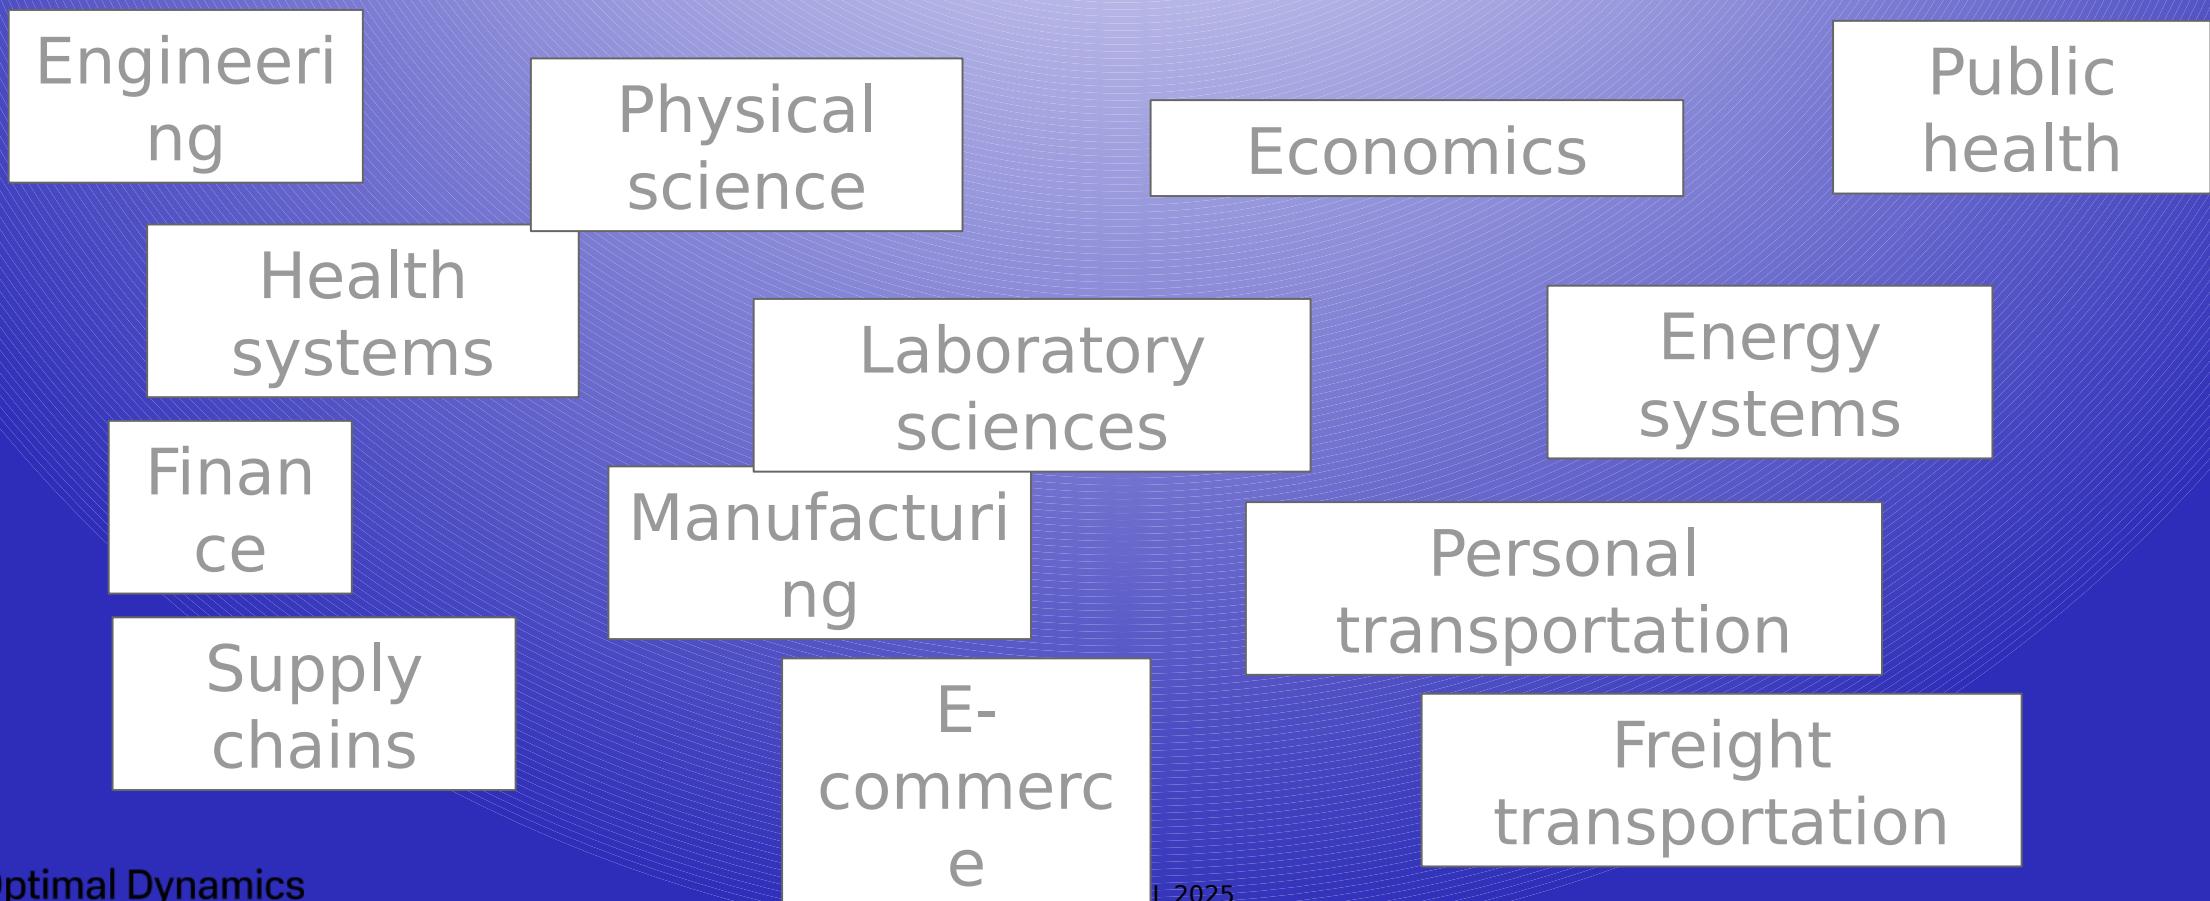


R



CHALLENGES

Virtually every problem in the domain of human processes combines decisions and uncertainty



Financial applications

Portfolio
balancing

Designing
hedging
contracts

Supply
chain
finance

Pricing
options

Pricing for e-
commerce

Managing cash
balances

Tuning trading
policies

Liquidating
assets

Lines of
credit

Currency
hedges

GOALS & OBJECTIVES

- » Increase profits
- » Reduce costs
- » Increase yield
- » Maximize productivity
- » Minimize waste
- » Maximize revenue
- » Increase strength
- » Reduce risk
- » Reduce carbon footprint
- » Reduce illness

Improving performance

If you want to run a better
to make better decisions. you have

- What decisions do you have to make?
 - How do we make better decisions?

OUTLINE

- The seven levels of artificial intelligence
- The universal modeling framework
- Designing policies
- Mutual fund cash balance optimization
- Choosing the best policy
- A new educational field: sequential decision analytics

OUTLINE

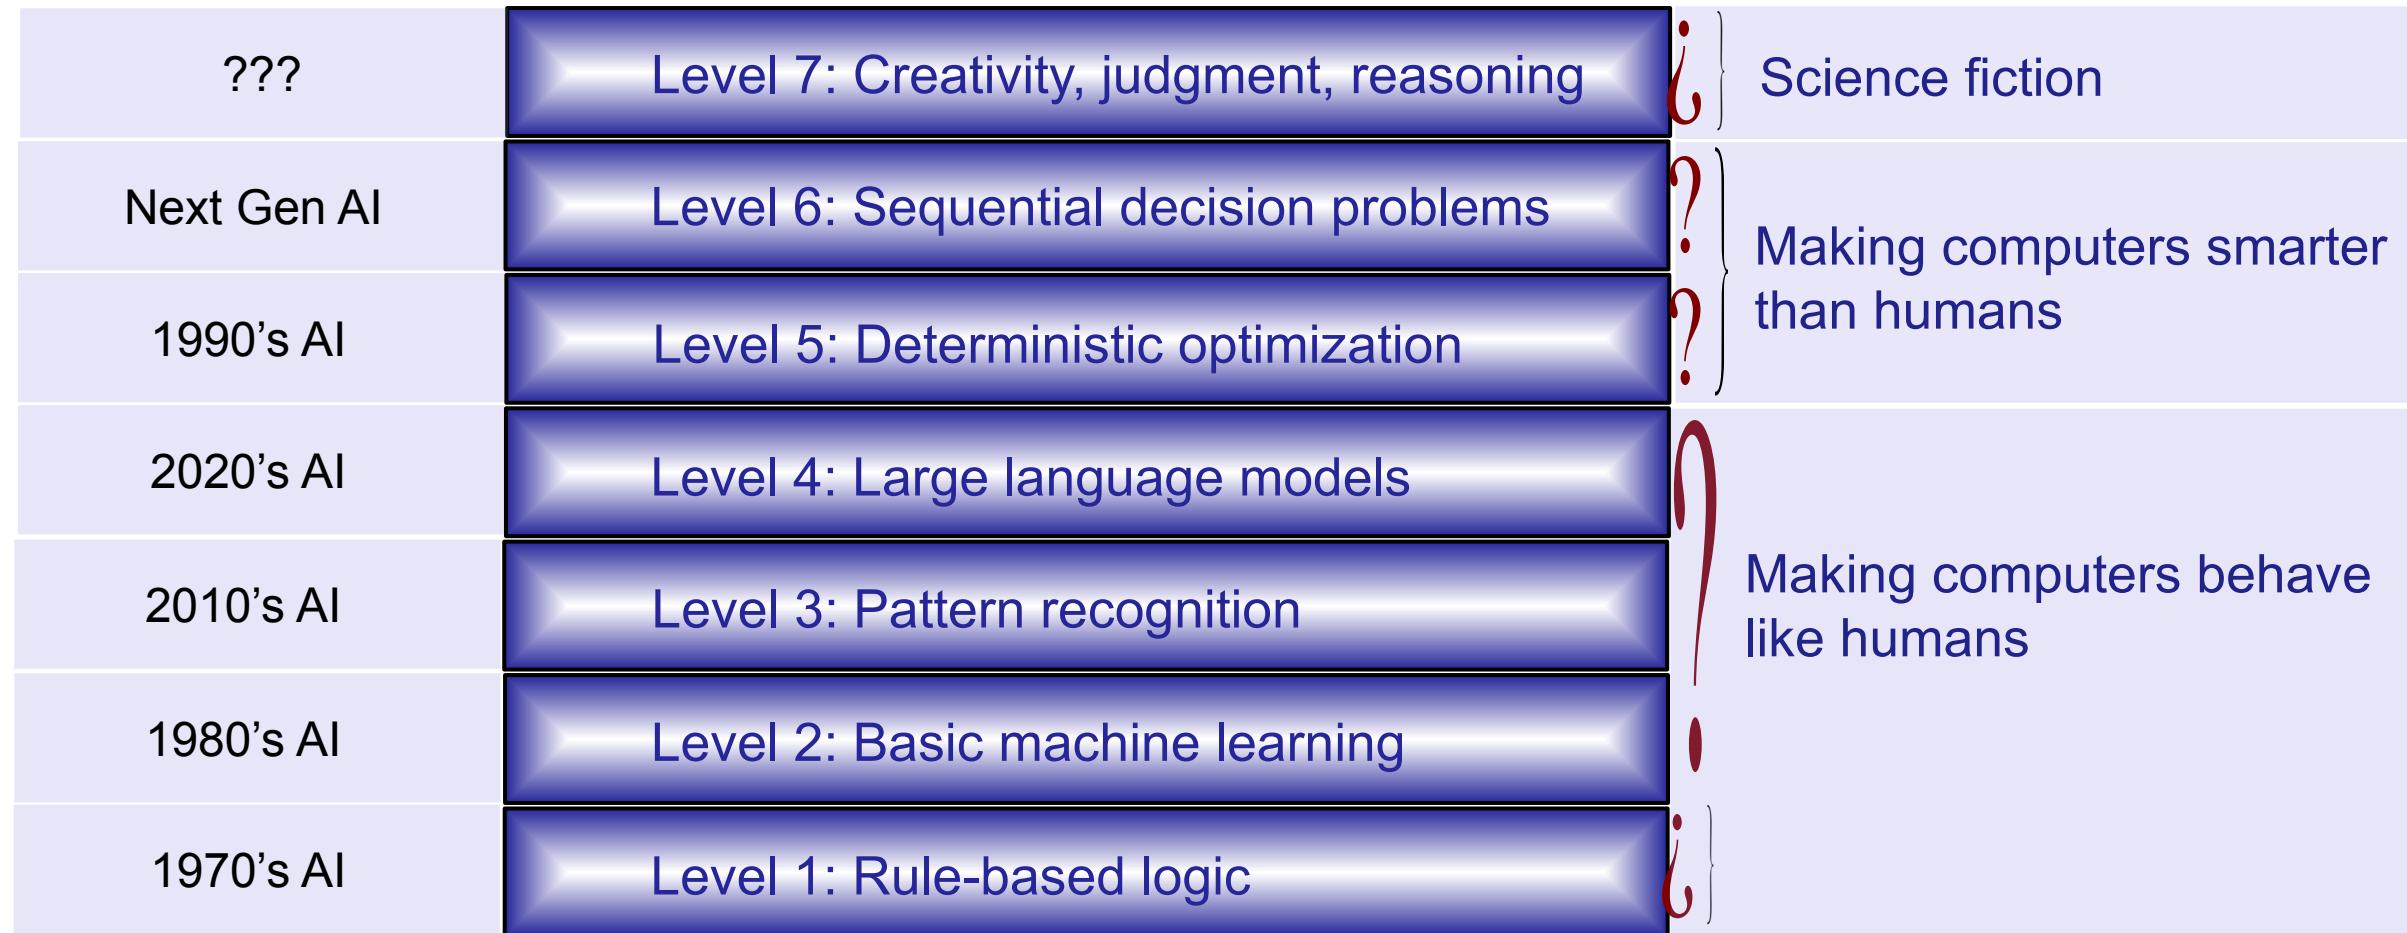
- The seven levels of artificial intelligence
- The universal modeling framework
- Designing policies
- Mutual fund cash balance optimization
- Choosing the best policy
- A new educational field: sequential decision analytics

Seven levels of artificial intelligence



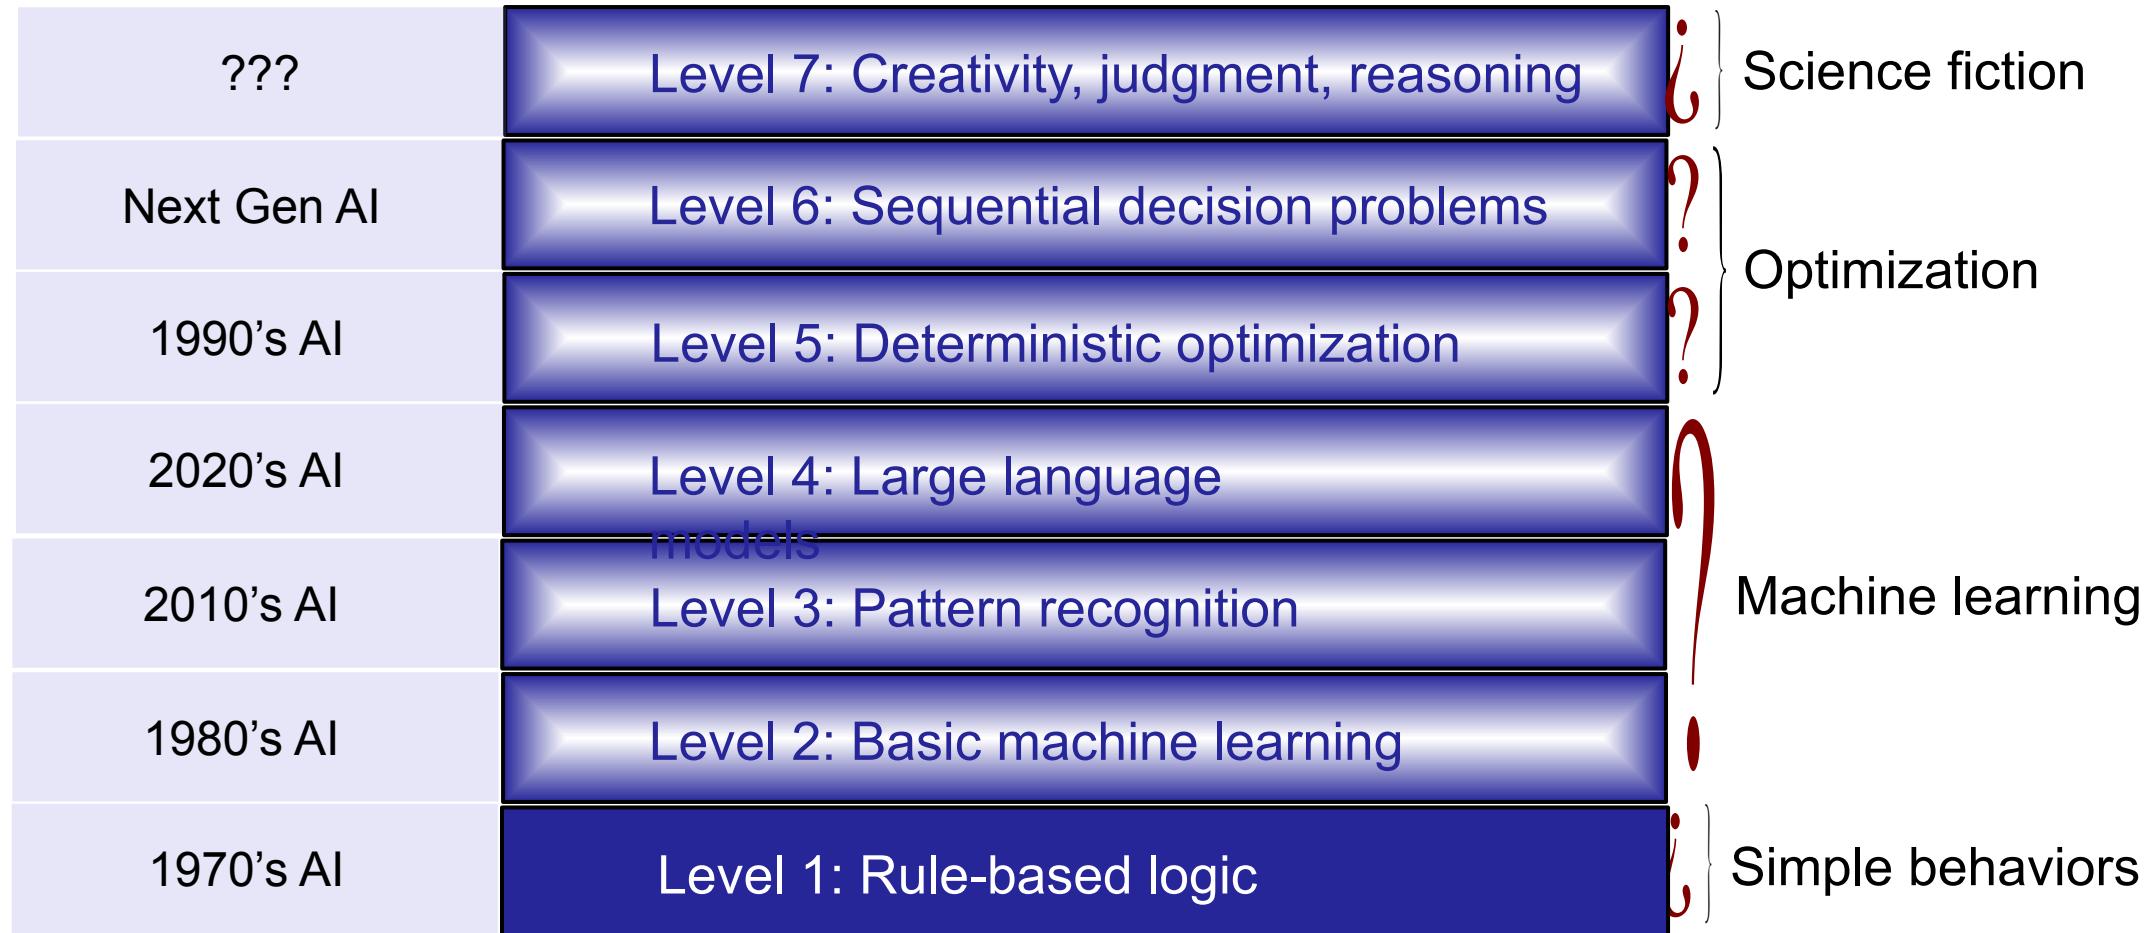
<https://tinyurl.com/7levelsofAI/>

Seven levels of artificial intelligence



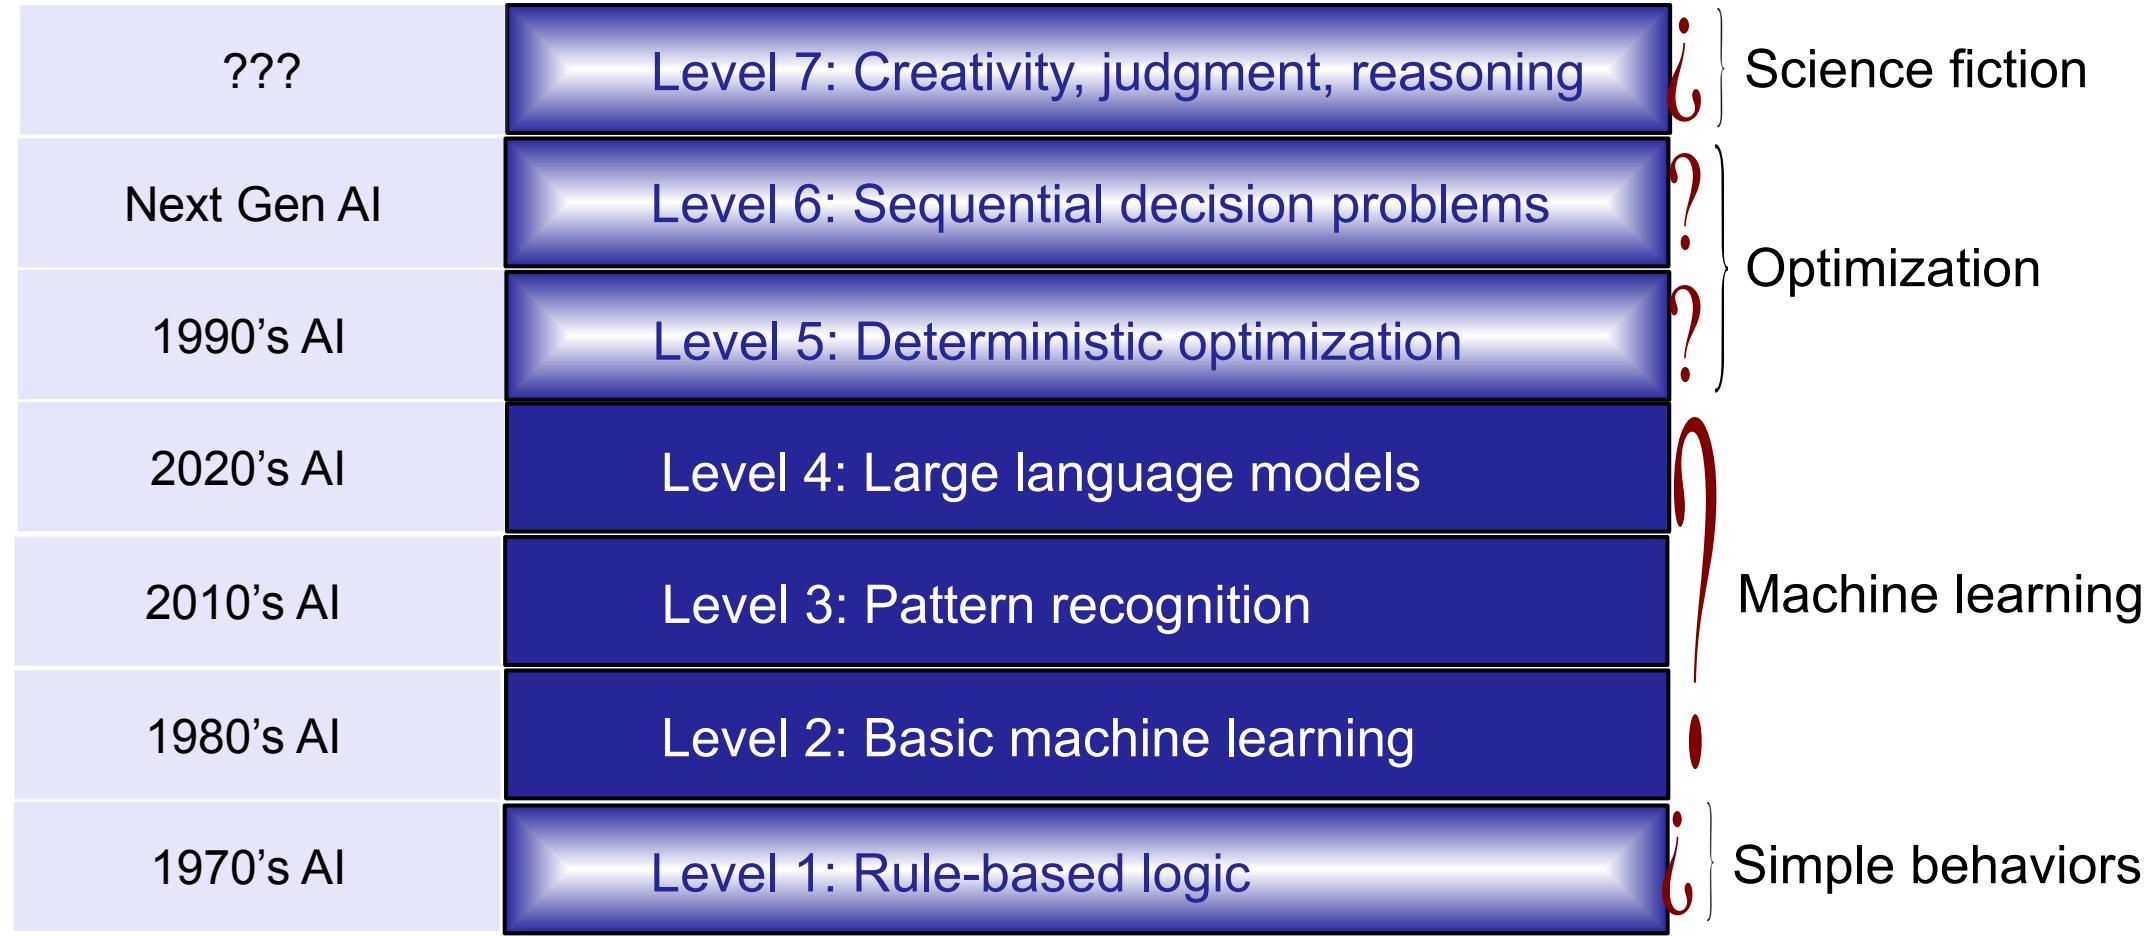
<https://tinyurl.com/7levelsofAI/>

Seven levels of artificial intelligence



<https://tinyurl.com/7levelsofAI/>

Seven levels of artificial intelligence



<https://tinyurl.com/7levelsOfAI/>

Level 2

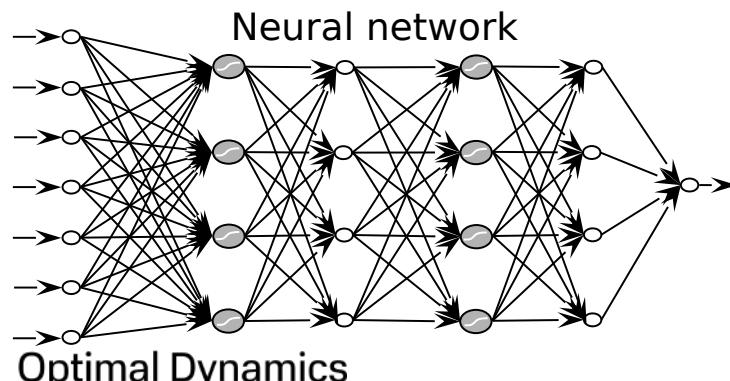
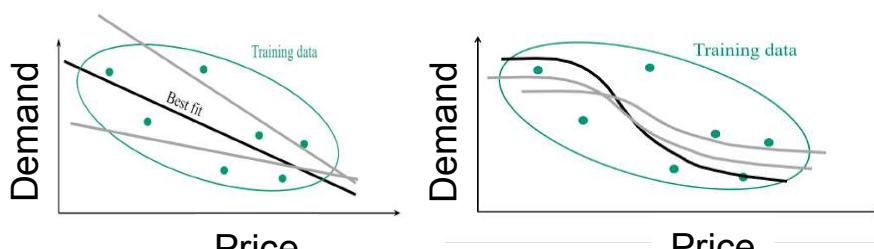
Basic machine learning (1900s)

➤ Sample applications

- Estimation
- Inference
- Forecasting

➤ Major classes of methods

- Parametric models
- Time-series models
- Early neural networks
(10k-1m parameters)



Level 3

Pattern recognition (2010s)

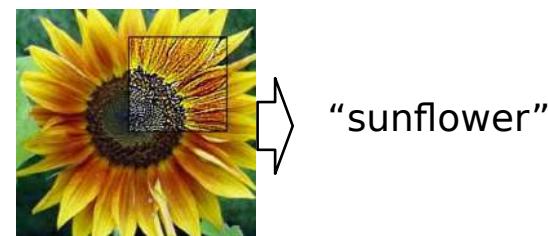
➤ Image recognition

➤ Voice recognition

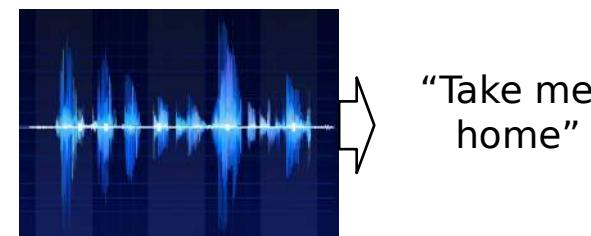
➤ Music identification

➤ Deep neural networks

- 10m-100m parameters
- First need for large labelled datasets



"sunflower"



"Take me home"

Level 4

Large language models (2020s)

➤ Text extension

- From seed input to entire paragraphs or documents

➤ Natural language processing

➤ Colossal neural networks

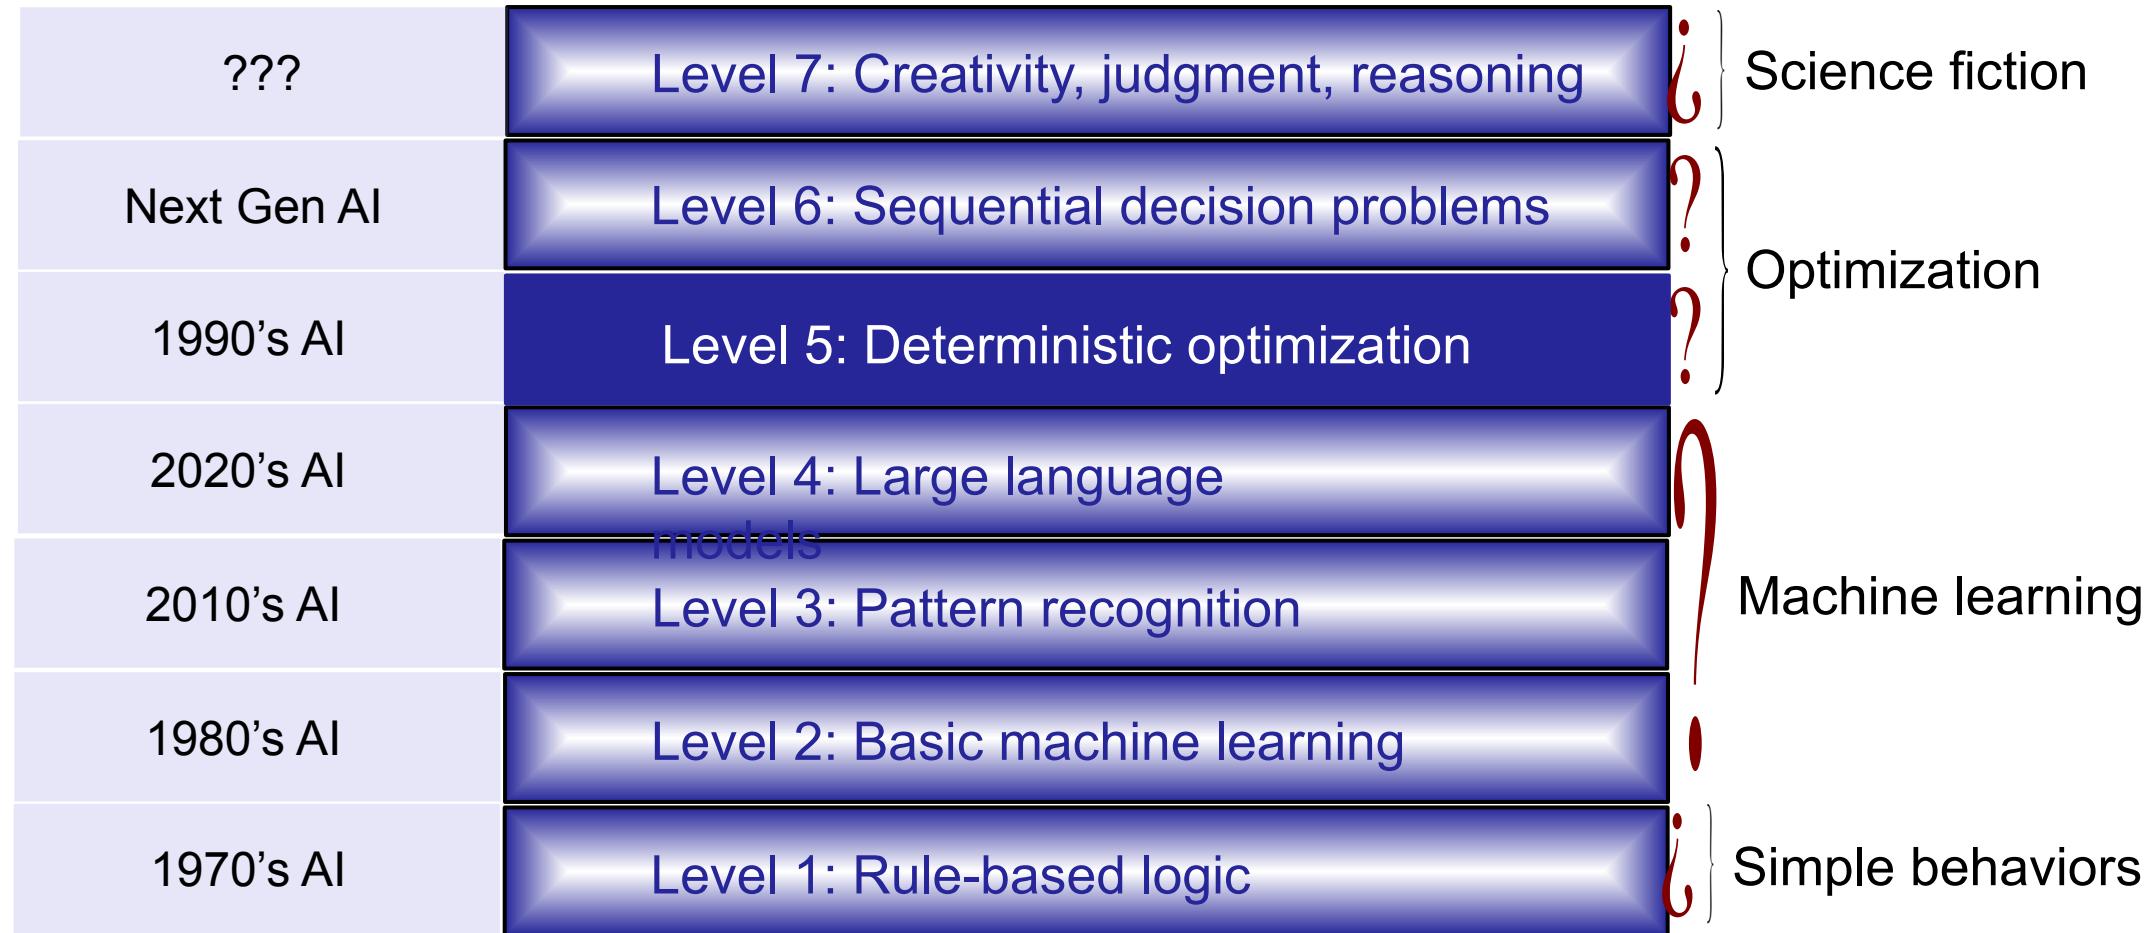
- 10 billion - 1 trillion parameters
- Trained on massive text datasets

Jack went to university

Jack went to the university in the early 1970s as a graduate student and started the company that would become McDonald's. Though he was a long time away from his family, he was an avid cook. "My mom is a huge foodie," he says. "She likes to know things about what I'm eating." Her recipes helped him develop an idea of how to go about the food business and eventually made him a fast-food millionaire in 1993. "I made sure that my family was always watching," he says. "And I always liked to tell them about everything I was doing on a regular basis." That early obsession with recipes led to his own idea of what a McDonald's menu would look like: a menu that would make people say hello to their food, even though they've been away.

Model output

Seven levels of artificial intelligence



<https://tinyurl.com/7levelsofAI/>



Decisions

What price to accept for a spot load?

Which load to accept now to move next week

Where should drivers be domiciled?

When should I refill the customer's tank with liquid nitrogen

Which customer tanks should we fill when we are in the area

Which material handling jobs should be done by robots?

When should inventory be refilled at a fulfillment center?

Which driver should move a load?

How many dedicated drivers should we have?

Which physician should handle a procedure?

How many syringes should be sent to each vaccination site, and when?

How many nurses should we have to serve local hospitals?

Which nurse should visit this doctor's office today?

Where should a patient be assigned for specific treatment?

What bid should we place on Google for a set of ad-words?

Which fulfillment center should handle an order?

What is the best policy for high-frequency trading?

What is the value of a financial option?

How much battery storage is needed?

When should gas turbines be scheduled?

How many suppliers should you have for a particular part, and where?

How much energy should I purchase from the wind farm?

Which supplier should manufacture turbine blades?

How many jet engines should be made each day?

What contracts to sign for raw materials?

Which vendor should supply each part?

When should inventory be ordered?

What price should be charged

DECISIONS

Types of decisions.

Physical Decisions



- » Managing inventories
- » Scheduling nurses, assigning drivers and moving trucks
- » Locating facilities and choosing suppliers

Financial Decisions



- » Buying/selling assets
- » Purchasing hedges
- » Managing loans
- » Pricing assets

Informational Decisions

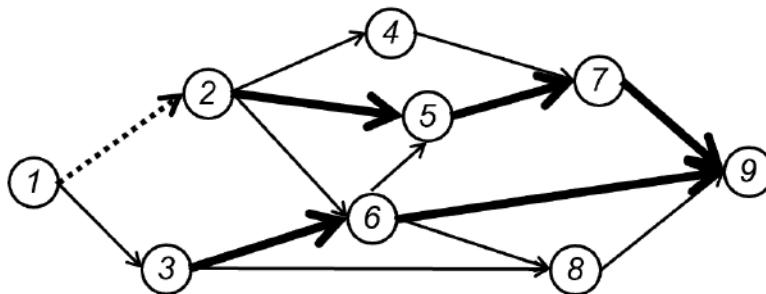


- » Purchasing reports
- » Testing new trading policies
- » Trying new traders
- » Investing in startups

DETERMINISTIC OPTIMIZATION

Planning a path to your destination

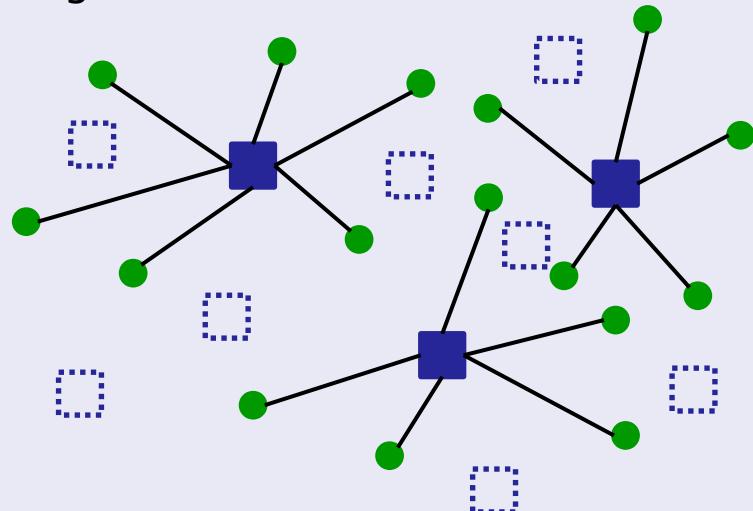
Low dimensional decisions



$$x_{ij} = \begin{cases} 1 & \text{If we move from node } i \text{ to node } j \\ 0 & \text{Otherwise} \end{cases}$$

Optimizing facility locations

High dimensional decisions



$$x_i = \begin{cases} 1 & \text{If we locate a facility at location } i \\ 0 & \text{Otherwise} \end{cases}$$

DETERMINISTIC OPTIMIZATION

Airline scheduling

Airlines



Optimization Model

$$\begin{aligned} & \min_x cx \\ & Ax = b \\ & x \geq 0 \end{aligned}$$

Airline Schedule

RTH	1120	[1]	109	P	6	1450	LGH	LGH	1615	D1	6966	1840	RLC	1940	[1]	6967	2215	LGH
RTH	1155	1214	189	P	1544	1552	LGH	LGH	1655	6966	1915	RLC	2022	[1]	6967	2230	LGH	
[3]	6814	235			1435	SSH	SSH	1555	[7]	6815	235		2125	LGH				
	6814	321	2		1488	1487	SSH	SSH	1545	1555		6815	224		2125	LGH	LGH	
RVN	1100		[X]	9511	176		1540	MAD	1640	511	T	1850	LTH					
RVN	1117	1125		9511	171		1612	1617	MAD	1785	511	T	LTH					
FRO	1130	[1]	2139		1425	MAD	MAD	1540	[1]	2706	1825	RGP	1925	[1]	2707	2215	MAD	
FRO	1130	1205	2139		1440	MAD	MAD	1550	1558	2706	1837	RGP	RGP		2025	2707	2259	
						MAD	MAD	1400	[1]	6652	326	1	2045	BRH	BRH	2238		
						1411	1418			6652	330		2024	BRH				
1035	[2]	4509	1315	LGH			LGH	1600	[4]	4746	1845	FAO	1945	[2]	4747	2215	LGH	
RGP	1135	4589	1402	LGH			LGH	1558	1612	4746	1843	FAO	1945	4747	217	2215	LGH	
1105	RCE	RCE	1230	[09]	4303	361		1640	MAD	MAD	1820			[1]	4330	337	2	
1117	RCE	RCE	1235	1245	4303	1621	1628	MAD	MAD	1830	1839				4330	336		
RGP	1135	[1]	575		1420	MAD	MAD	1540	D1	592	1820	RLC	1920	[7]	593	2155	MAD	
RGP	1248	1256	575		1530	MAD	MAD	1634	1640	592	1907	RLC	RLC	2024	593	2248		
RGP	1145	[1]	013		1425	LTH	LTH	1540	[1]	026	1800	RLC	1908	[4]	027	2138	LTH	
RGP	1268	013	1430		LTH			1540	026	1755	RLC	1895	1905	027	2138	LTH		

Airlines around the world use tools that depend on this mathematical model to perform strategic and operational planning.

Language of deterministic optimization

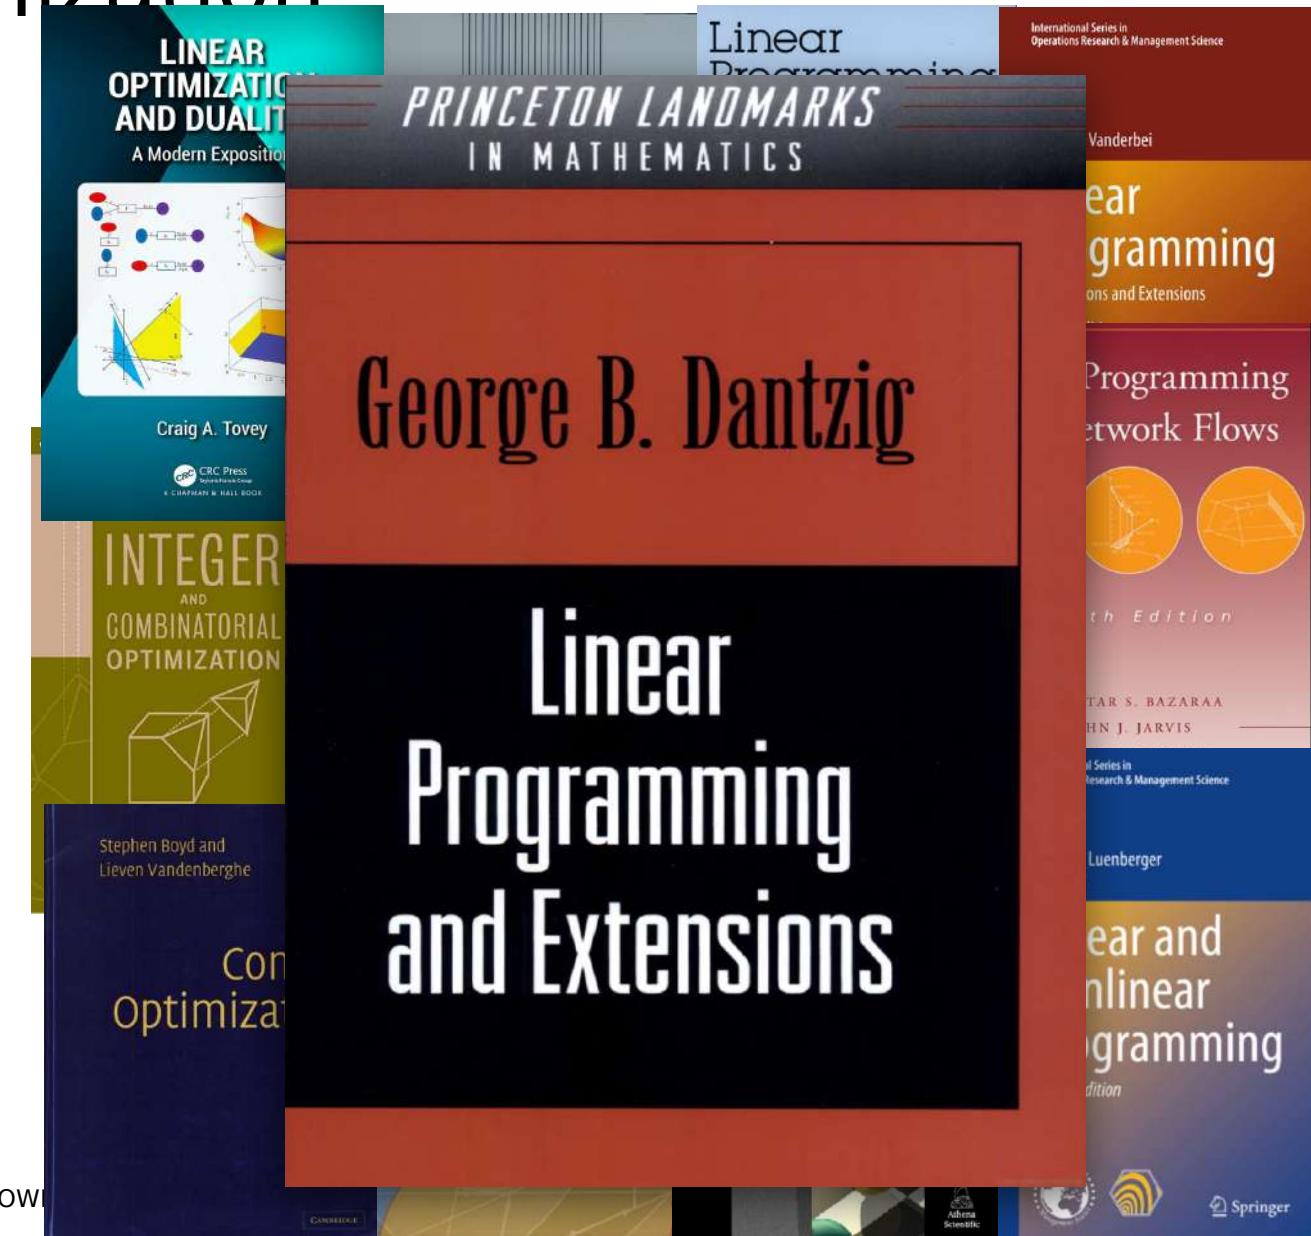
The language of deterministic optimization

$$\min_x c^T x$$

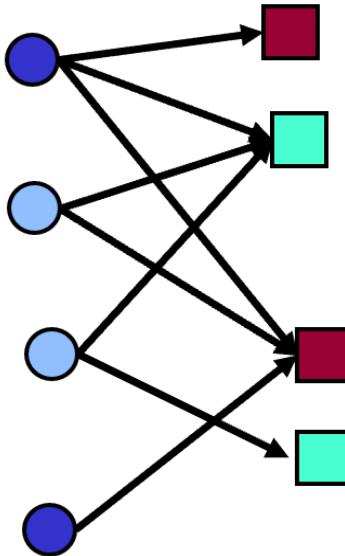
$$Ax = b$$

$$x \geq 0$$

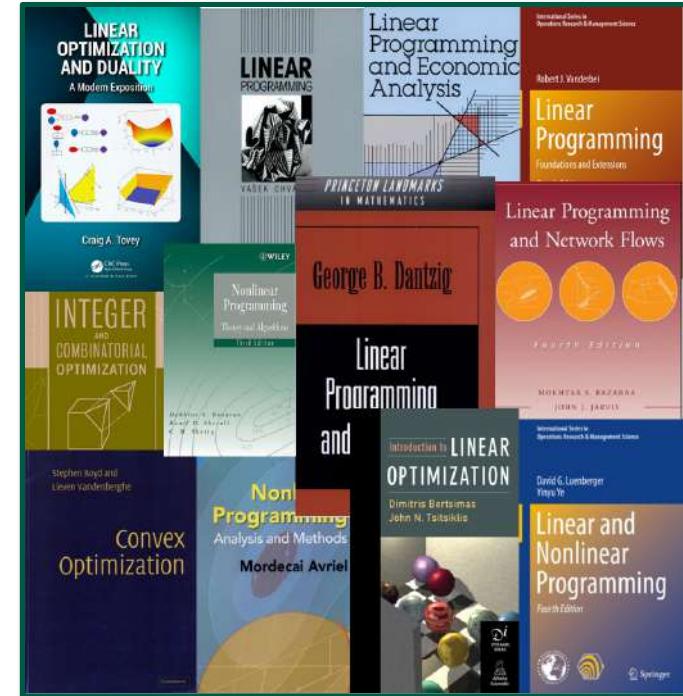
- » Spoken around the world.
- » Many books communicate the same core theory
- » Computer packages are available to solve realistic problems
- » Many graduate programs producing thousands of students each year.



Deterministic optimization



- Standard modeling strategy is to pose the optimization problem at a point in time.
- The next step is to design some algorithm (e.g. simplex) to solve the decision problem.
- We then declare that the solution is *optimal!*



The language of classical optimization does not capture anything related to the evolution of the system over time.





Decisions



Linear programs

Seven levels of artificial intelligence



<https://tinyurl.com/7levelsofAI/>



Decisions

What price to accept for a spot load?

Which load to accept now to move next week

Where should drivers be domiciled?

When should I refill the customer's tank with liquid nitrogen

Which customer tanks should we fill when we are in the area

Which material handling jobs should be done by robots?

When should inventory be refilled at a fulfillment center?

Which driver should move a load?

How many dedicated drivers should we have?

Which physician should handle a procedure?

How many syringes should be sent to each vaccination site, and when?

How many nurses should we have to serve local hospitals?

Which nurse should visit this doctor's office today?

Where should a patient be assigned for specific treatment?

What bid should we place on Google for a set of ad-words?

Which fulfillment center should handle an order?

What is the best policy for high-frequency trading?

What is the value of a financial option?

How much battery storage is needed?

When should gas turbines be scheduled?

How many suppliers should you have for a particular part, and where?

How much energy should I purchase from the wind farm?

Which supplier should manufacture turbine blades?

How many jet engines should be made each day?

What contracts to sign for raw materials?

Which vendor should supply each part?

When should inventory be ordered?

What price should be charged

Information

Market prices for spot freight

Offered loads by shipper, by lane

Driver application for jobs by region

Employment rate; unemployment filings

Patient arrivals and symptoms

Customer usage rate of liquid nitrogen and tank level.

Equipment failures at customer nitrogen tanks

Flow of different parts to each machining station

Flow of orders for a product by region around the country

Driver requests for loads

Changes in asset prices and trading volume

Production delays in order fulfillment

New COVID-19 cases by county; willingness to be vaccinated.

Requests for nurses from doctor's offices and hospitals.

Number of nurses calling in sick

Availability of specialists to treat a condition

Whether a bid wins an ad-click auction

Orders for a product from different regions

Prices of raw materials by region

Quality of orders provided by a vendor

Transportation and port unloading delays

Wind generation from a wind farm

Competitor prices and availability

Electricity prices on the grid and generator failures

Capacity shutdowns at suppliers due to labor or political problems

Lead times required by each manufacturer

The amount of energy that is generated from wind.

Daily production of new jet engines

SEQUENTIAL DECISIONS

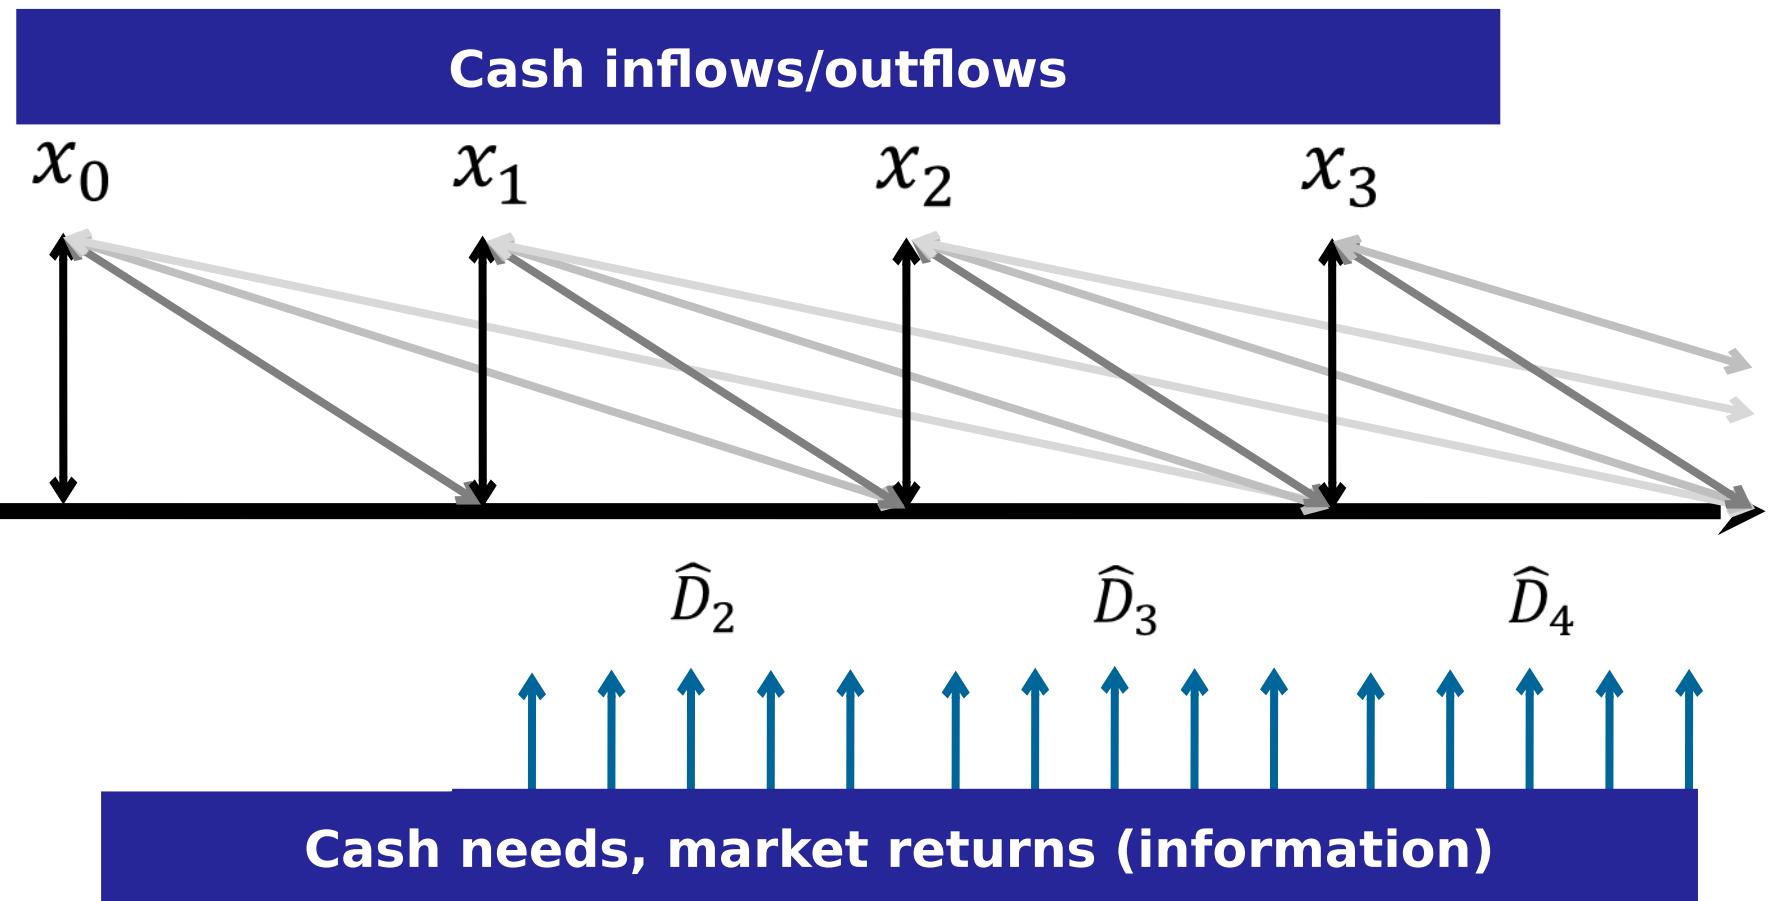
In most settings, decisions are made over time...

Decisions

Information that arrives after a decision is made is not known when we made the decision.

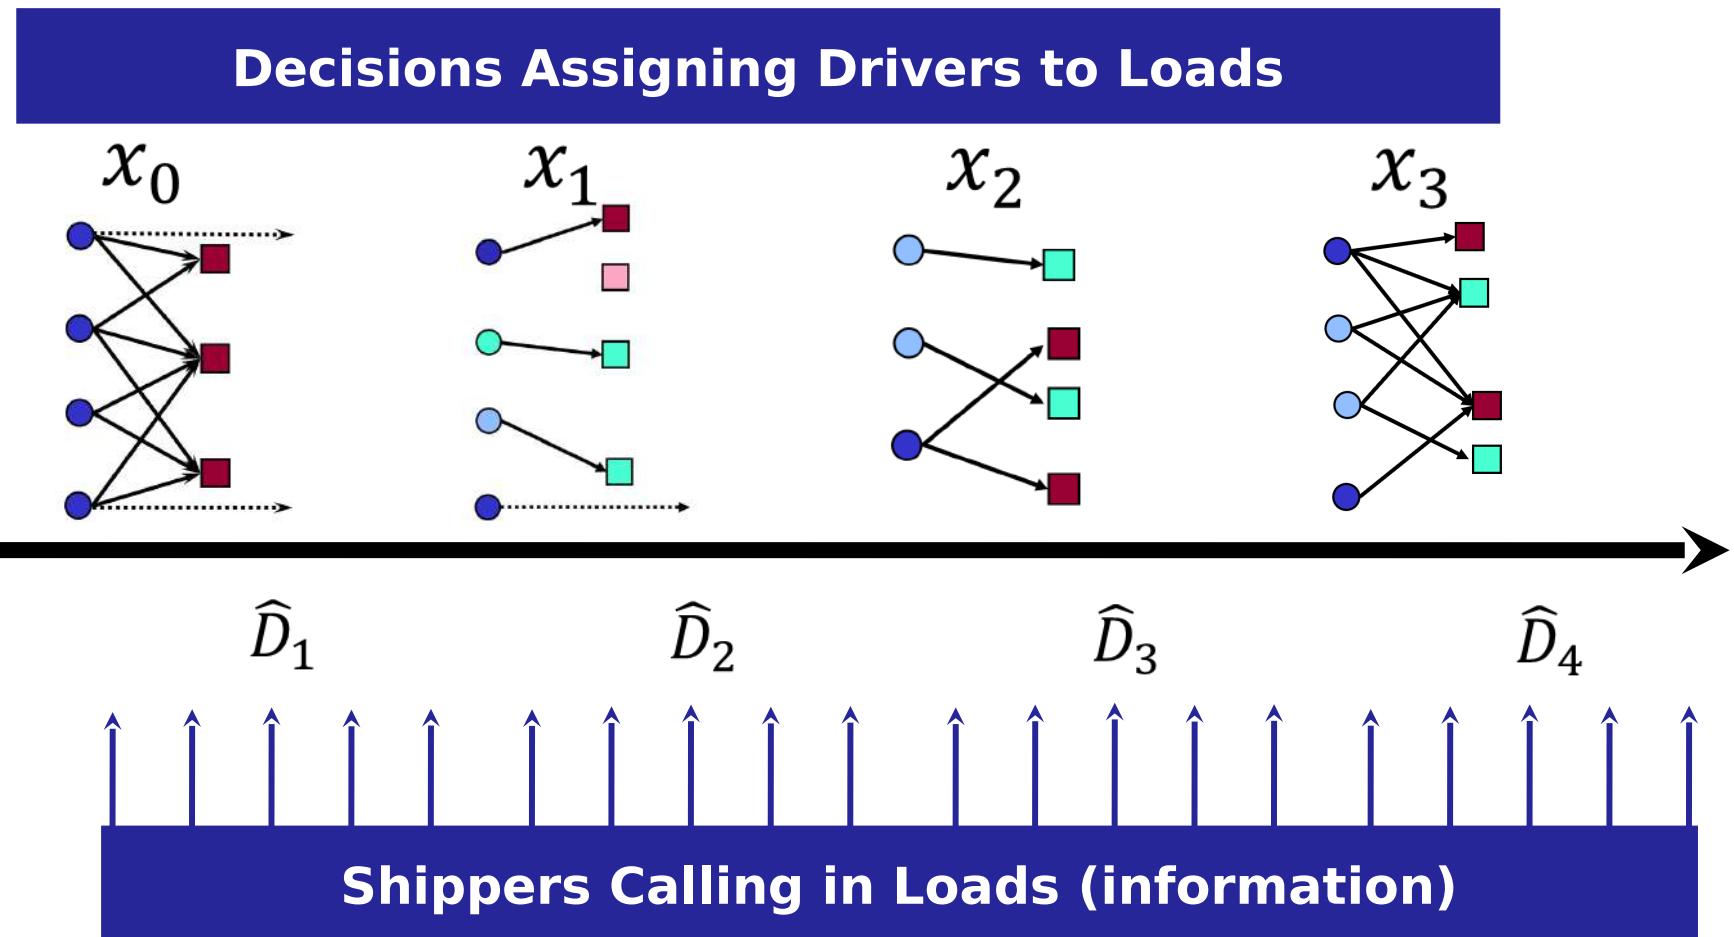
SEQUENTIAL DECISIONS

Cash management



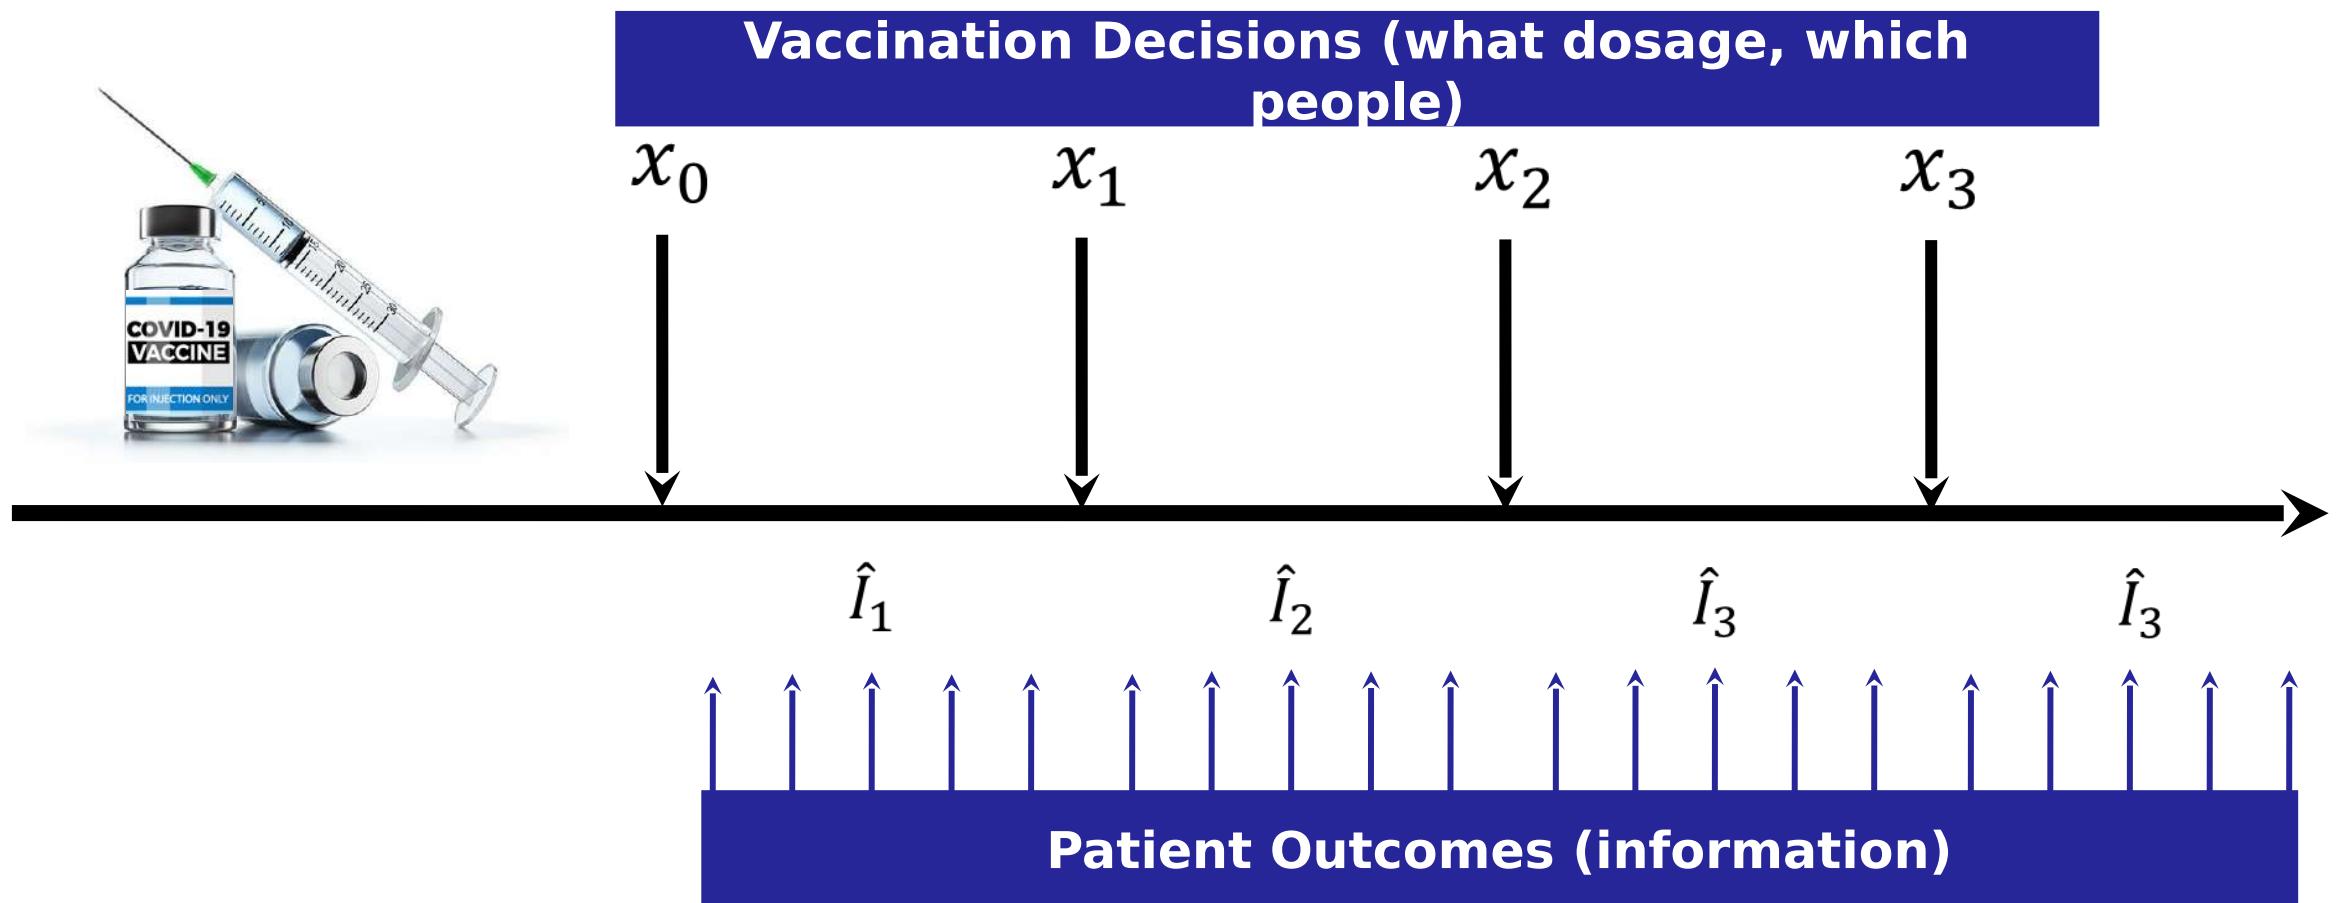
SEQUENTIAL DECISIONS

Driver dispatch for truckload trucking



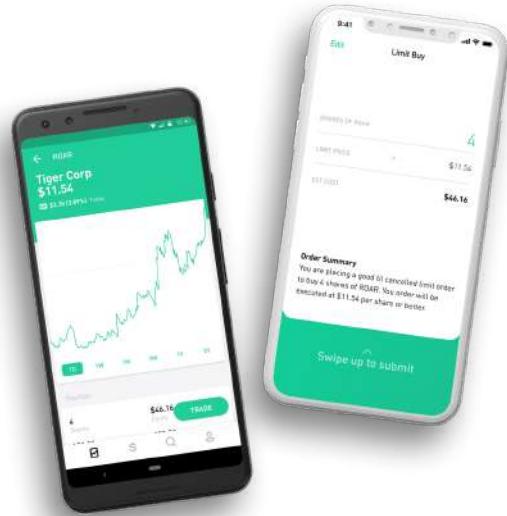
SEQUENTIAL DECISIONS

Testing new vaccines

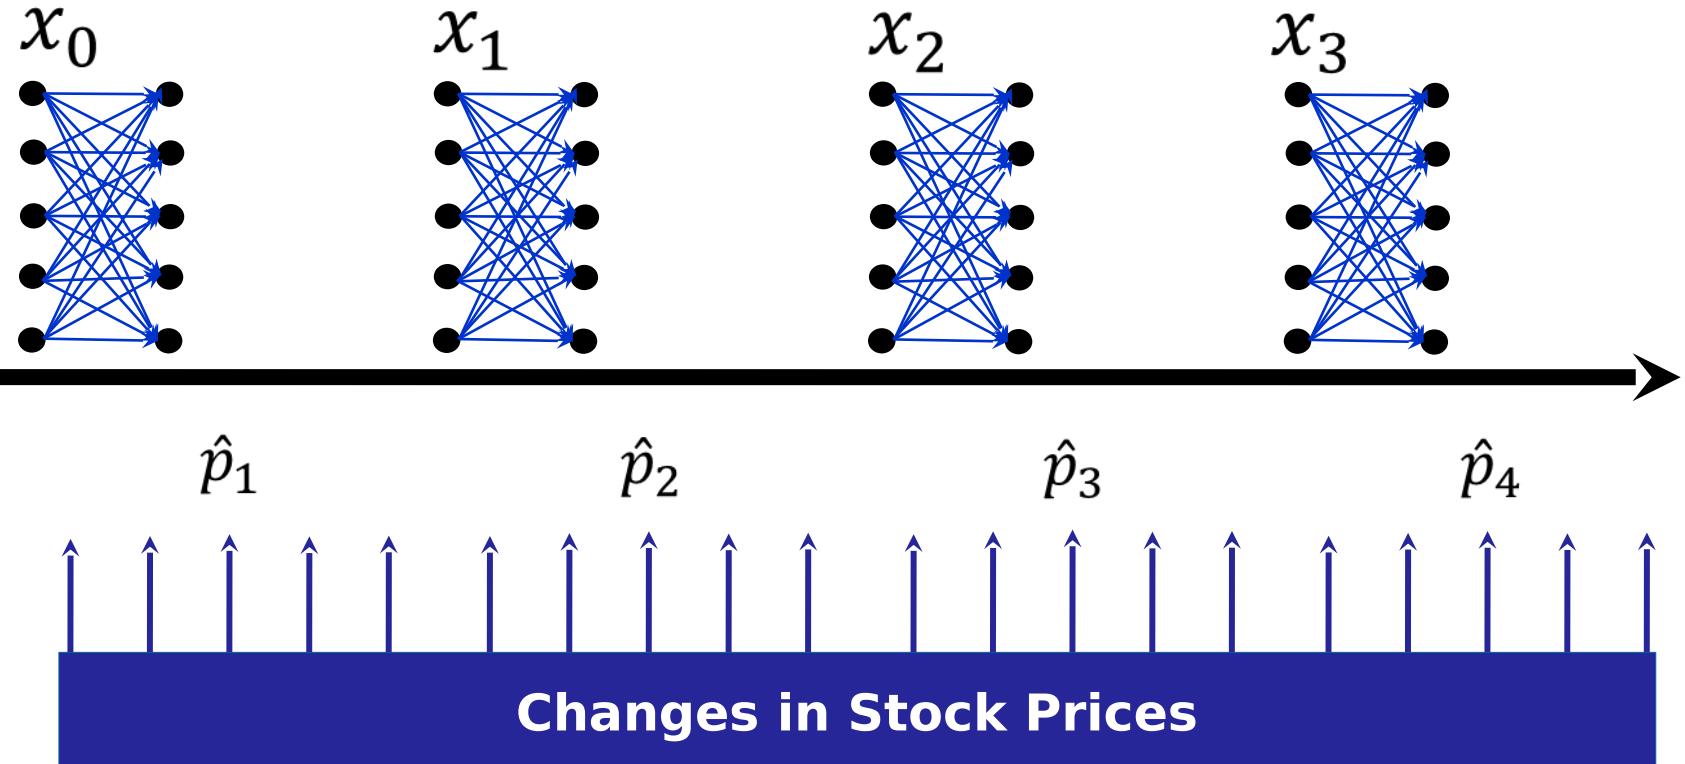


SEQUENTIAL DECISIONS

Financial Trading



Buy-sell decisions (what assets, how much)



OUTLINE

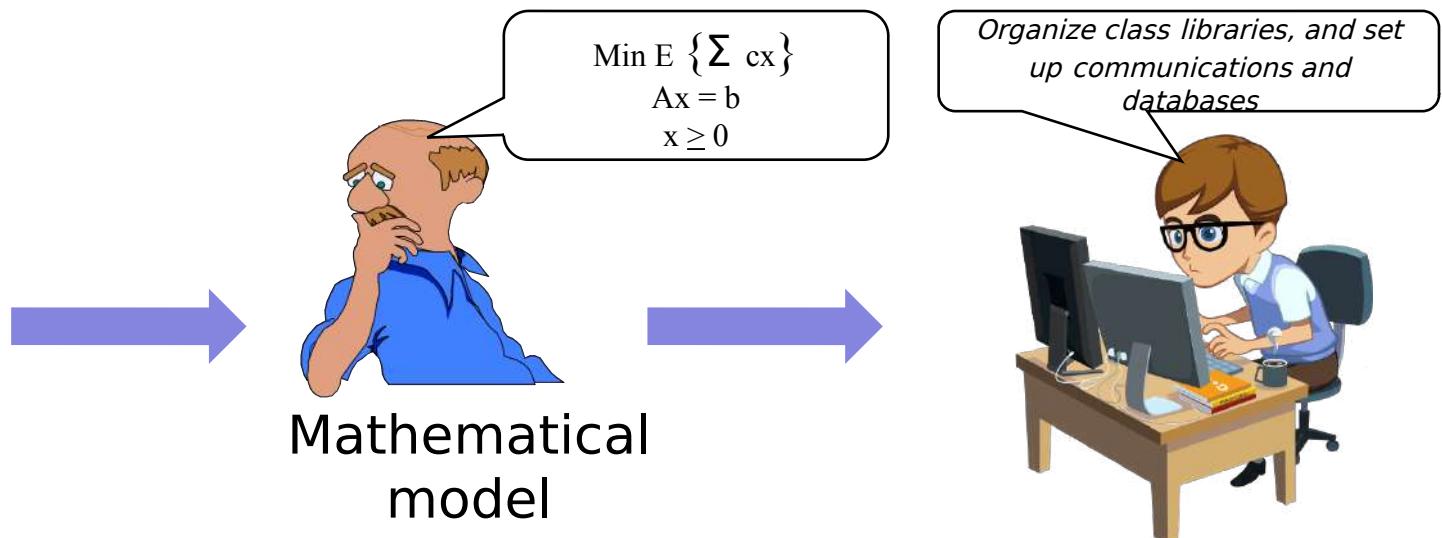
- The seven levels of artificial intelligence
- The universal modeling framework
- Designing policies
- Mutual fund cash balance optimization
- Choosing the best policy
- A new educational field: sequential decision analytics

Modeling sequential decision problems

The biggest challenge when making decisions under uncertainty is ***modeling***.



Everyone writing out a deterministic optimization model, or machine learning model, knows how to write out their problem mathematically...



...we lack a standard modeling framework for sequential decisions.



Stochastic
programming

Optimal
learning

Bandit
problems

Stochastic
control

Simulation
optimization

Model
predictive
control

Reinforcement
learning

Robust
optimization

Optimal
control

Markov
decision
processes

Decision
analysis

Dynamic
Programming
and
control

Stochastic
optimization

Approximate
dynamic
programming

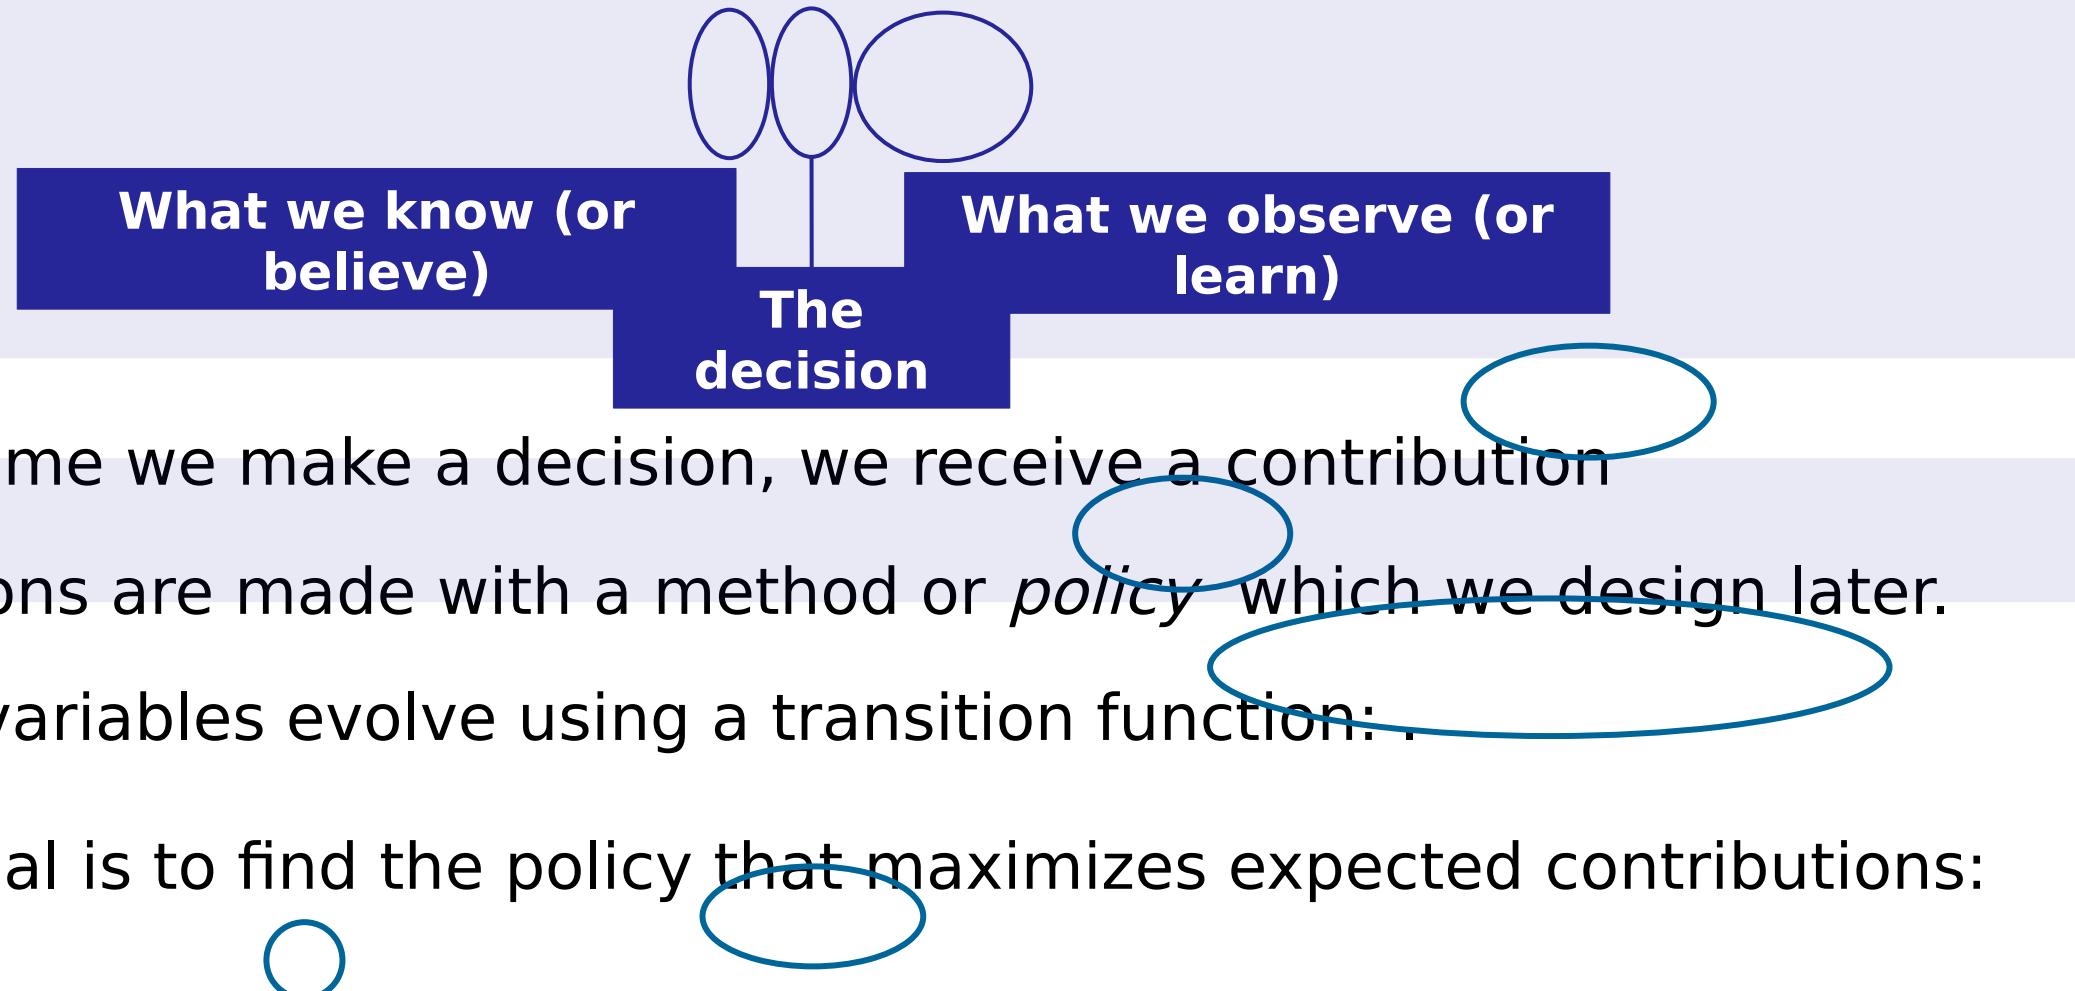
Stochastic
search

Online
computation



Modeling sequential decision problems

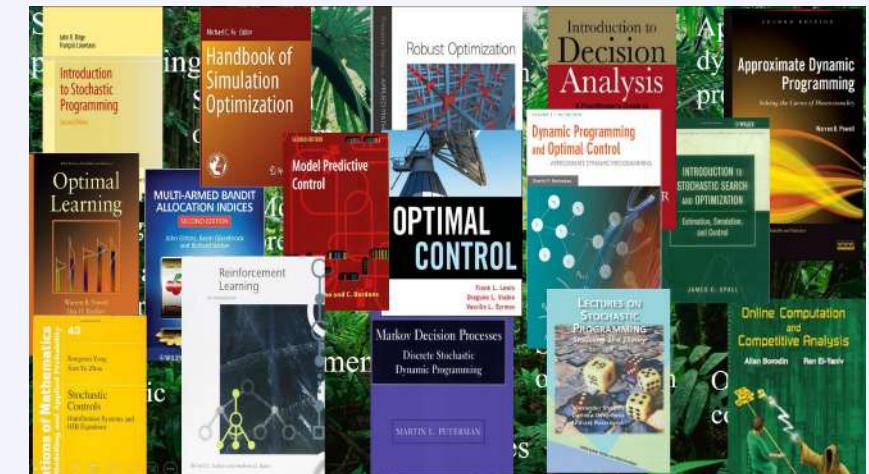
- Any sequential decision problem can be written:



Modeling sequential decision problems

Every sequential decision problem can be modeled using 5 core components

- 1) State variables
 - Physical state , other information , beliefs .
- 2) Decision variables (or action , or control)
 - Decisions are determined by a policy
- 3) Exogenous variables
 - This is new information that arrives between and .
- 4) Transition function
 - This is how our state variable evolves given and .
- 5) Objective function for finding the best policy



This is the universal modeling framework

Modeling sequential decision problems

- Objective functions
 - » Cumulative reward (“online learning”)

$$\max_{\pi} \mathbb{E} \left\{ \sum_{t=0}^T C_t(S_t, X_\square^\pi(S_t)) \vee S_0 \right\}$$

- » Final reward $r(x_t) \in \{r_e(\text{Learning}^N), r_{S_0}\}$

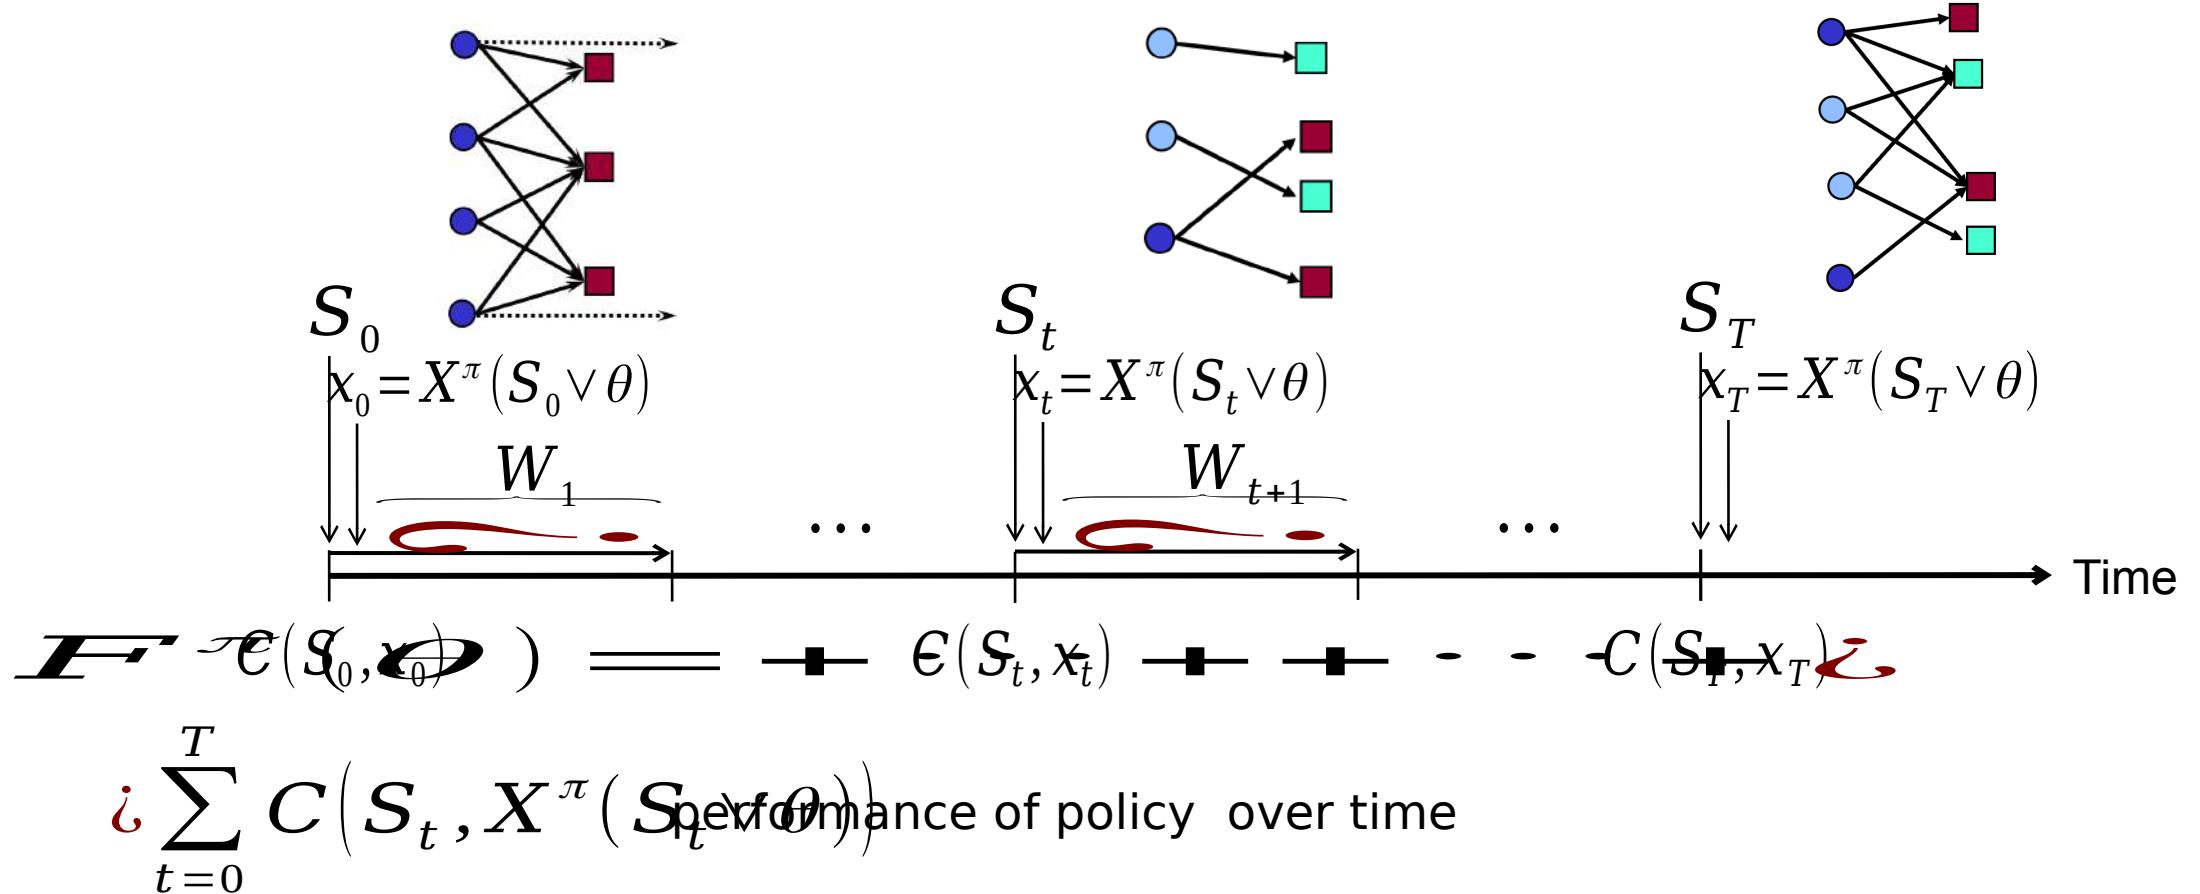
$$max_{\pi} \mathbb{E} \left\{ \sum_{t=0}^T C_t(S_t, X_{\square}^{\pi}(S_t)) \vee S_0 \right\} - \eta \varrho \left\{ C(S_0, X_0^{\pi}(S_0)), \dots, C(S_T, X_T^{\pi}(S_T)) \vee S_0 \right\}$$

— Base objective — Risk metric

- » Risk

Deterministic optimization

Evaluating performance over time:



Modeling sequential decision problems

The objective function

We might average over simulations:

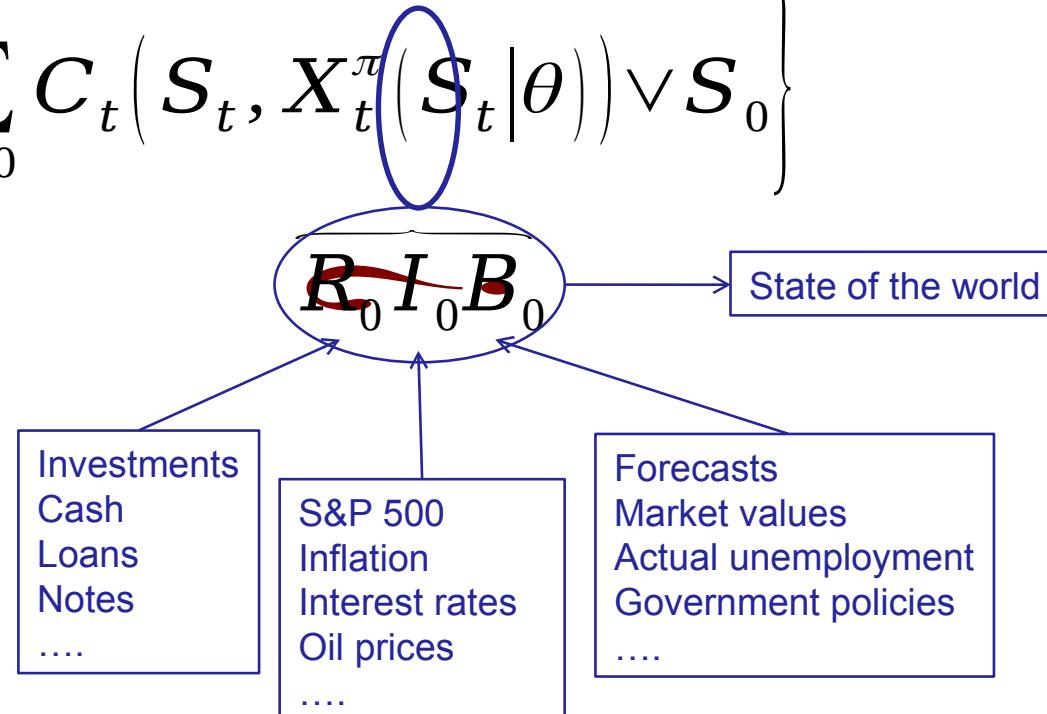
$$\bar{F}^\pi = \frac{1}{N} \sum_{n=1}^N \sum_{t=0}^T C_t(S_t(\omega^n), X_t^\pi(S_t(\omega^n) | \theta))$$

Modeling sequential decision problems

The objective function

.... which is an approximation of the expectation:

$$F^\pi = \mathbb{E} \left\{ \sum_{t=0}^T C_t(S_t, X_t^\pi(S_t | \theta)) \vee S_0 \right\}$$



Modeling sequential decision problems

The objective function

.... which is an approximation of the expectation:

$$F^\pi = \mathbb{E} \left\{ \sum_{t=0}^T C_t(S_t, X_t^\pi(S_t | \theta)) \right\} + S_0$$

Evaluating policies

Theoretically
Optimality proofs
Asymptotic optimality
Regret bounds

Computer simulation



Field testing



Modeling sequential decision problems

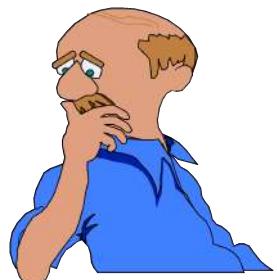
The objective function

We now want to find the best policy

$$\max_{\pi=(f \in F, \theta \in \Theta^f)} F^\pi = \mathbb{E} \left\{ \sum_{t=0}^T C_t(S_t, X_t^\pi(S_t | \theta)) \vee S_0 \right\}$$

Type of policy
(structure of function)

Algorithmic tuning



$$\theta^{n+1} = \theta^n + \alpha_n \nabla_\theta F(\theta^n, W^{n+1})$$

Modeling sequential decision problems

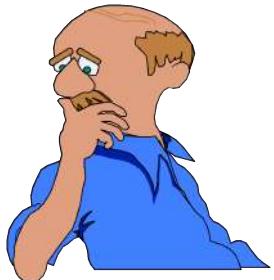
The objective function

We now want to find the best policy

$$\max_{\pi=(f \in F, \theta \in \Theta^f)} F^\pi = \mathbb{E} \left\{ \sum_{t=0}^T C_t(S_t, X_t^\pi(S_t | \theta)) \mid S_0 \right\}$$

Type of policy
(structure of function)

Algorithmic tuning



$$\theta^{n+1} = \theta^n + \alpha_n \nabla_\theta F(\theta^n, W^{n+1})$$

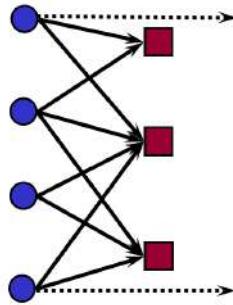
Evaluating policies

- » Solution quality
 - On average
 - Worst case performance
- » Computational
 - Speed
 - Reliability
- » Transparency
- » Methodological complexity
- » Data requirements

Modeling sequential decision problems

If our policy involves solving an optimization problem, we need to recognize two optimization problems:

The surrogate optimization model – used for making decisions



$$x_t = X^\pi(S_t \vee \theta)$$

The objective function, used for evaluating policies.

$$F^\pi(\theta) = \mathbb{E} \left\{ \sum_{t=0}^T C_t(S_t, X_t^\pi(S_t | \theta)) \vee S_0 \right\}$$

An optimal solution to the surrogate optimization model is almost never an optimal policy, and may perform quite poorly.

The problem solving process

A
p
p
l
i
c
a
t
i
o
n

Step 1

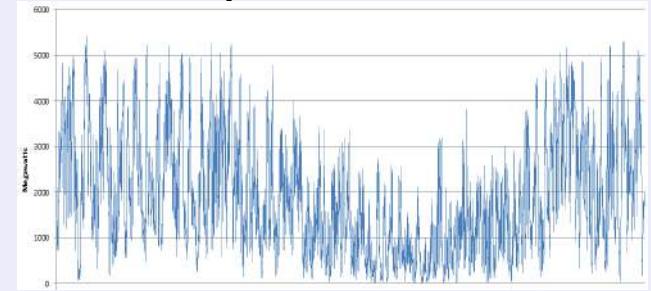
Identify:

1. Performance metrics
2. Types of decisions
3. Sources of uncertainty

Step 2: Mathematical model:

- » State variables $S_t = (R_t, I_t, B_t)$
 - Physical state R_t , other information I_t , belief state B_t .
- » Decision variables (x_t, a_t, u_t)
 - Made with policy $X^\pi(S_t|\theta)$ (or $A^\pi(S_t)$ or $U^\pi(S_t)$)
- » Exogenous information W_{t+1}
 - What do we learn for the first time between t and $t + 1$?
- » Transition function $S_{t+1} = S^M(S_t, x_t, W_{t+1})$
 - How do the state variables evolve over time?
- » Objective function
 - $\max_{\pi} \mathbb{E}_{S_0} \mathbb{E}_{W_1, \dots, W_T | S_0} \sum_{t=0}^T C(S_t, X^\pi(S_t))$

Step 3: Uncertainty modeling



Step 4: Designing policies

$$\max_{\pi} \mathbb{E} \left\{ \sum_{t=0}^T C(S_t, X^\pi(S_t)) \mid S_0 \right\}$$

Step 5: Computer model



Step 6: Implementation/analysis



The problem solving process

A
p
p
l
i
c
a
t
i
o
n

Step 1

Identify:

1. Performance metrics
2. Types of decisions
3. Sources of uncertainty

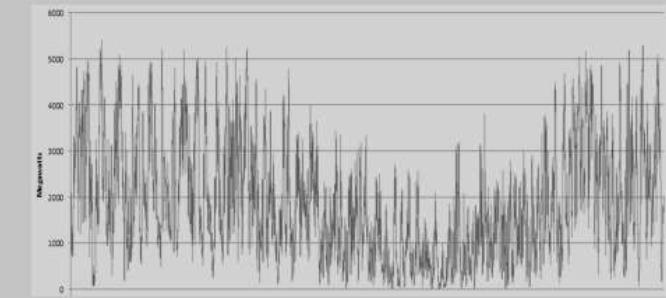
Step 2:

Mathematical model:

- » State variables $S_t = (R_t, I_t, B_t)$
 - Physical state R_t , other information I_t , belief state B_t .
- » Decision variables (x_t, a_t, u_t)
 - Made with policy $X^\pi(S_t|\theta)$ (or $A^\pi(S_t)$ or $U^\pi(S_t)$)
- » Exogenous information W_{t+1}
 - What do we learn for the first time between t and $t + 1$?
- » Transition function $S_{t+1} = S^M(S_t, x_t, W_{t+1})$
 - How do the state variables evolve over time?
- » Objective function
 - $\max_{\pi} \mathbb{E}_{S_0} \mathbb{E}_{W_1, \dots, W_T | S_0} \sum_{t=0}^T C(S_t, X^\pi(S_t))$

Step 3:

Uncertainty modeling



Step 4:

Designing policies

$$\max_{\pi} \mathbb{E} \left\{ \sum_{t=0}^T C(S_t, X^\pi(S_t)) | S_0 \right\}$$

Step 5:

Computer model



Step 6:

Implementation/analysis



The problem solving process

Step 1

If you want to run a better you have to
make better decisions.

The path to improvement starts by answering three questions:

1. What are the performance metrics?
2. What decisions do you have to make (and who makes them)?
3. What are the sources of uncertainty?

Download from <https://tinyurl.com/DoingBetterppt/>

OUTLINE

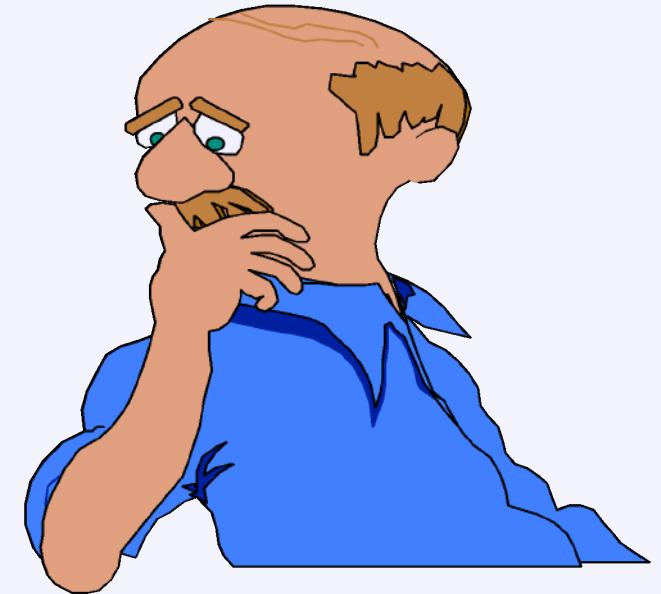
- The seven levels of artificial intelligence
- The universal modeling framework
- Designing policies
- Mutual fund cash balance optimization
- Choosing the best policy
- A new educational field: sequential decision analytics

Designing policies

What is a policy?

A policy is method that makes a decision using the information in the state variable.

... any method.



Designing policies

Policies and the English language

Algorithm	Format	Precedent
Behavior	Formula	Prejudice
Belief	Grammar	Principle
Bias	Habit	Procedure
Canon	Heuristics	Process
Code	Laws/bylaws	Protocols
Commandment	Manner	Recipe
Conduct	Method	Ritual
Control law	Mode	Rule
Convention	Mores	Strategy
Culture	Norm	Style
Customs	Orthodox	Syntax
Doctrine	Patterns	Technique
Duty	Plans	Template
Etiquette	Policies	Tenet
Fashion	Practice	Tradition
Way of life		

<http://tinyurl.com/policiesanddecisions>

BRIDGING MACHINE LEARNING & SEQUENTIAL DECISIONS

Machine
learning

$$\min_{f \in F, \theta \in \Theta^f} \frac{1}{N} \sum_{n=1}^N (y^n - f(x^n | \theta))^2$$

Searching over
functions

“Big
dataset”

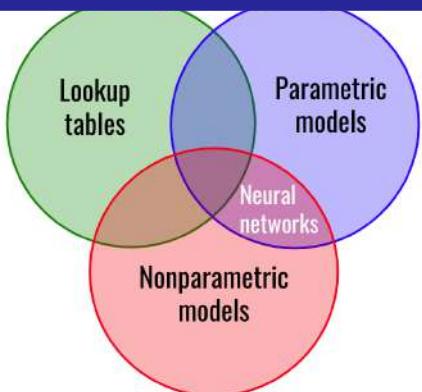
Sequential decisions

$$\max_{\pi = (f \in F, \theta \in \Theta^f)} \frac{1}{N} \sum_{n=1}^N \sum_{t=0}^T C(S_t^n, X^\pi(S_t^n | \theta))$$

$$S_{t+1} = S^M(S_t, x_t, W_{t+1})$$

Searching over
policies

System model



POLICIES:

Designing policies

There are two fundamental strategies for designing policies

Policy search - Search over a class of methods for making decisions to optimize some metric over time.

- » Finding the best class of policy.
- » Finding the best policy within the class.

Lookahead approximations - Approximate the impact of a decision now on the future.

- » The contribution from the first period, plus
- » An approximation of the sum of contributions in future time periods resulting from the first decision.

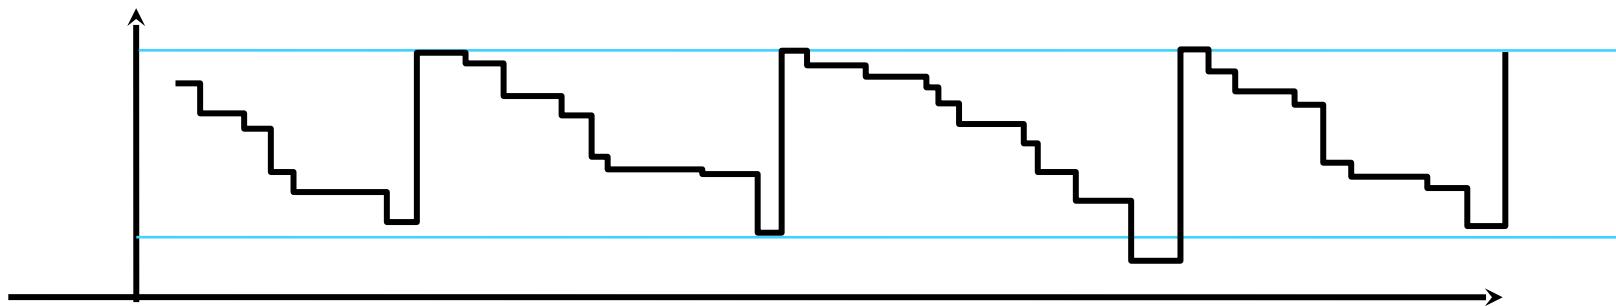
Policy search

1) Policy function approximation (PFA)

These are analytical functions that specify what to do given what we know.

Examples:

- a) Order-up-to inventory policy



- b) Buy when the price goes **below** and sell when it goes **above**
- c) Lookup tables, linear/nonlinear models, neural networks, nonparametric models, ... any function we might use in machine learning.

Policy search

2) Cost function approximations (CFAs)

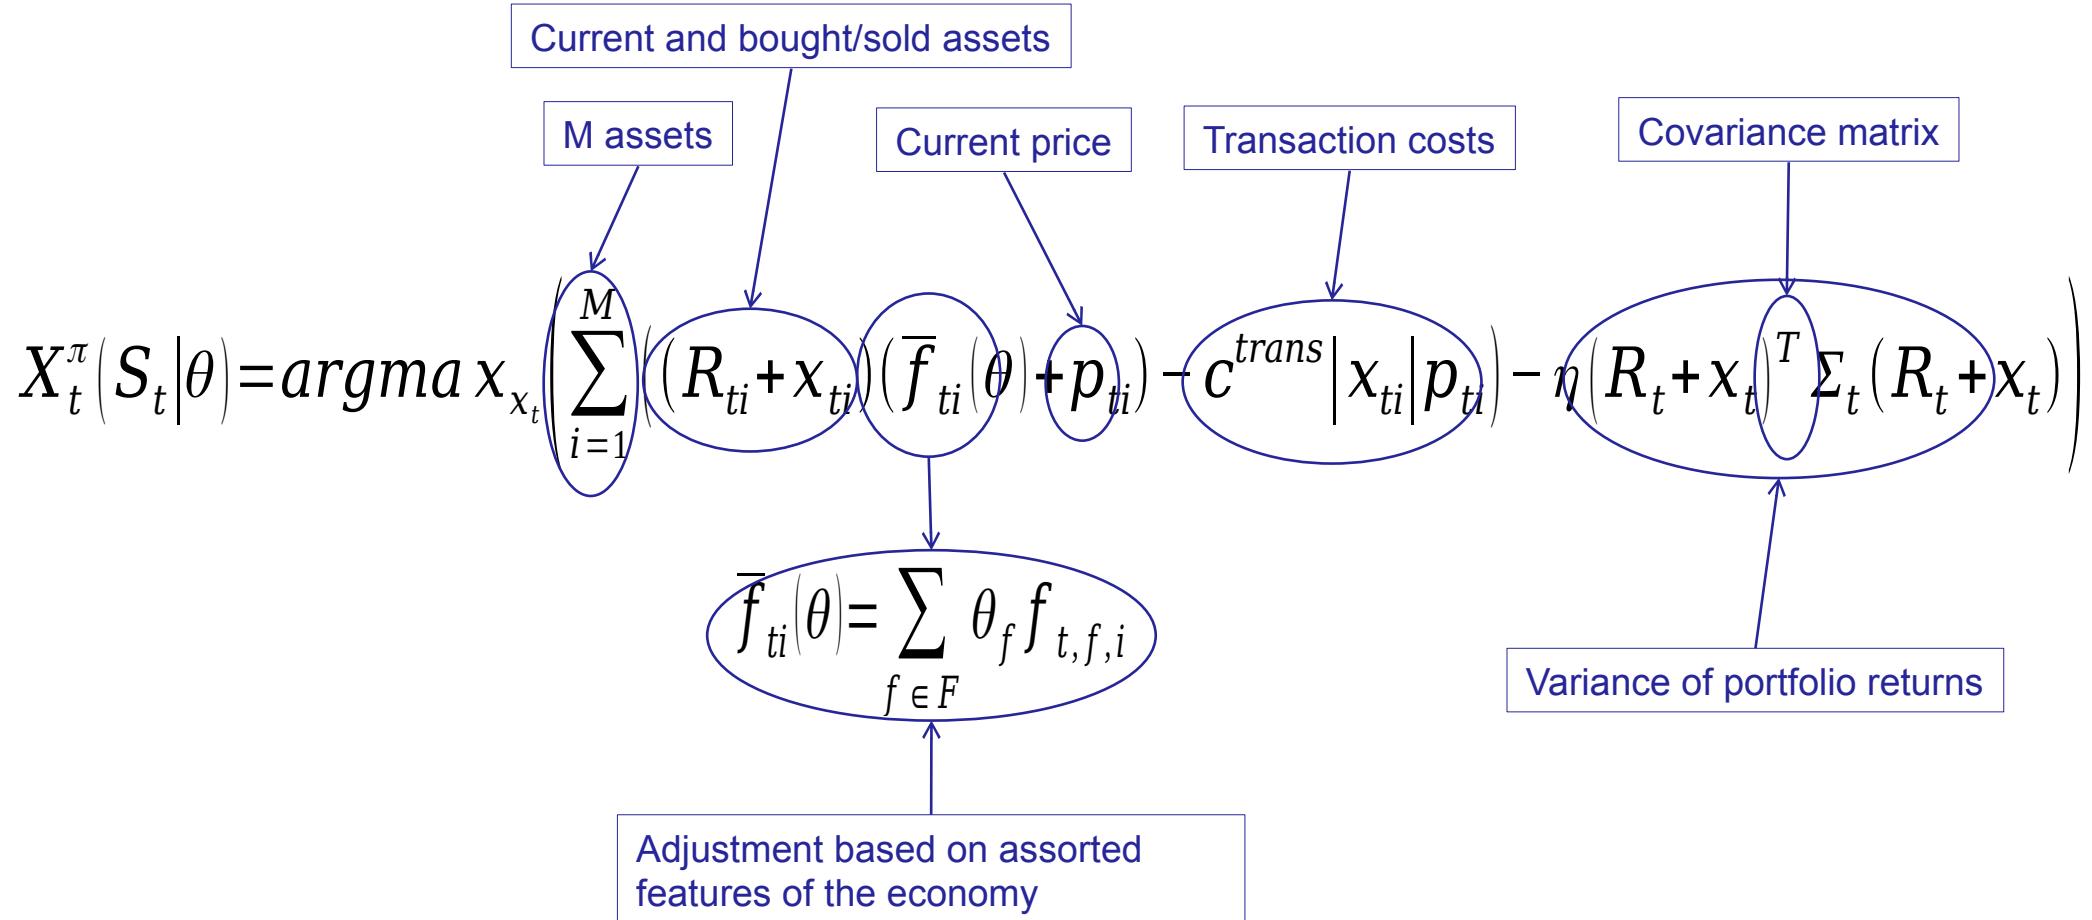
These are parameterized optimization problems:

- a) Find the shortest path to a destination, but add a buffer (e.g. 15 minutes) to make sure you arrive on time.
- b) Schedule nurses for hours per week, which allows for unforeseen emergencies.
- c) Advertise the product which solves:

Parametric CFAs are widely used in industry yet dismissed by the academic research community. This is actually quite a powerful strategy.

Portfolio optimization

Portfolio optimization:

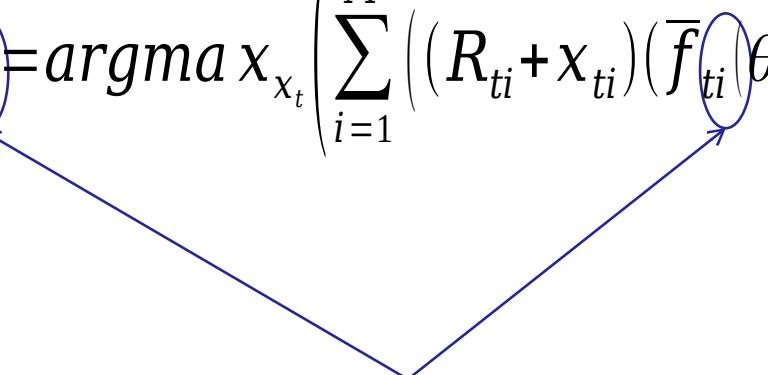


Portfolio optimization

Portfolio optimization:

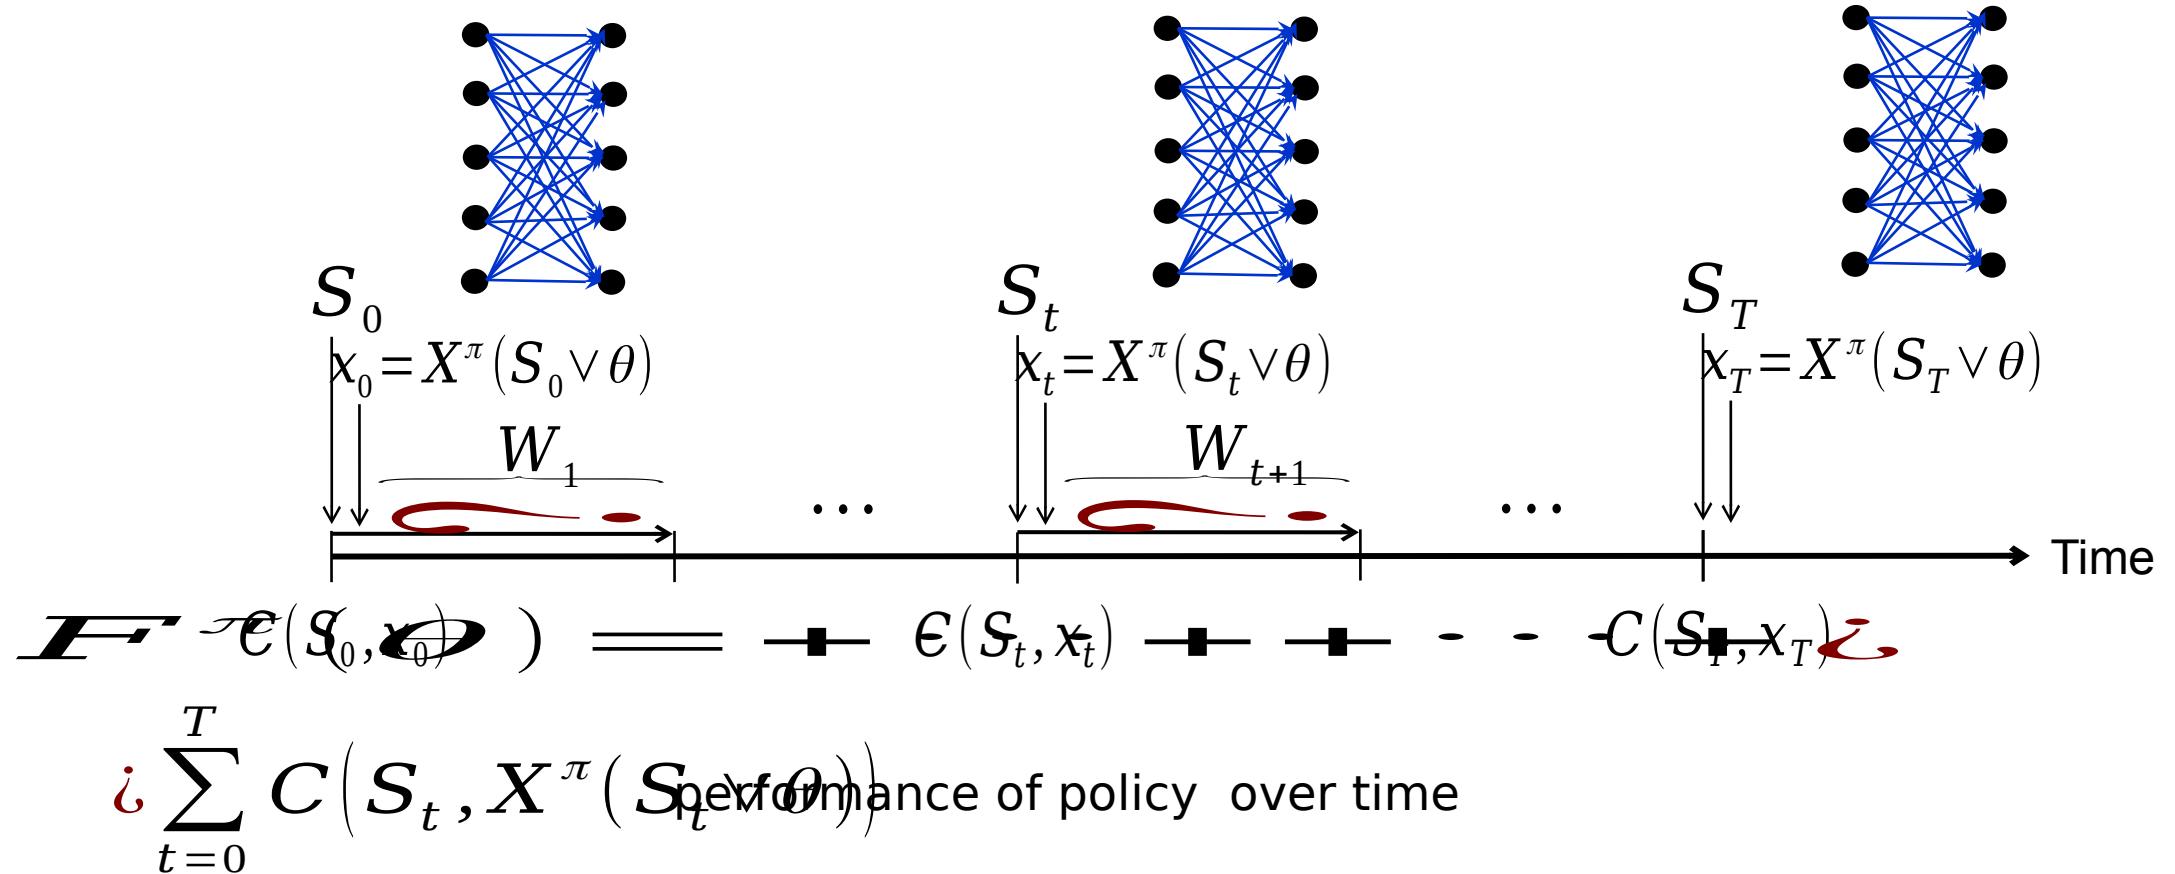
$$X_t^\pi(S_t|\theta) = \arg\max_{x_t} \left(\sum_{i=1}^M \left((R_{ti} + x_{ti})(\bar{f}_{ti}(\theta) + p_{ti}) - c^{trans}|x_{ti}|p_{ti} \right) - \eta(R_t + x_t)^T \Sigma_t (R_t + x_t) \right)$$

Tunable parameters



Portfolio optimization

Evaluating performance over time:



Derivative-based stochastic search

- The policy (deterministic nonlinear program)

$$X_t^\pi(S_t|\theta) = \arg\max_{x_t} \left(\sum_{i=1}^M \left((R_{ti} + x_{ti})(\bar{f}_{ti}(\theta) + p_{ti}) - c^{trans}|x_{ti}|p_{ti} \right) - \eta(R_t + x_t)^T \Sigma_t (R_t + x_t) \right)$$

- The objective function (evaluated using simulation)
where

Historical prices

- Stochastic optimization problem

$$\max_{\pi=(f,\theta)} \mathbb{E}_W F^\pi(\theta, W)$$

- Stochastic gradient algorithm

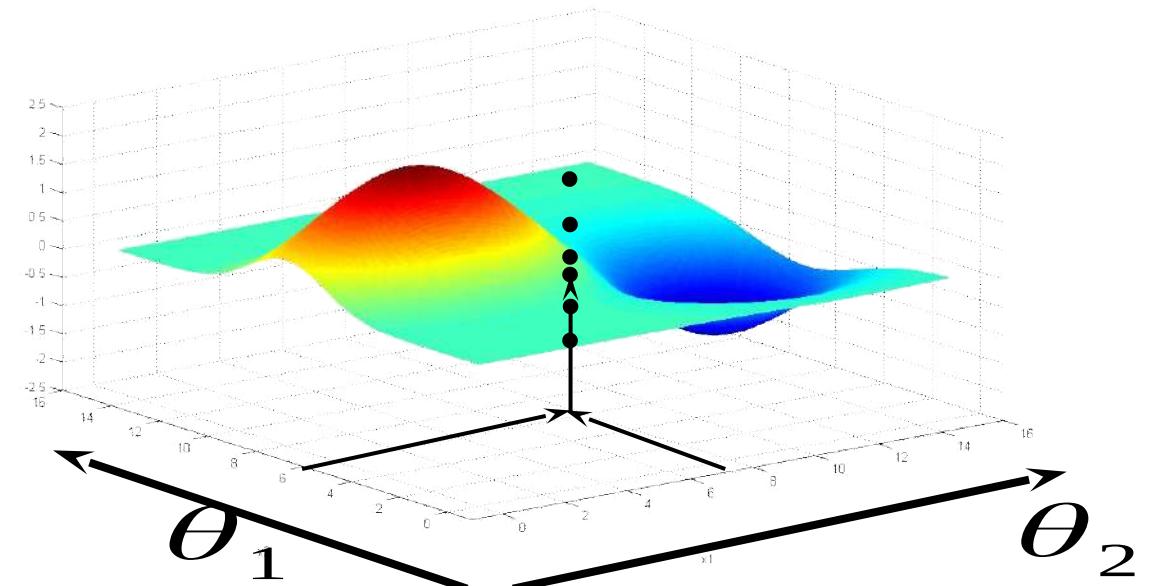
$$\theta^{n+1} = \theta^n + \alpha_n \nabla F^\pi(\theta^n, \omega^{n+1})$$

Cost function approximations

- CFAs are widely used in industry (in an ad-hoc manner)
 - » Airlines optimizing schedules with schedule slack to handle weather uncertainty.
 - » Manufacturers using buffer stocks to hedge against production delays and quality problems.
 - » Optimizing the location of distribution centers and warehouses, allowing for the benefits of inventory pooling.
 - » Grid operators scheduling extra generation capacity in case of outages.

Policy search

- Both PFAs and CFAs have tunable parameters which have to be tuned. We write this mathematically as
- There are two ways to evaluate a policy:
 - » In a simulator – This allows us to perform extensive testing in a controlled environment.
 - » In the field – This is “learning by doing”

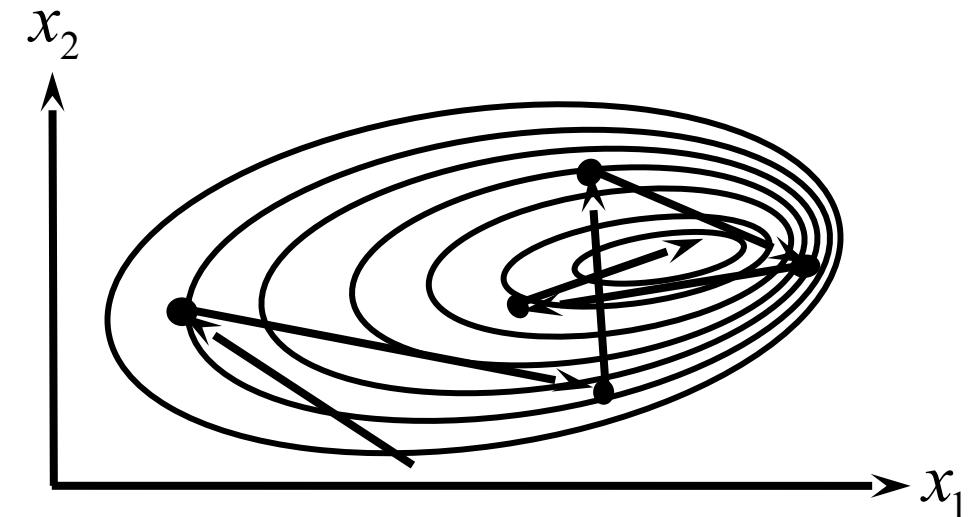


Policy search

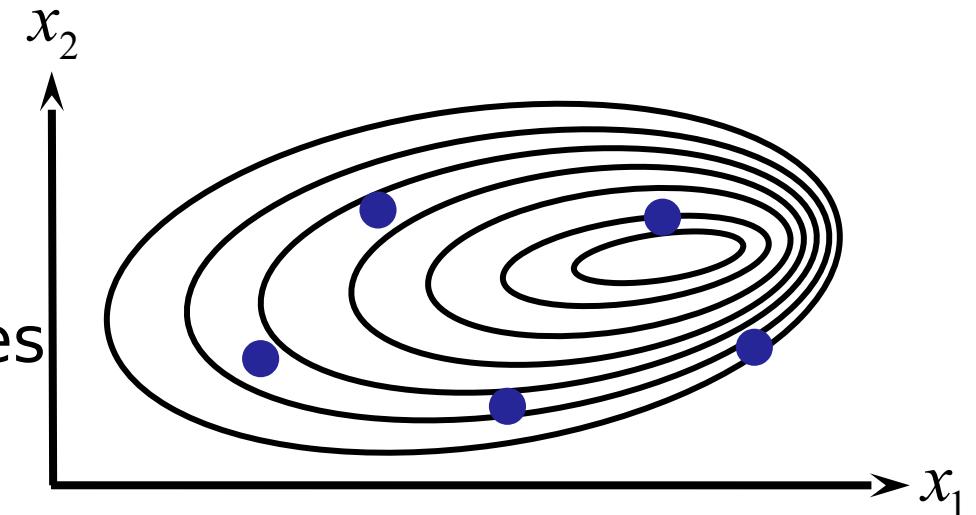
- How do we search for the best
 - » Derivative-based
 - Stochastic gradient methods:

$$\theta^{n+1} = \theta^n + \alpha_n \nabla_{\theta} F(\theta^n, W^{n+1})$$

Decision



- » Derivative-free
 - Build a belief model that approximates our function.



Designing policies

There are two fundamental strategies for designing policies

Policy search - Search over a class of methods for making decisions to optimize some metric over time.

- » Finding the best class of policy.
- » Finding the best policy within the class.

Lookahead approximations - Approximate the impact of a decision now on the future.

- » The contribution from the first period, plus
- » An approximation of the sum of contributions in future time periods resulting from the first decision.

Lookahead approximations

- Lookahead approximations combine:
 - » The immediate contribution (or cost) of a decision made now...
 - » ... and an approximation of future contributions (or costs)



Lookahead approximations

Lookahead policies are based on solving

$$X_t^*(S_t) = \arg\max_x \left(C(S_t, x_t) + \mathbb{E} \left[\max_{\pi} \left(\mathbb{E} \sum_{t'=t+1}^T C(S_t, X_t^\pi(S_t)) \mid S_{t+1} \right) \mid S_t, x_t \right] \right)$$

Contribution we receive now

Future contributions

- » This looks like scary mathematics, but it is what all of us are doing when we make decisions now that consider what might happen in the future.
- » The challenge is ... *how to compute it!!!*

Lookahead approximations

Lookahead approximations

Approximate the impact of a decision now on the future

$$X_t^*(S_t) = \operatorname{argmax}_{x_t} \left(C(S_t, x_t) + \mathbb{E} \left[\max_{\pi} \left\{ \mathbb{E} \sum_{t'=t+1}^T C(S_{t'}, X_t^\pi(S_{t'})) \mid S_{t+1} \right\} \mid S_t, x_t \right] \right)$$

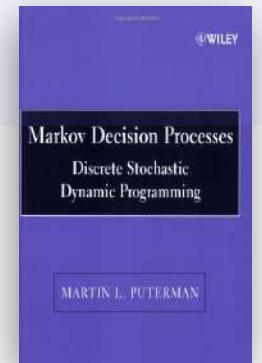
3) Value function approximations (VFAs)

$$X_t^*(S_t) = \operatorname{argmax}_{x_t} \left(C(S_t, x_t) + \mathbb{E} \left\{ V_{t+1}(S_{t+1}) \mid S_t, x_t \right\} \right)$$

$$X_t^{VFA}(S_t) = \operatorname{argmax}_{x_t} \left(C(S_t, x_t) + \mathbb{E} \left\{ \bar{V}_{t+1}(S_{t+1}) \mid S_t, x_t \right\} \right)$$

$$= \operatorname{argmax}_{x_t} \left(C(S_t, x_t) + \bar{V}_t^x(S_t^x) \right)$$

$$= \operatorname{argmax}_{x_t} \bar{Q}_t(S_t, x_t) \quad ("Q\text{-learning}")$$



Lookahead approximations

Lookahead approximations

Approximate the impact of a decision now on the future

$$X_t^*(S_t) = \arg\max_{x_t} \left(C(S_t, x_t) + \mathbb{E} \left\{ \max_{\pi} \left\{ \mathbb{E} \sum_{t'=t+1}^T C(S_{t'}, X_t^\pi(S_{t'})) \mid S_{t'+1} \right\} \mid S_t, x_t \right\} \right)$$

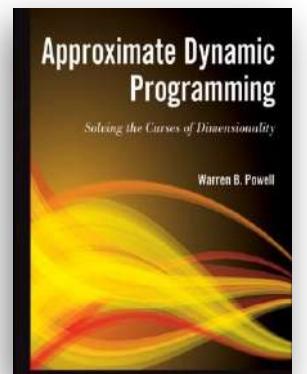
3) Value function approximations (VFAs)

$$X_t^*(S_t) = \arg\max_{x_t} \left(C(S_t, x_t) + \mathbb{E} \left\{ V_{t+1}(S_{t+1}) \mid S_t, x_t \right\} \right)$$

$$X_t^{VFA}(S_t) = \arg\max_{x_t} \left(C(S_t, x_t) + \mathbb{E} \left\{ \bar{V}_{t+1}(S_{t+1}) \mid S_t, x_t \right\} \right)$$

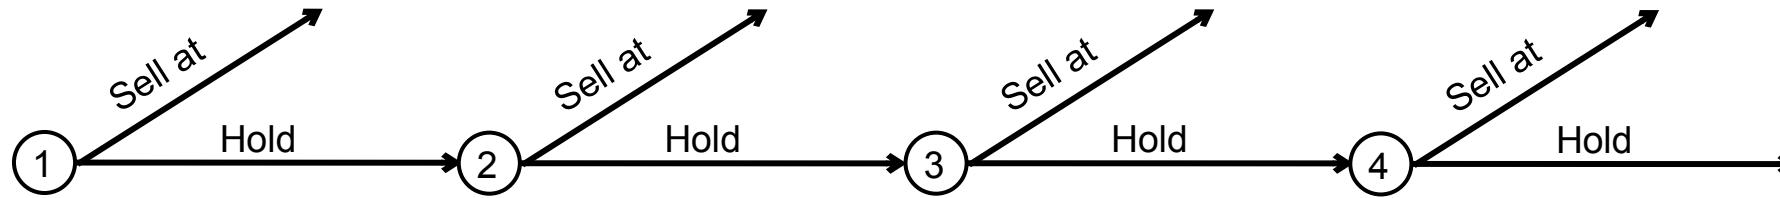
$$= \arg\max_{x_t} \left(C(S_t, x_t) + \bar{V}_t^x(S_t^x) \right)$$

$$= \arg\max_{x_t} \bar{Q}_t(S_t, x_t) \quad ("Q\text{-learning}")$$



Value function approximations

Longstaff and Schwartz - Pricing an American option

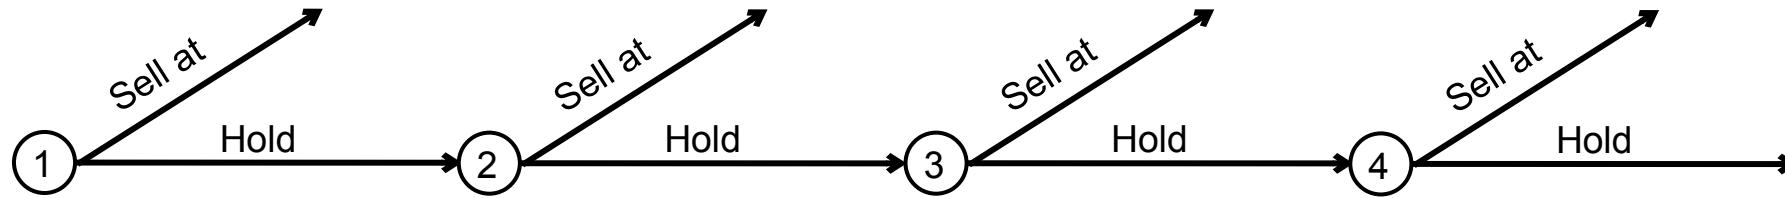


Outcome	Stock prices			
	1	2	3	4
1	1.21	1.08	1.17	1.15
2	1.09	1.12	1.17	1.13
3	1.15	1.08	1.22	1.35
4	1.17	1.12	1.18	1.15
5	1.08	1.15	1.10	1.27
6	1.12	1.22	1.23	1.17
7	1.16	1.14	1.13	1.19
8	1.22	1.18	1.21	1.28
9	1.08	1.11	1.09	1.10
10	1.15	1.14	1.18	1.22



Value function approximations

Longstaff and Schwartz - Pricing an American option

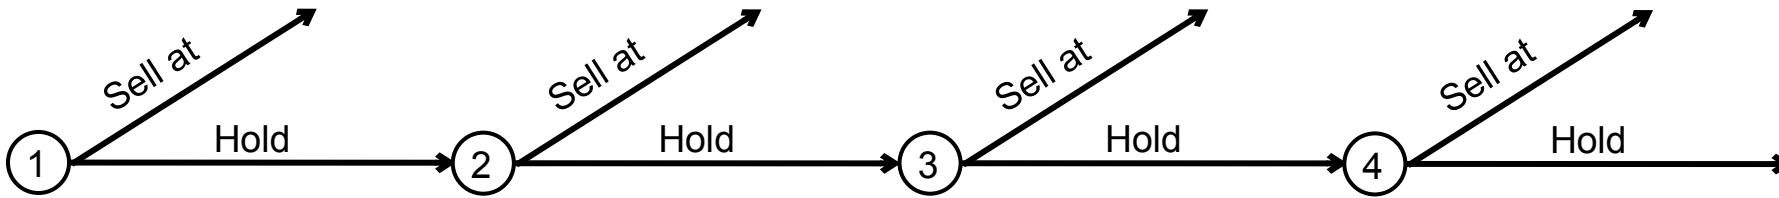


Outcome	Decision	
	Sell	Hold
1	0.03	$0.04155 \times .95 = \mathbf{0.03947}$
2	0.03	$0.03662 \times .95 = \mathbf{0.03479}$
3	0.00	$0.02397 \times .95 = \mathbf{0.02372}$
4	0.02	$0.03346 \times .95 = \mathbf{0.03178}$
5	0.10	$0.05285 \times .95 = \mathbf{0.05021}$
6	0.00	$0.00414 \times .95 = \mathbf{0.00394}$
7	0.07	$0.00899 \times .95 = 0.00854$
8	0.00	$0.01610 \times .95 = \mathbf{0.01530}$
9	0.11	$0.06032 \times .95 = 0.05731$
10	0.02	$0.03099 \times .95 = \mathbf{0.02944}$

Prices at time 2

Value function approximations

Longstaff and Schwartz - Pricing an American option



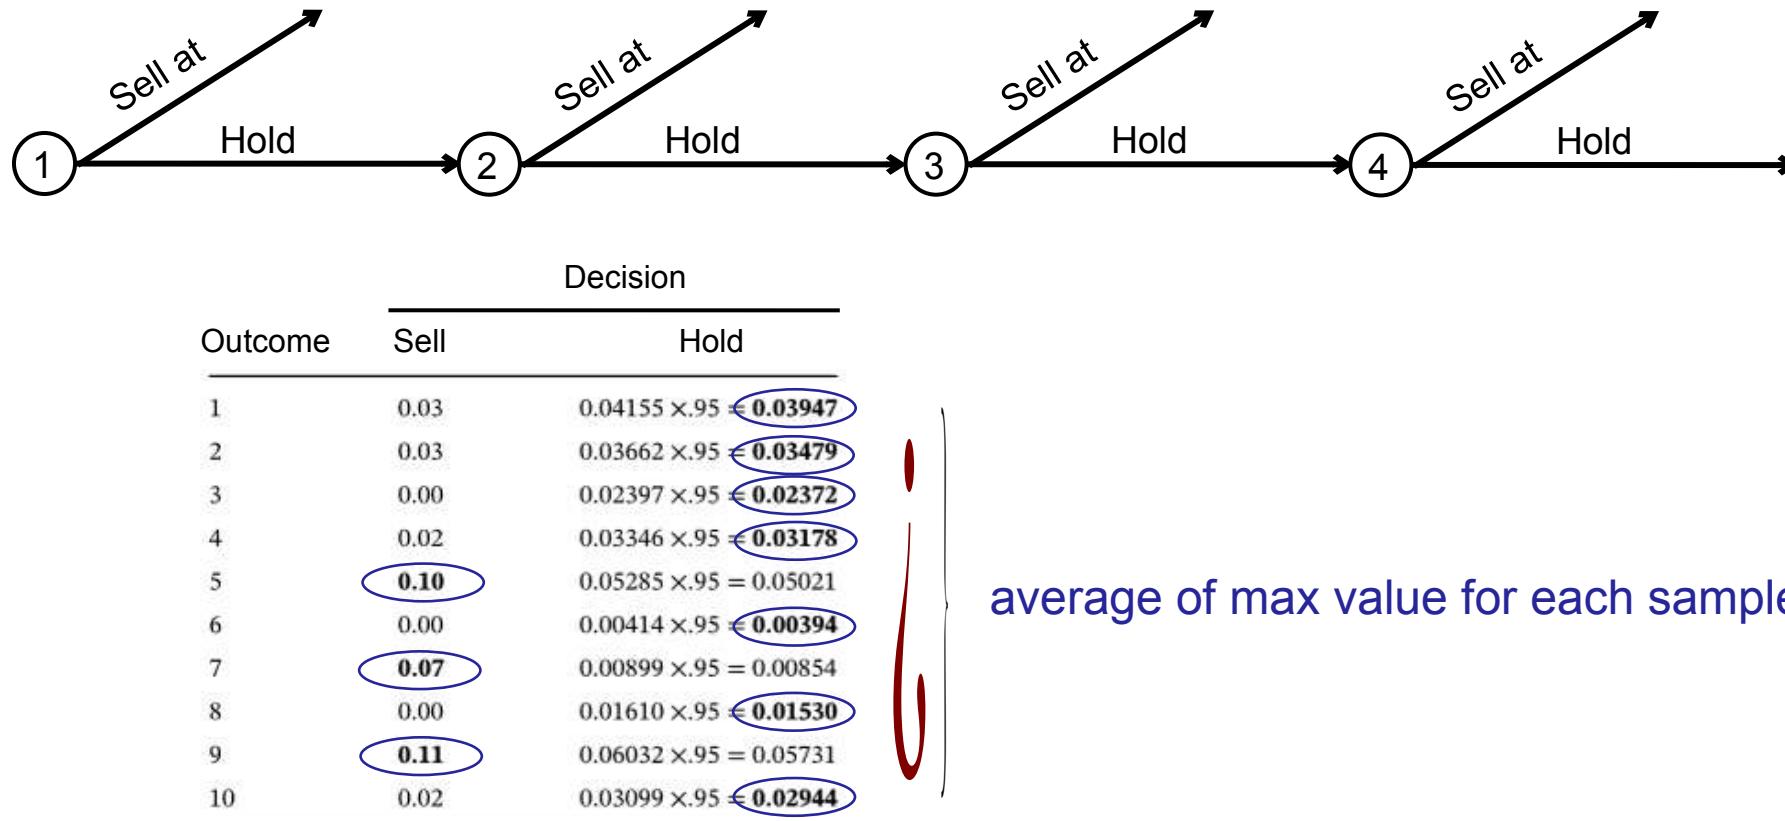
Outcome	Decision	
	Sell	Hold
1	0.03	$0.04155 \times .95 = \mathbf{0.03947}$
2	0.03	$0.03662 \times .95 = \mathbf{0.03479}$
3	0.00	$0.02397 \times .95 = \mathbf{0.02372}$
4	0.02	$0.03346 \times .95 = \mathbf{0.03178}$
5	0.10	$0.05285 \times .95 = \mathbf{0.05031}$
6	0.00	$0.00414 \times .95 = \mathbf{0.00394}$
7	0.07	$0.00899 \times .95 = \mathbf{0.00854}$
8	0.00	$0.01610 \times .95 = \mathbf{0.01530}$
9	0.11	$0.06032 \times .95 = \mathbf{0.05731}$
10	0.02	$0.03099 \times .95 = \mathbf{0.02944}$

Expected value of holding the asset:

$$\mathcal{V}_t = \theta_{t0} p_{t-2} + \theta_{t1} p_{t-1} + \theta_{t2} p_t + \theta_{t3} (p_t)^2$$
$$\hat{\mathcal{V}}_t = 0.0056 - 0.1234 p_1 + 0.6011 p_2 - 0.3903 (p_2)^2$$

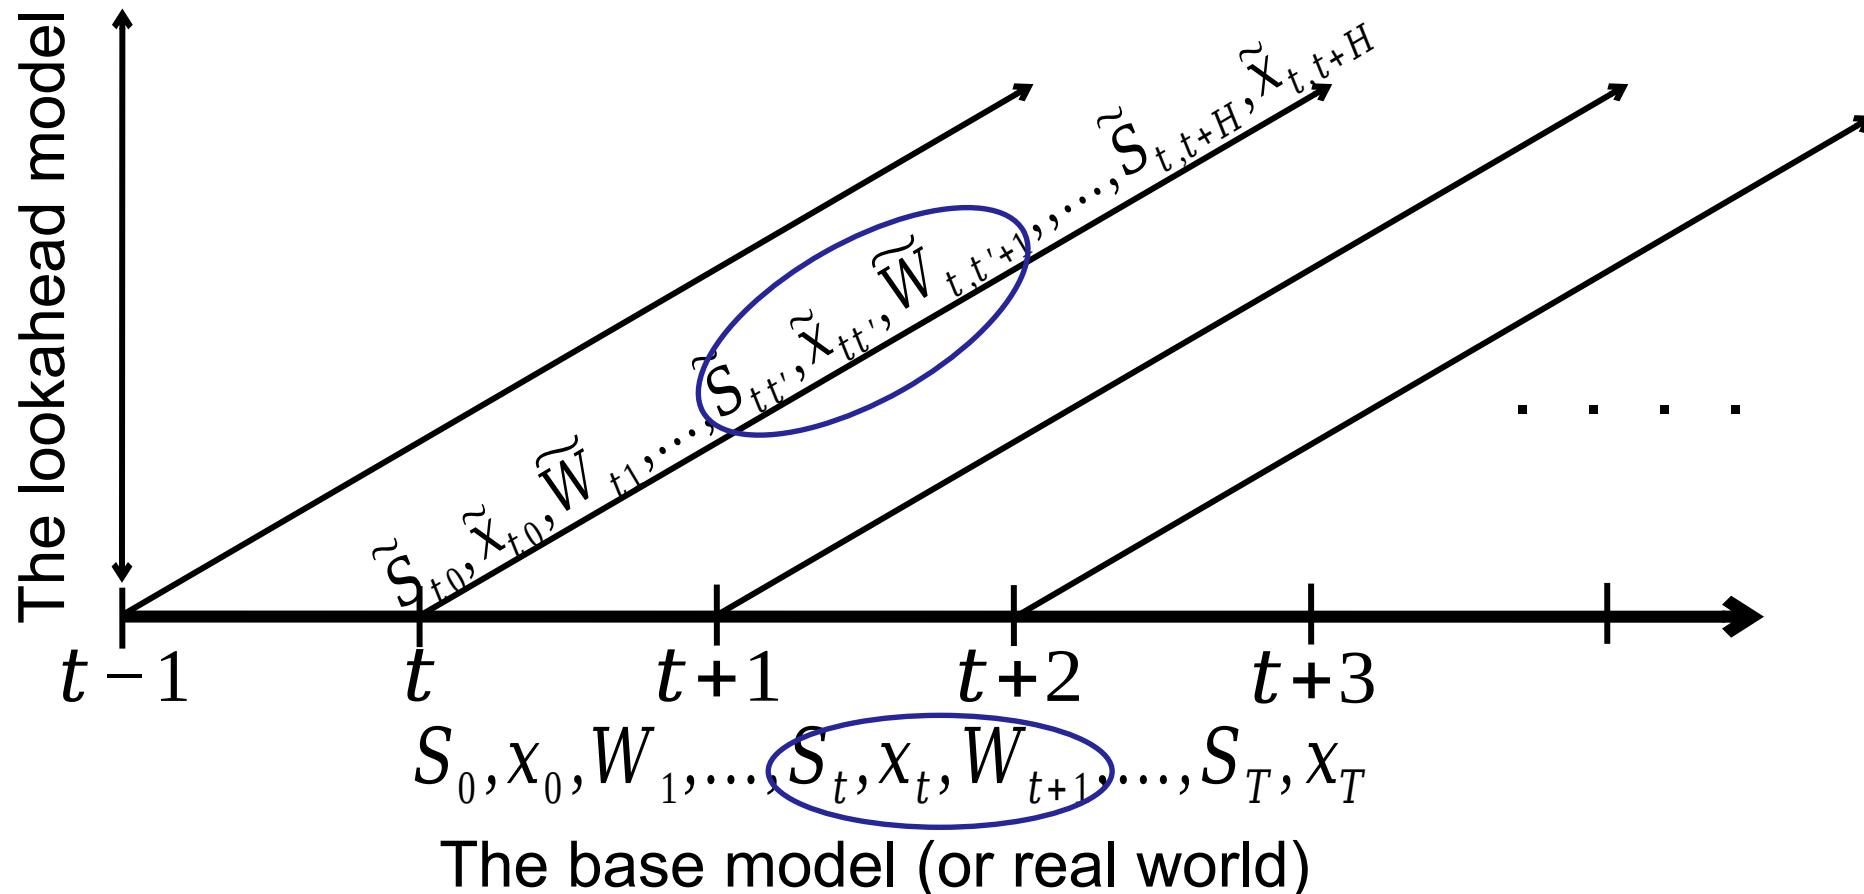
Value function approximations

Longstaff and Schwartz - Pricing an American option



Direct lookahead policies

4) **Direct lookahead policies (DLAs)** – Here we create an approximation called the *approximate lookahead model*:



Direct lookahead policies

4) Direct lookahead policies (DLAs) – Here we create an approximation called the *approximate lookahead model*:

$$(\tilde{S}_{tt}, \tilde{x}_{tt}, \tilde{W}_{t,t+1}, \tilde{S}_{t,t+1}, \tilde{x}_{t,t+1}, \tilde{W}_{t,t+2}, \dots, \tilde{S}_{tt'}, \tilde{x}_{tt'}, \tilde{W}_{t,t'+1}, \dots)$$

There are six classes of approximations we can introduce.
Our direct lookahead policy now requires solving:

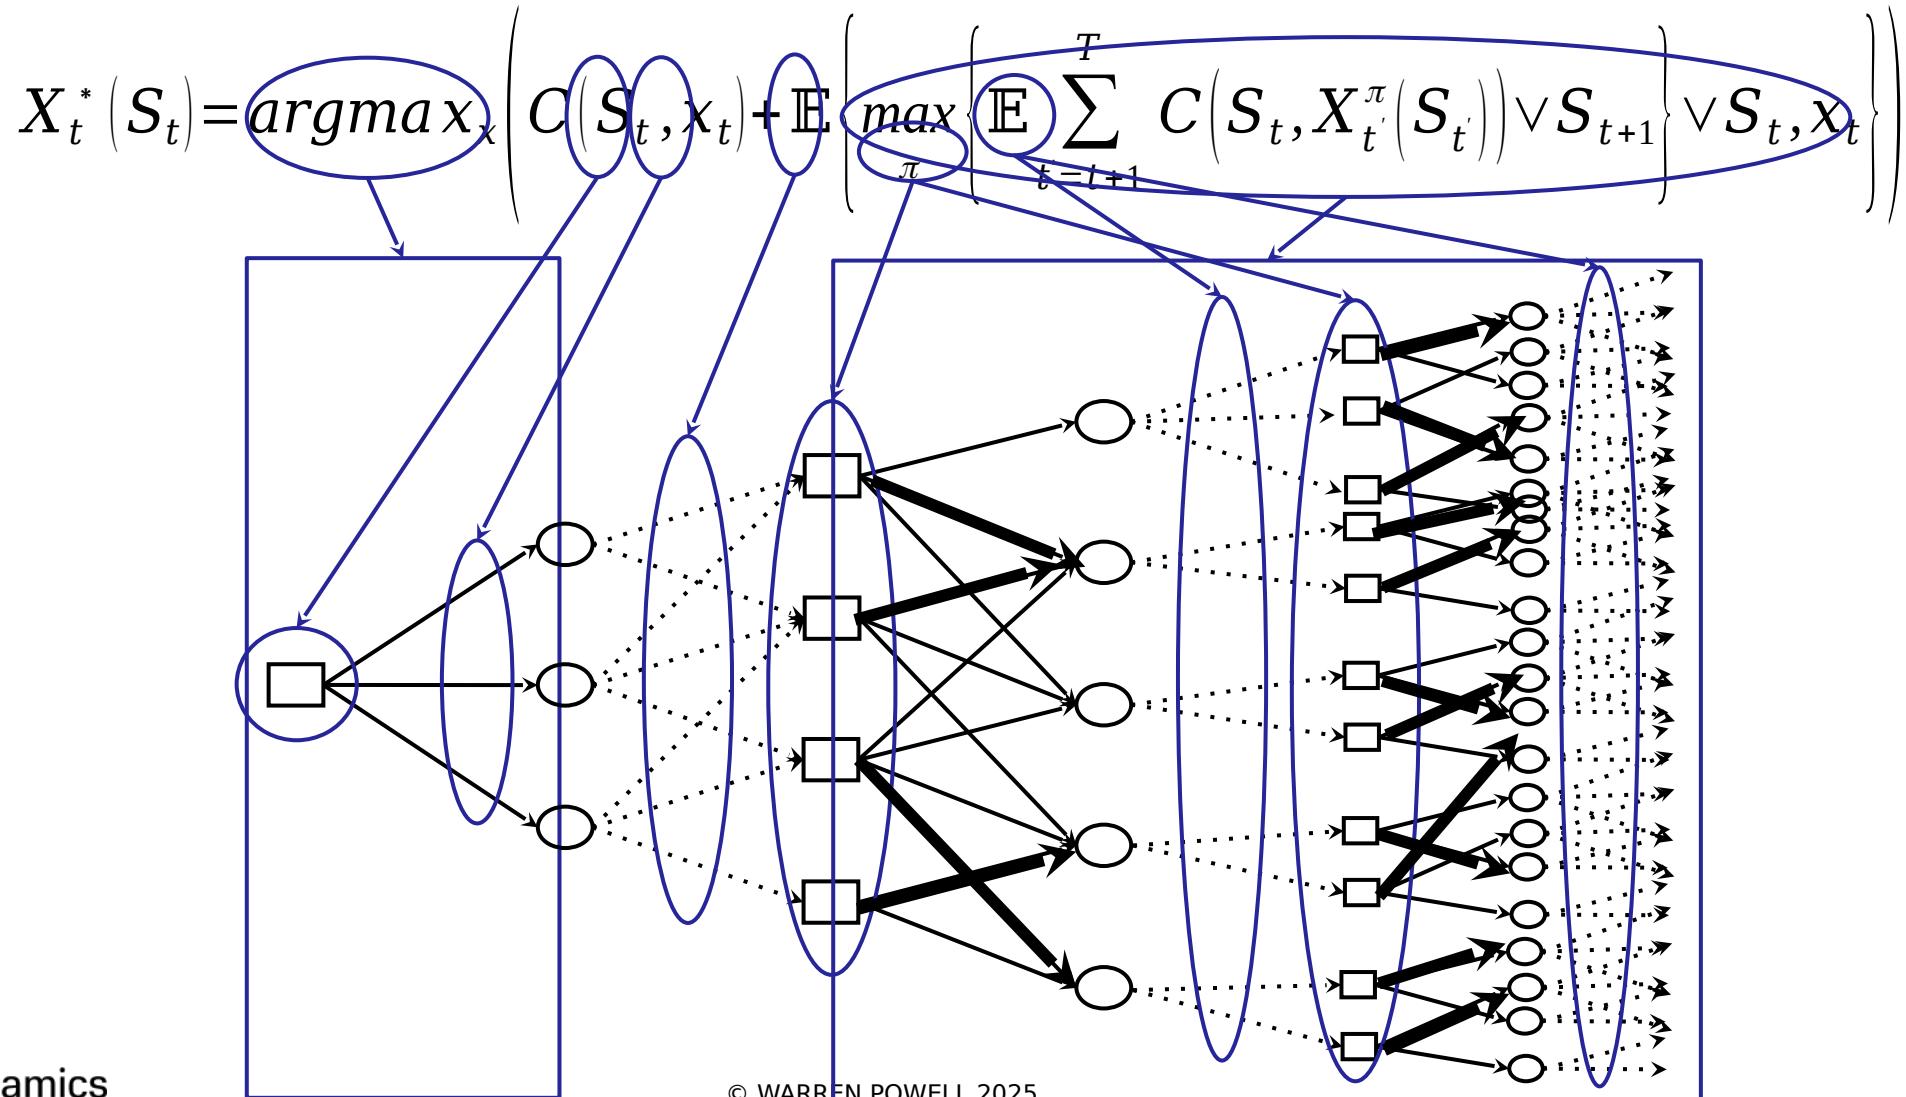
$$X_t^{DLA}(S_t \vee \theta) = \operatorname{argmax}_x \left(C(S_t, x_t) + \tilde{E} \left\{ \max_{\tilde{\pi}} \tilde{E} \left\{ \sum_{t'=t+1}^{t+H} C(\tilde{S}_{t'}, \tilde{X}_{t'}^{\tilde{\pi}}(\tilde{S}_{t'})) \vee \tilde{S}_{t+1} \right\} | S_t, x_t \right\} \right)$$

The diagram illustrates the components of the lookahead policy equation. It shows arrows pointing from the equation to six labeled boxes:

- Sampled information process (points to the sampled information term \tilde{E})
- Restricted horizon (points to the restricted horizon term $\sum_{t'=t+1}^{t+H}$)
- Limited decisions (points to the limited decisions term $C(\tilde{S}_{t'}, \tilde{X}_{t'}^{\tilde{\pi}}(\tilde{S}_{t'})) \vee \tilde{S}_{t+1}$)
- Simplified policies (points to the simplified policies term $\max_{\tilde{\pi}}$)
- Reduced state variable (points to the reduced state variable term $C(S_t, x_t)$)

Lookahead approximations

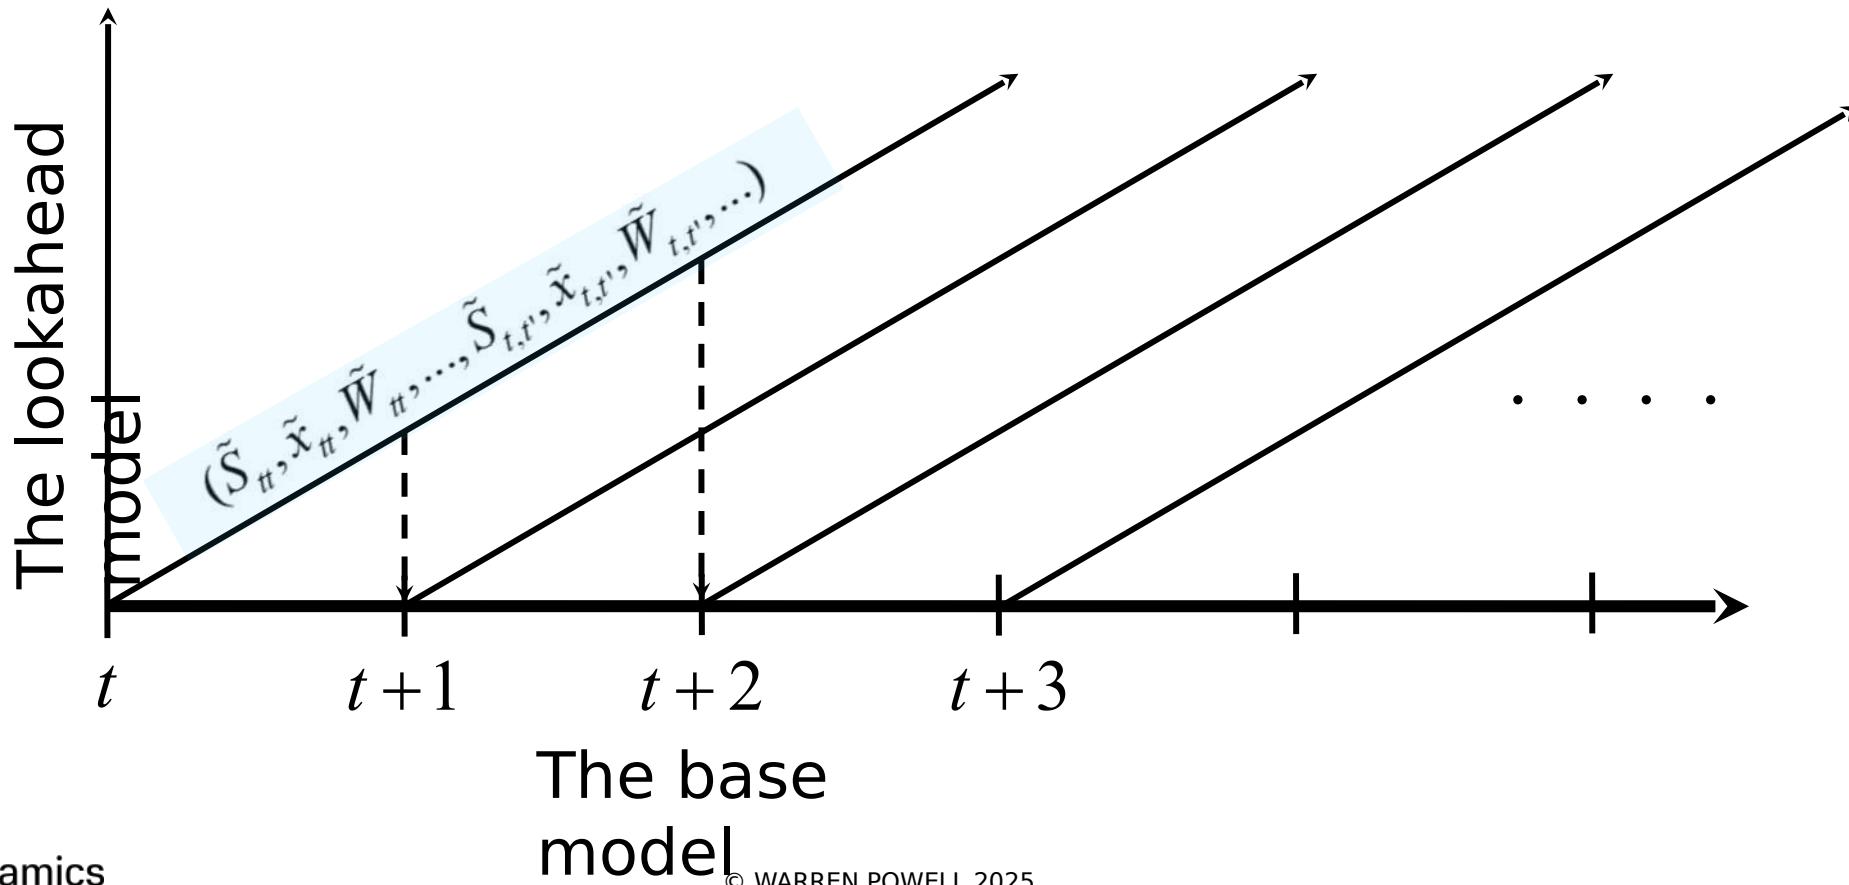
Lookahead policies are based on solving



Direct lookahead policies

Direct Lookahead Policies (DLAs)

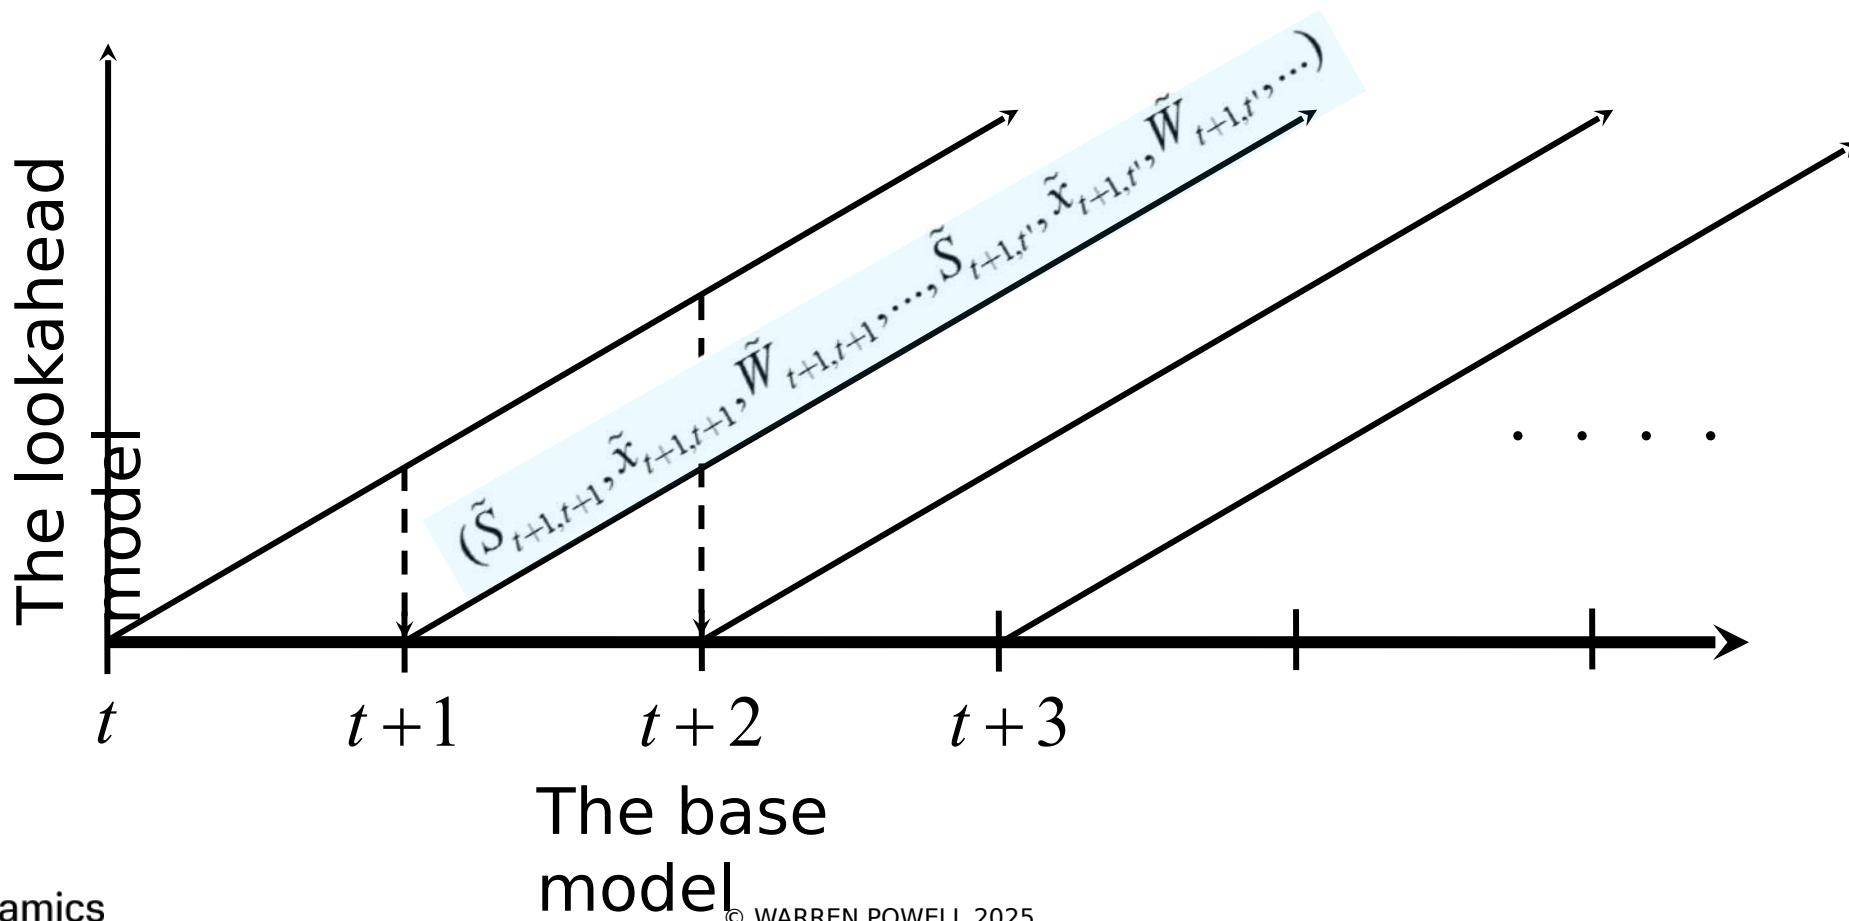
- » Tilde variables are used to model approximate lookahead



Direct lookahead policies

Direct Lookahead Policies (DLAs)

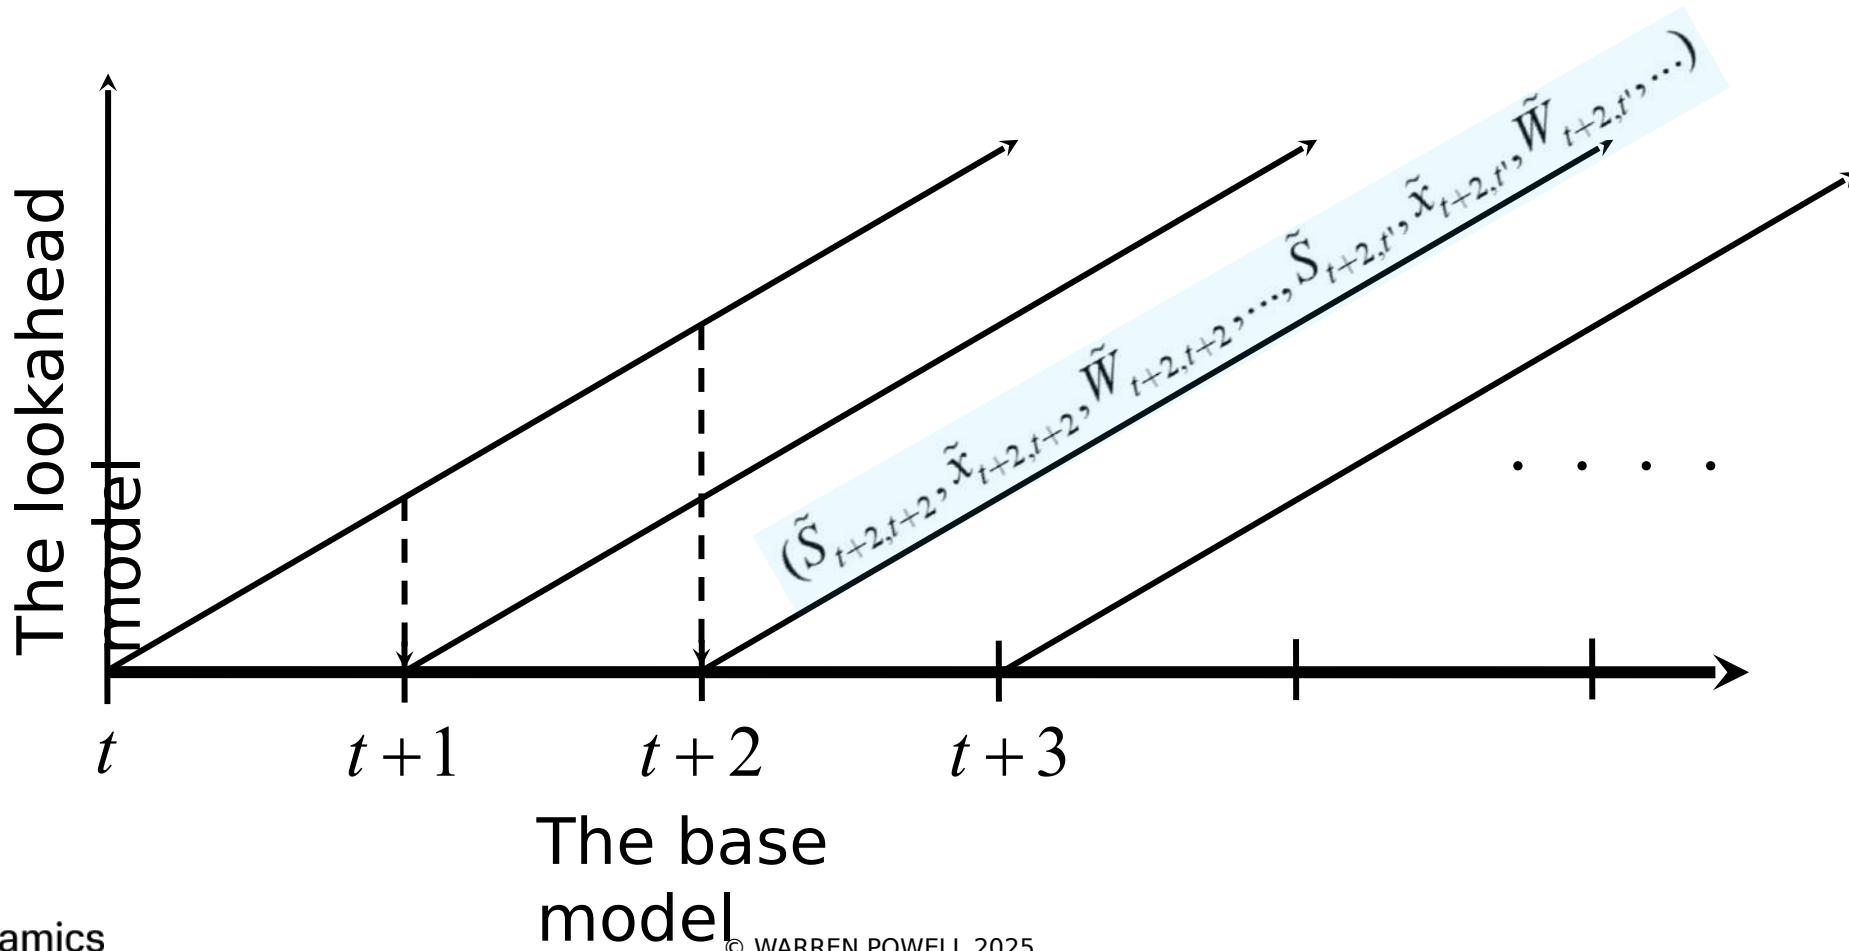
- » Tilde variables are used to model approximate lookahead



Direct lookahead policies

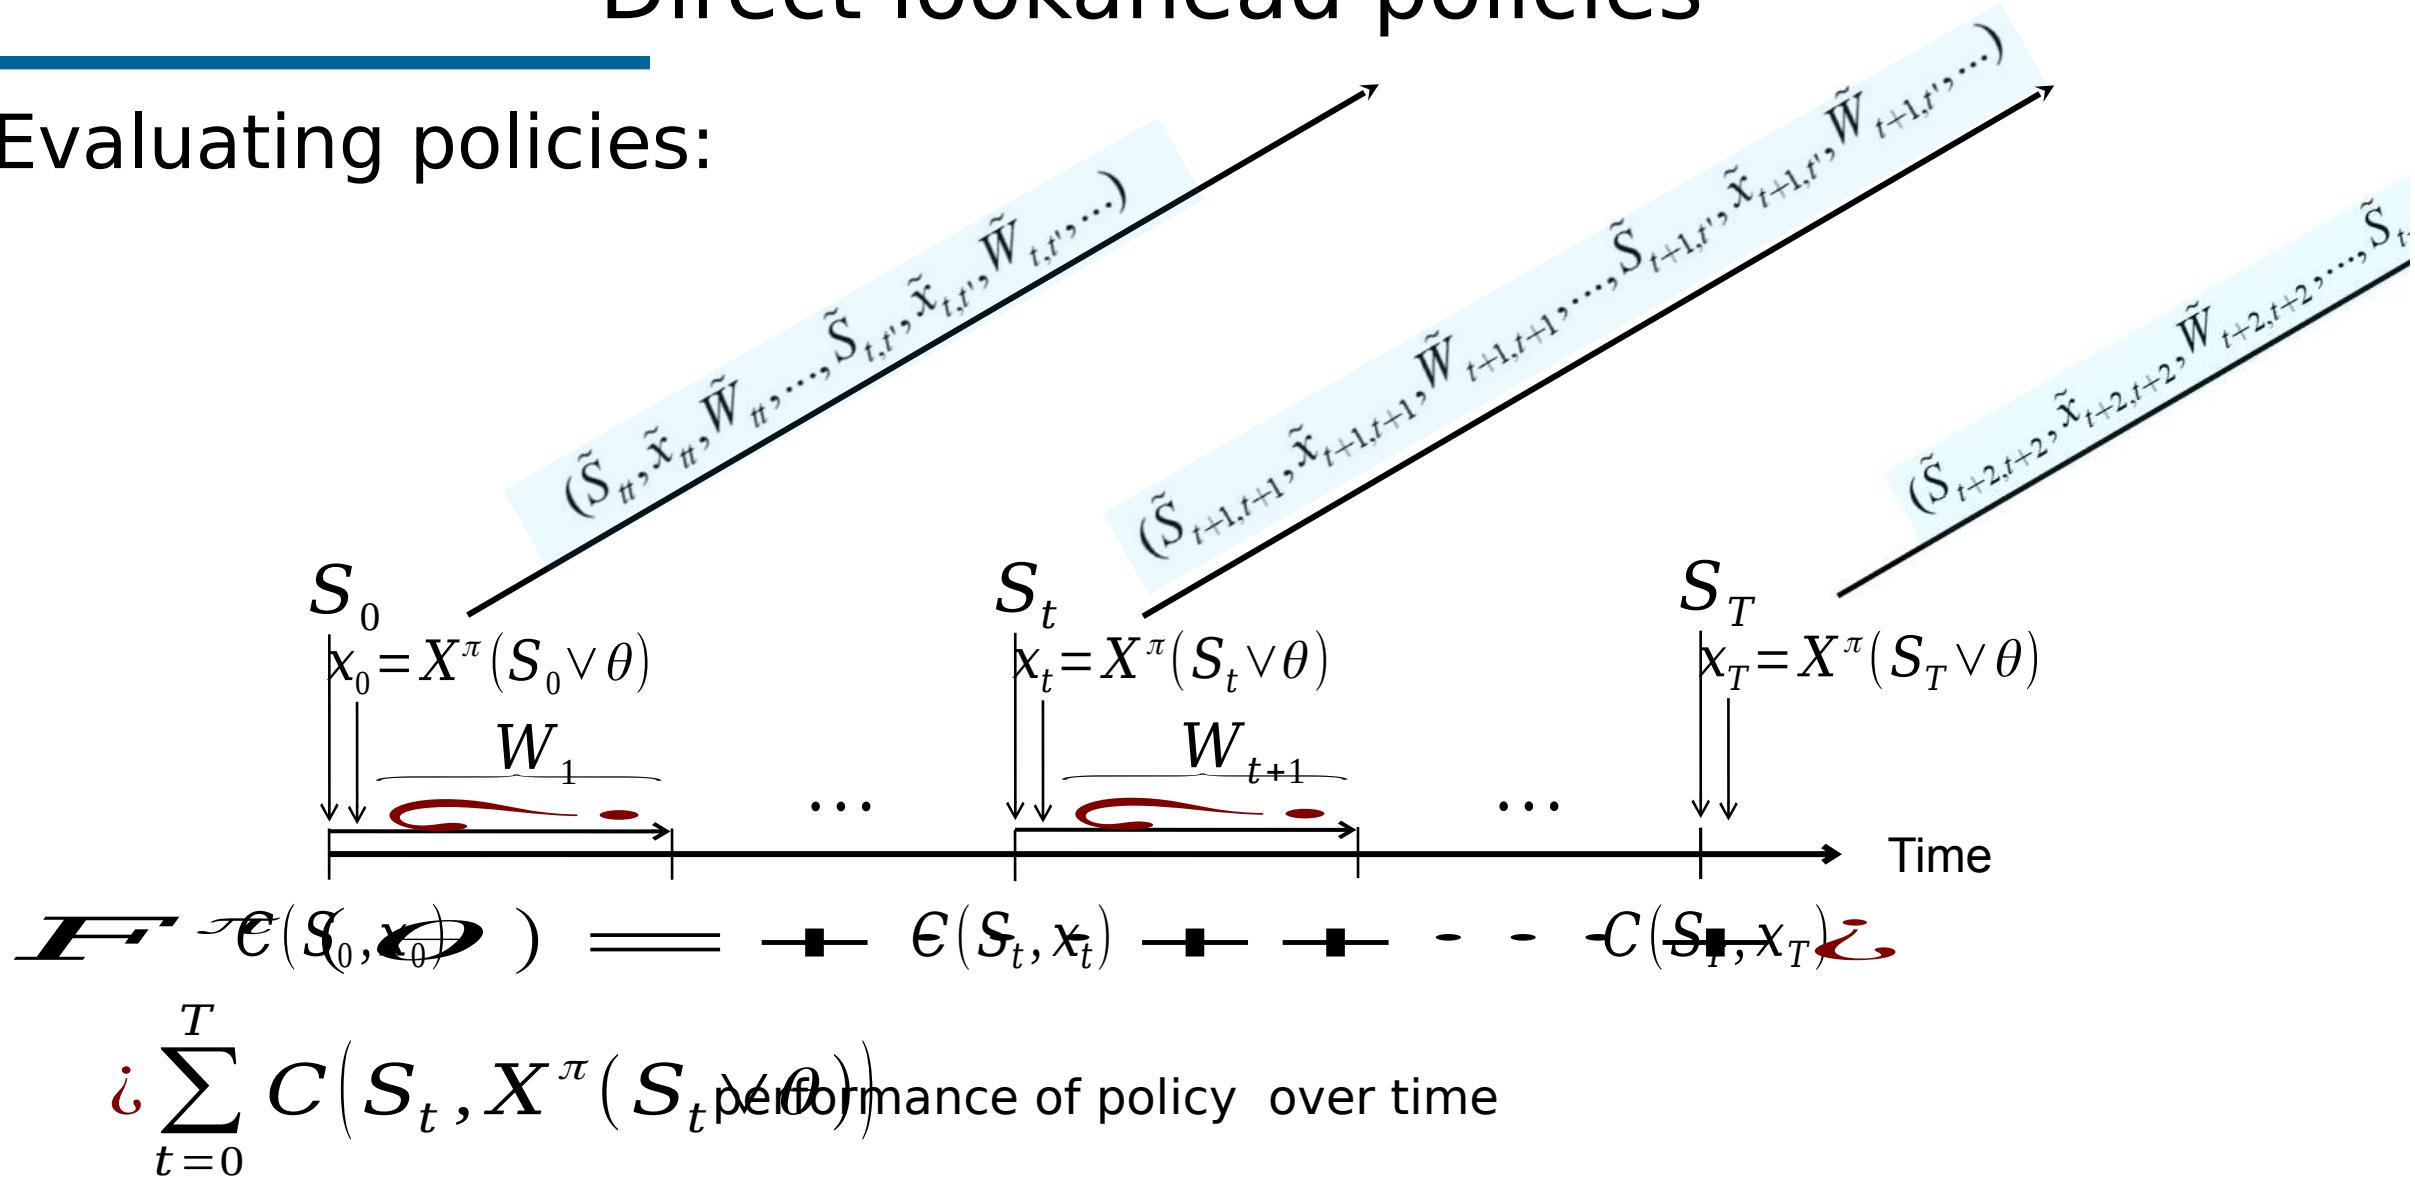
Direct Lookahead Policies (DLAs)

- » Tilde variables are used to model approximate lookahead



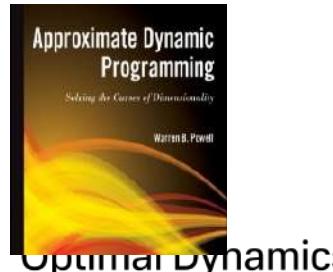
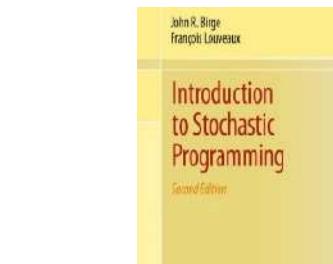
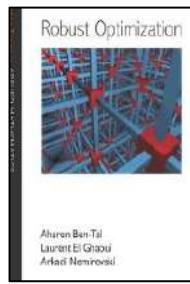
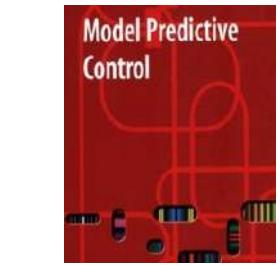
Direct lookahead policies

Evaluating policies:

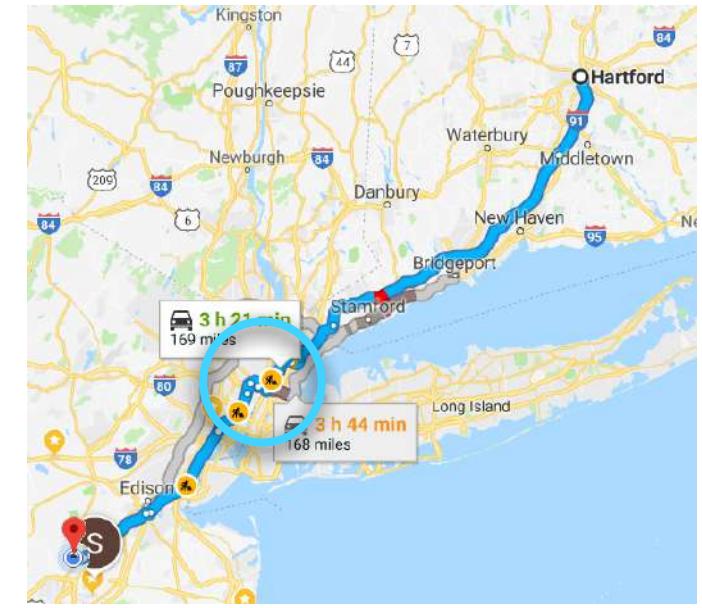


Direct lookahead policies

Examples of Lookahead Models



- » **The deterministic lookahead model**
 - This is what is most widely used in practice.
 - Standard approach is to use a “best estimate” (which means deterministic) of travel times in the future.
 - This is often referred to as “model predictive control”
- » **Robust optimization** - We could use the 90th percentile of travel times.
- » **Stochastic programming** – We represent the future using, say, 20 samples.
- » **Approximate dynamic programming applied to approximate lookahead model**
- » **Chance constrained programming** – Impose constraint on the probability of



Designing policies

Policy search policies

Policy function approximations (PFAs)

- » Simple rules, functions
- » Examples:
 - Order up to inventory policies
 - Buy low, sell high
 - Linear/nonlinear functions
 - Any function that might be used in machine learning

Cost function approximations (CFAs)

- » Parameterized optimization models
- » Examples
 - Schedule slack for scheduling models
 - Buffer stocks in supply chain models
 - Power reserves for energy planning

Lookahead policies

Value function approximations (VFAs)

- » Making a decision now using the value of being in a future state
- » Examples:
 - The value of a truck driver
 - The value of holding an asset

Direct lookaheads (DLAs)

- » Models that optimize over a planning horizon (deterministically/stochastically)
- » Examples:
 - Deterministic lookahead: Google maps
 - Stochastic lookahead: Decision trees

Designing policies

Policy search policies

Policy function approximations (PFAs)

- » Simple rules, functions
- » Examples:
 - Order up to inventory policies
 - Buy low, sell high
 - Linear/nonlinear decision rules

Cost function approximations (CFAs)

- » Parameterized optimization models
- » Examples
 - Schedule slack for scheduling models
 - Buffer stocks in supply chain models
 - Power reserves for energy planning

Lookahead policies

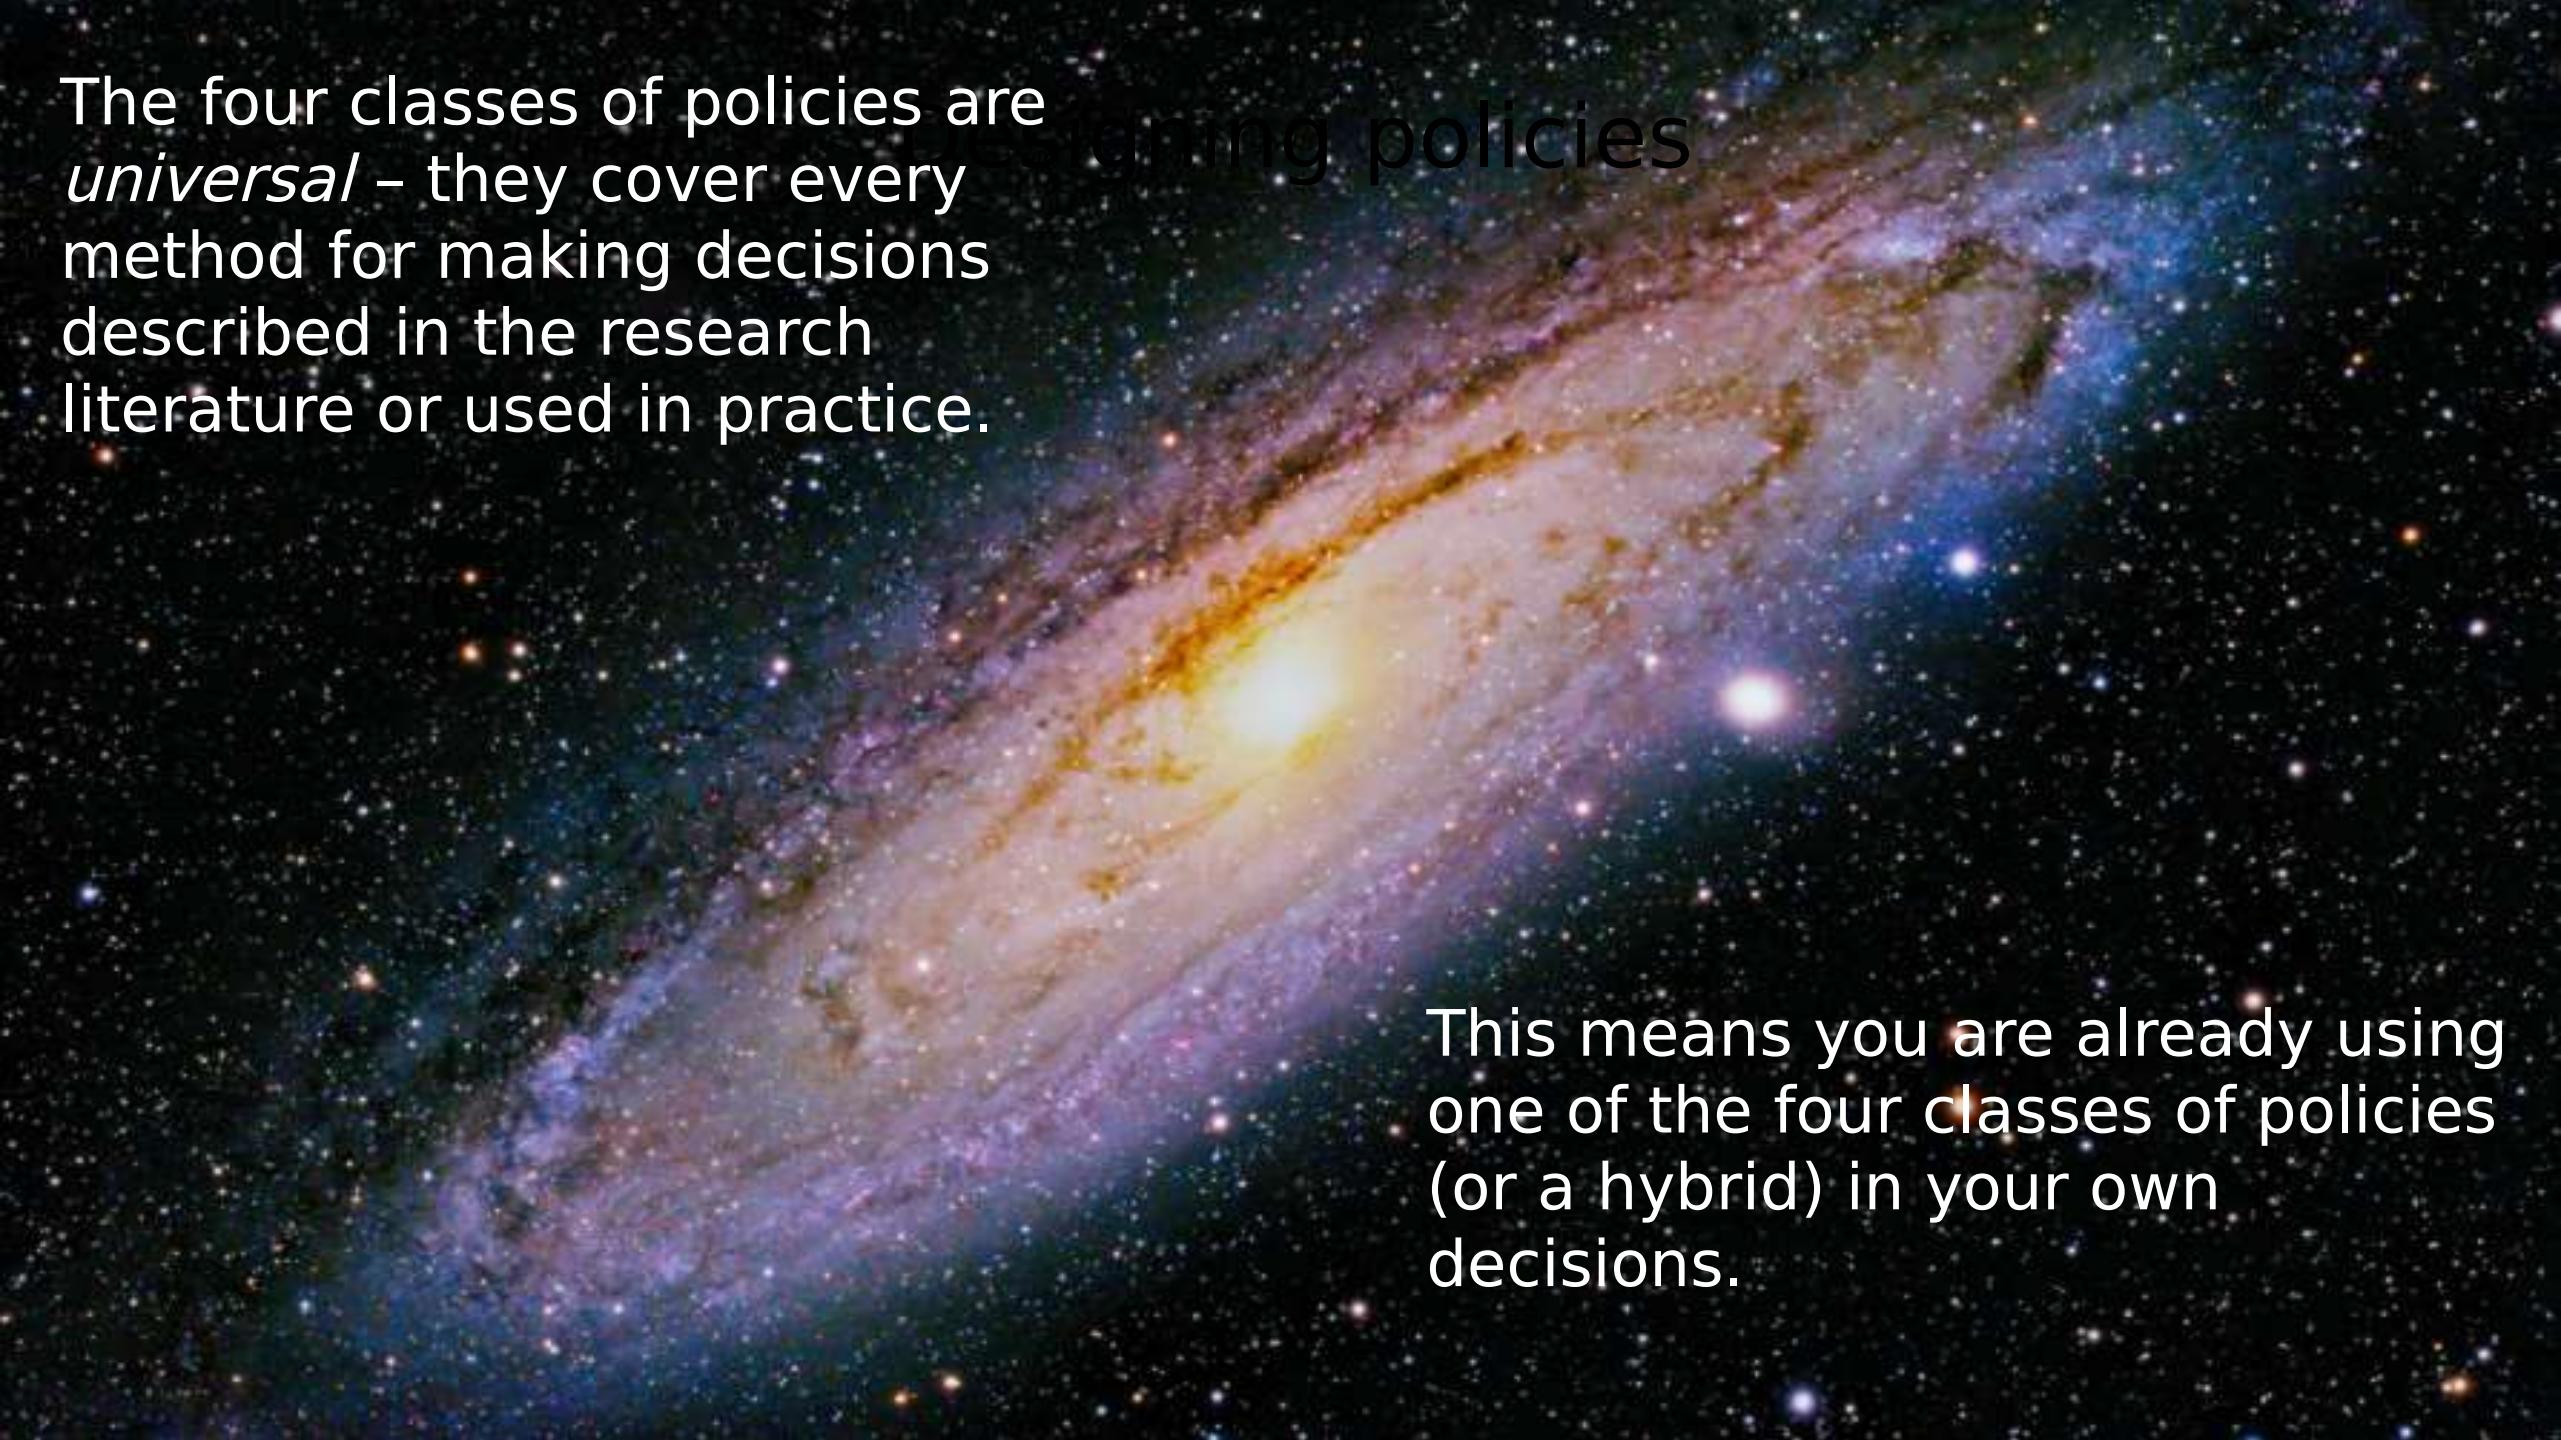
Value function approximations (VFAs)

- » Making a decision now using the value of being in a future state
- » Examples:
 - The value of a truck driver
 - The value of holding an asset

Direct lookaheads (DLAs)

- » Models that optimize over a planning horizon (deterministically/stochastically)
- » Examples:
 - Deterministic lookahead: Google maps
 - Stochastic lookahead: Decision trees

Three of the four classes of policies involve solving optimization problems.

The background of the slide features a detailed image of a spiral galaxy, likely the Milky Way, showing its characteristic spiral arms and central bright nucleus. The colors range from deep blues and purples to bright yellows and reds, representing different stellar populations and gas density.

The four classes of policies are *universal* - they cover every method for making decisions described in the research literature or used in practice.

Universal policies

This means you are already using one of the four classes of policies (or a hybrid) in your own decisions.

Designing policies

From the four classes, we can create hybrids:

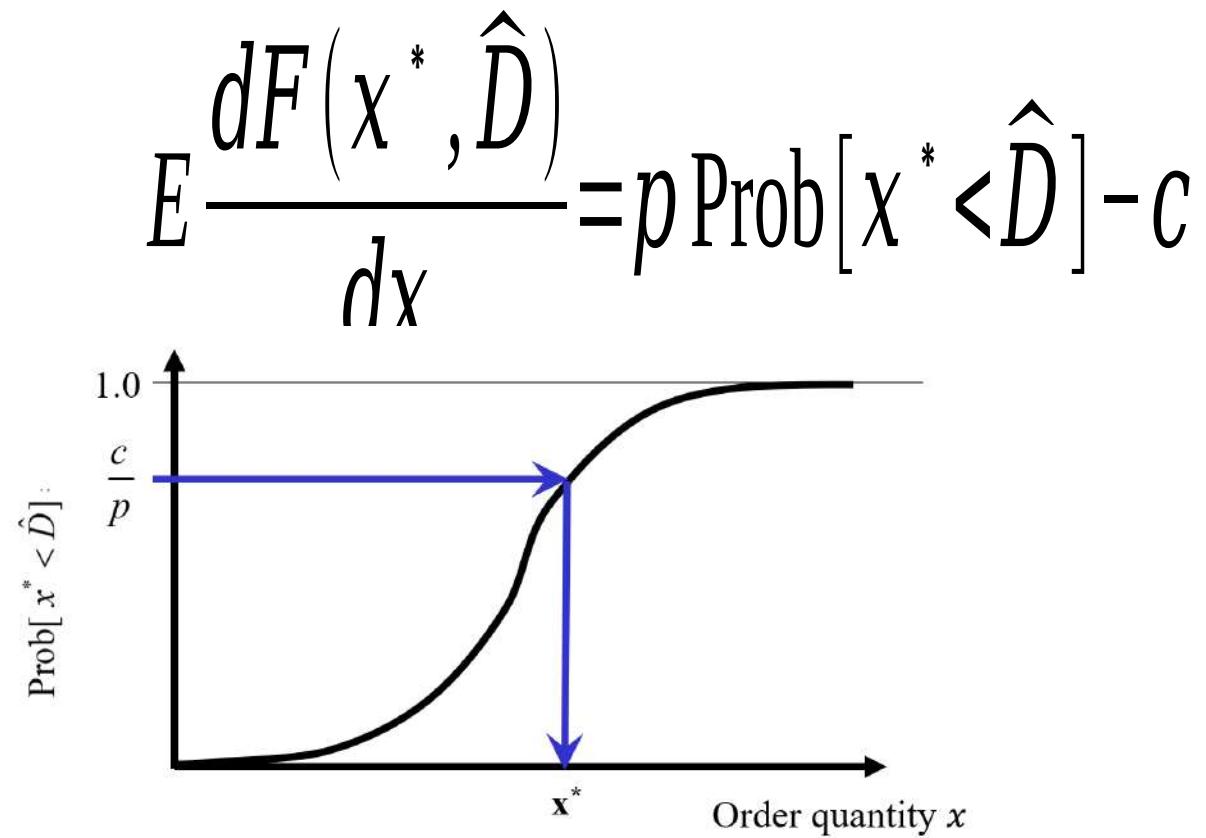
- CFAs with PFAs
 - Choosing least cost supplier, but with rules to exclude high-risk companies.
- Lookahead (DLA) with VFA
 - Optimize seasonal production plan, with functions capturing value of ending inventories
- Parameterized deterministic direct lookaheads (DLA/CFA)
 - Plan seasonal production plan, using -percentile (say, 80th percentile) demand forecasts.
- VFA policy using PFA
 - Distribution planning using VFAs to value inventory at each warehouse, but using rules (PFAs) to force deliveries to specific locations.
- VFA with CFA
 - Start with a VFA-based policy with a linear model, and then tune the parameters of the linear VFA to get the best results using a simulator.

The only way to create effective hybrids is to know all four classes of policies.

OUTLINE

- The seven levels of artificial intelligence
- The universal modeling framework
- Designing policies
- Mutual fund cash balance optimization
- Choosing the best policy
- A new educational field: sequential decision analytics

The newsvendor problem



Th



- The student was taught a stylized problem (the “newsvendor problem”) which the professor solved in class.
- The student was not taught how to *think* about the problem in a way that allowed him to go after the real-world problems that arise in practice.

Th



Professor – The newsvendor inventory problem we learned in class looks like a way to optimize how much cash I should hold in the mutual fund I manage.

The problem

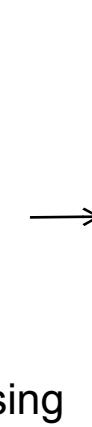
Mutual funds hold a certain percentage of their assets in cash to meet redemption requests from investors.

We can start by using their existing policy (a form of PFA):

$$X^\pi(S_t | \theta) = \theta^{\text{asset}} \cdot \text{asset } s_t$$

... which can be tuned in a simulator or the field:

$$\max_{\theta^{\text{asset}}} \mathbb{E} \left\{ \sum_{t=0}^T C(S_t, X^\pi(S_t | \theta^{\text{asset}})) | S_0 \right\}$$



Computer simulation



Field experience



Computer simulations can be performed using historical data. Simulations, as with field testing, can be completely general. There is no need for simplifying assumptions.

Th

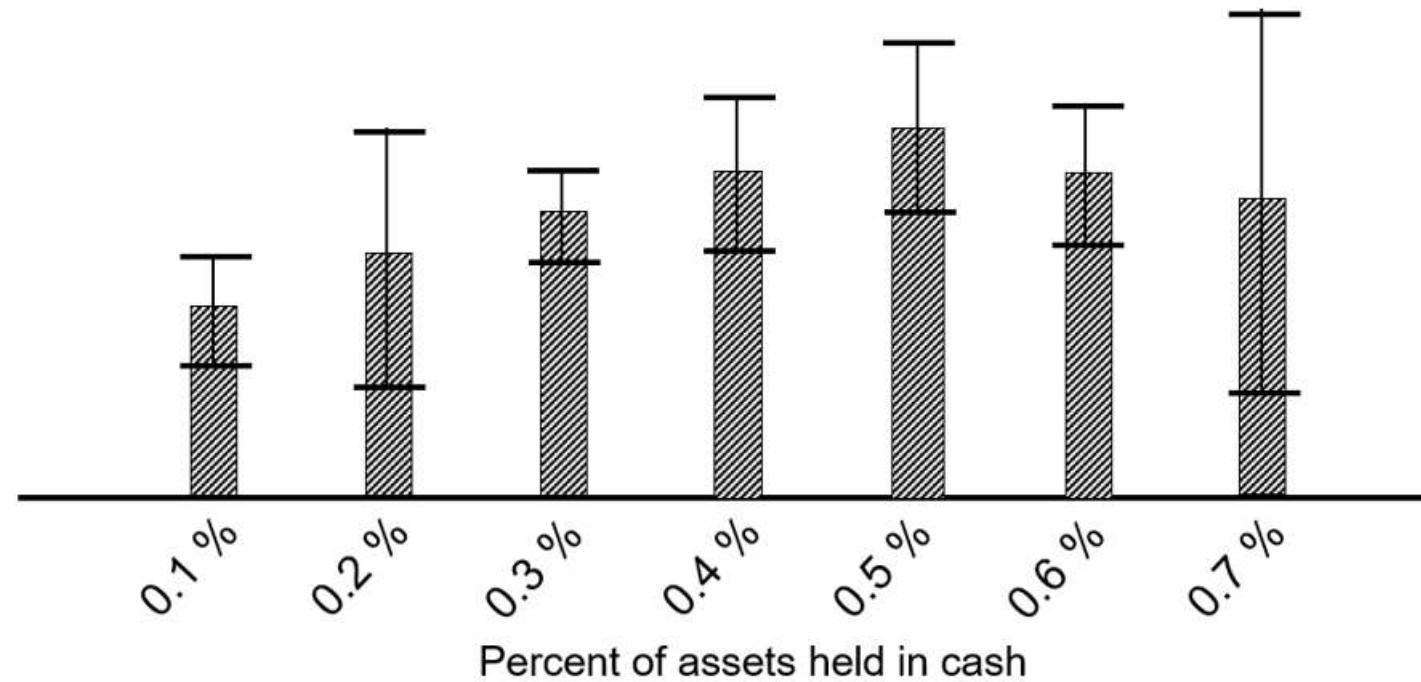
Professor – The newsvendor inventory problem we learned in class looks like a way to optimize how much cash I should hold in the mutual fund I manage.



The problem

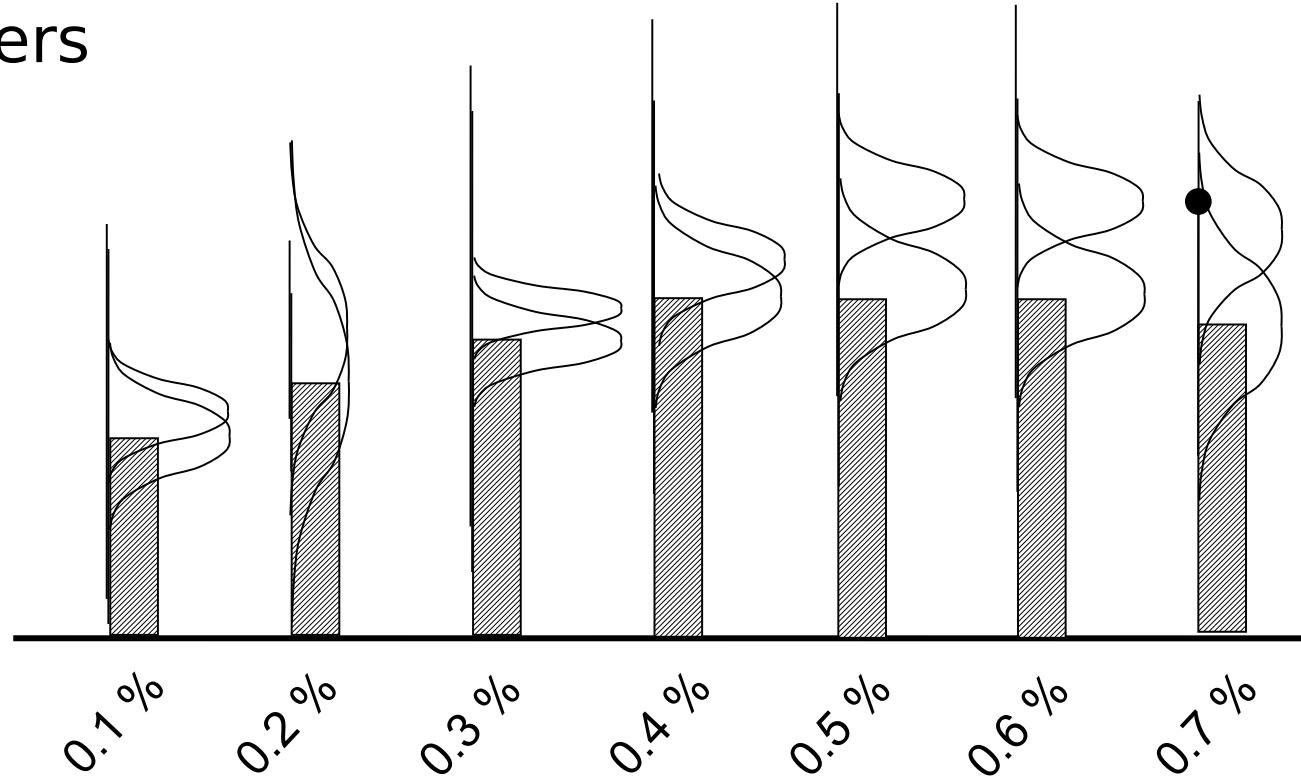
Mutual funds hold a certain percentage of their assets in cash to meet redemption requests from investors.

Intelligent trial and error (optimal learning)



Cash management problem

Correlated beliefs: Testing one parameter value teaches us about parameters



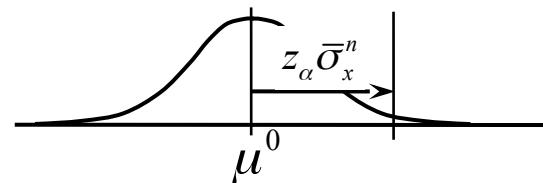
Cash management problem

Cost function approximations (CFA)

- » Upper confidence bounding

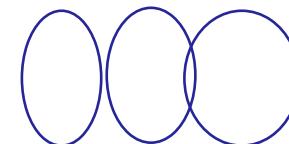
$$X^{UCB}(S^n | \theta^{UCB}) = \arg \max_x \left(\bar{\mu}_x^n + \theta^{UCB} \sqrt{\frac{\log n}{N_x^n}} \right)$$

- » Interval estimation



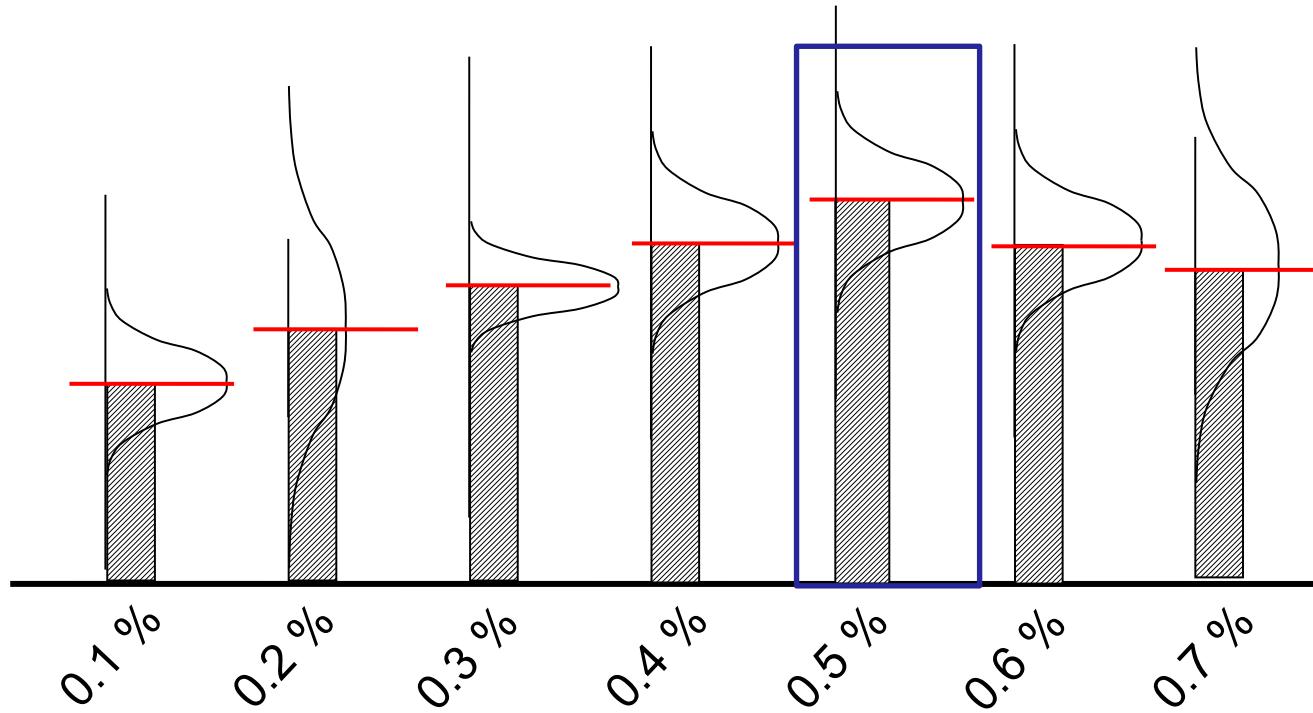
$$X^{IE}(S^n | \theta^{IE}) = \arg \max_x \left(\bar{\mu}_x^n + \theta^{IE} \bar{\sigma}_x^n \right)$$

- » Thompson sampling



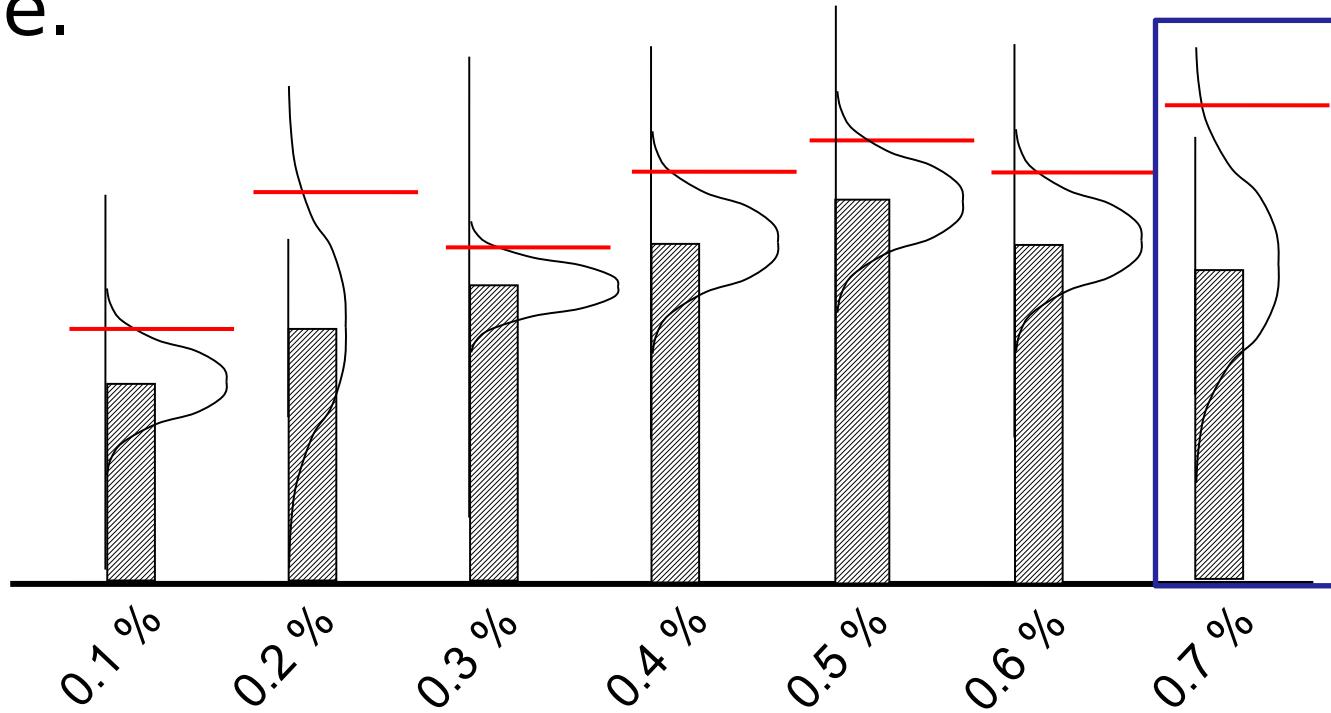
Cash management problem

Picking means we are evaluating each choice at the mean.



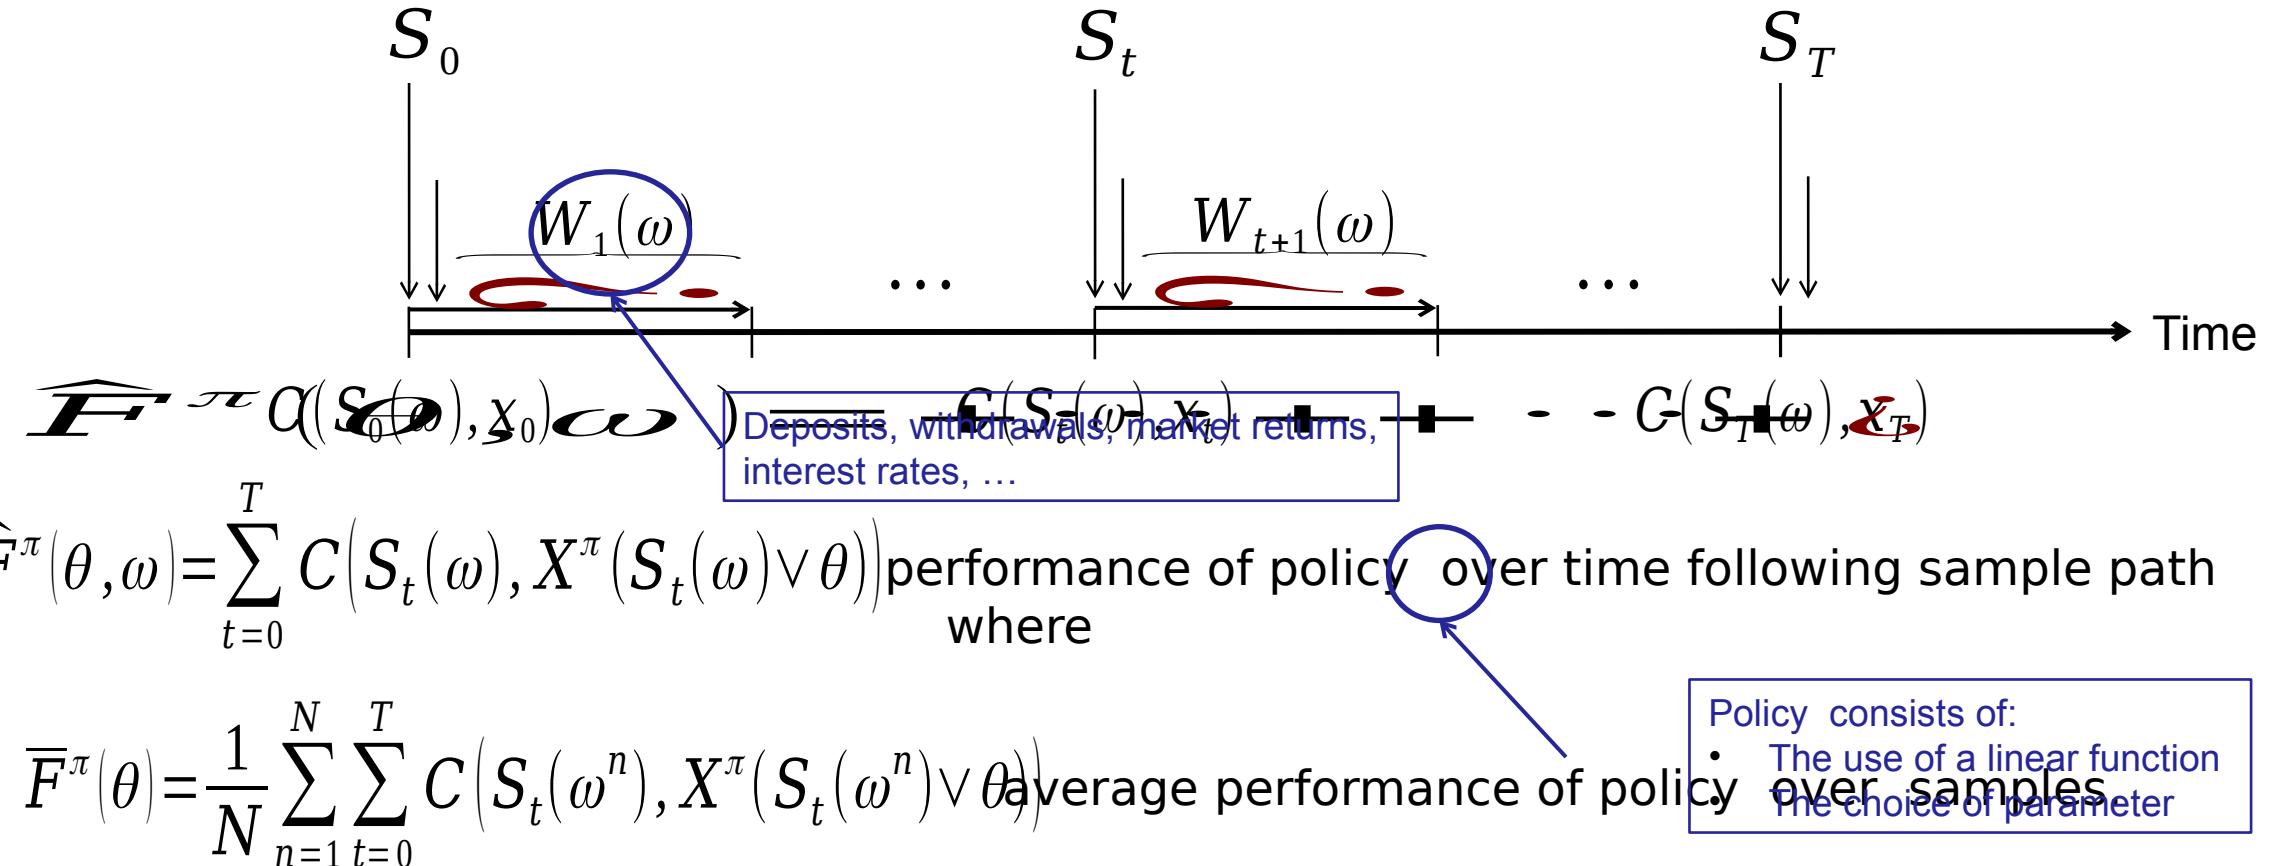
Cash management problem

Picking means we are evaluating each choice at the 95th percentile.



Cash management problem

Evaluating performance over time:



Cash management problem

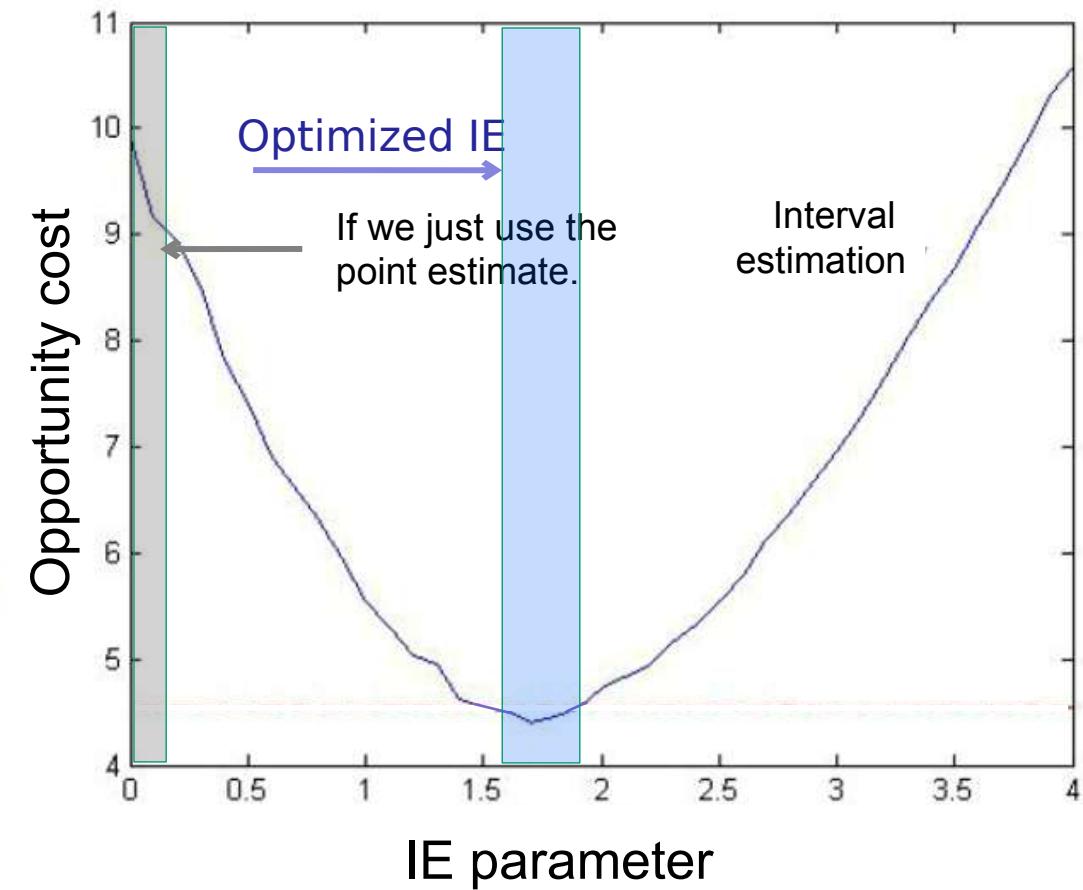
- Optimizing the IE policy
 - » We optimize to maximize:

$$\max_{\theta^{IE}} F(\theta^{IE}) = \mathbb{E}F(x^{\pi, N}, W)$$

where

$$x^n = X_{\square}^{IE}(S^n \vee \theta^{IE}) = \operatorname{argmax}_x (\bar{\mu}_x^n + \theta^{IE} \bar{\sigma}_x^n)$$

- Notes:
 - » Interval estimation is a form of cost function approximation (CFA)



The price of simplicity is tunable parameters ... *and tuning is hard!*

Th



Professor – The newsvendor inventory problem we learned in class looks like a way to optimize how much cash I should hold in the mutual fund I manage.

The problem

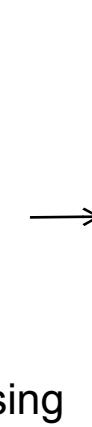
Mutual funds hold a certain percentage of their assets in cash to meet redemption requests from investors.

We can start by using their existing policy...

$$X^\pi(S_t | \theta) = \theta^{retail} \cdot asset s_t^{retail} + \theta^{inst} \cdot asset s_t^{inst}$$

... which can be tuned in a simulator or the field:

$$\max_{\theta^{asset}} \mathbb{E} \left\{ \sum_{t=0}^T C(S_t, X^\pi(S_t | \theta^{asset})) | S_0 \right\}$$



Computer simulation



Field experience



Computer simulations can be performed using historical data. Simulations, as with field testing, can be completely general. There is no need for simplifying assumptions.

Th

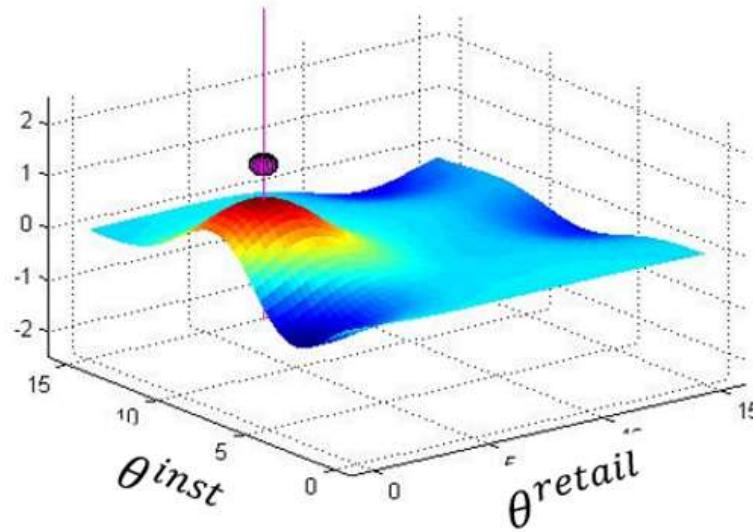
Professor – The newsvendor inventory problem we learned in class looks like a way to optimize how much cash I should hold in the mutual fund I manage.



The problem

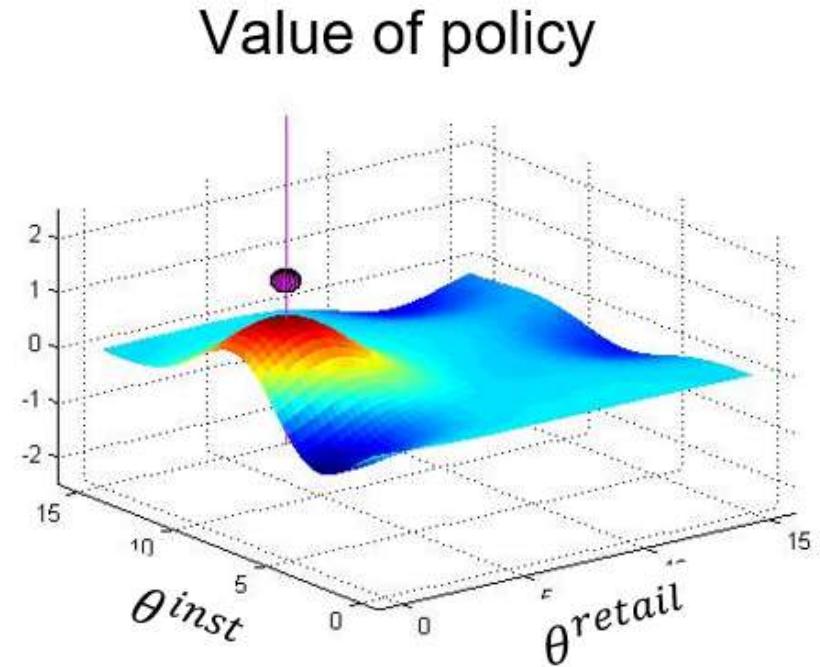
Mutual funds hold a certain percentage of their assets in cash to meet redemption requests from investors.

Value of policy



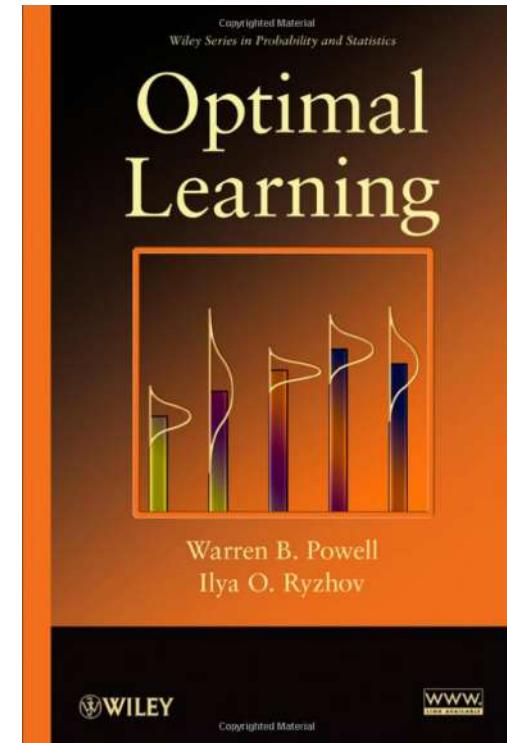
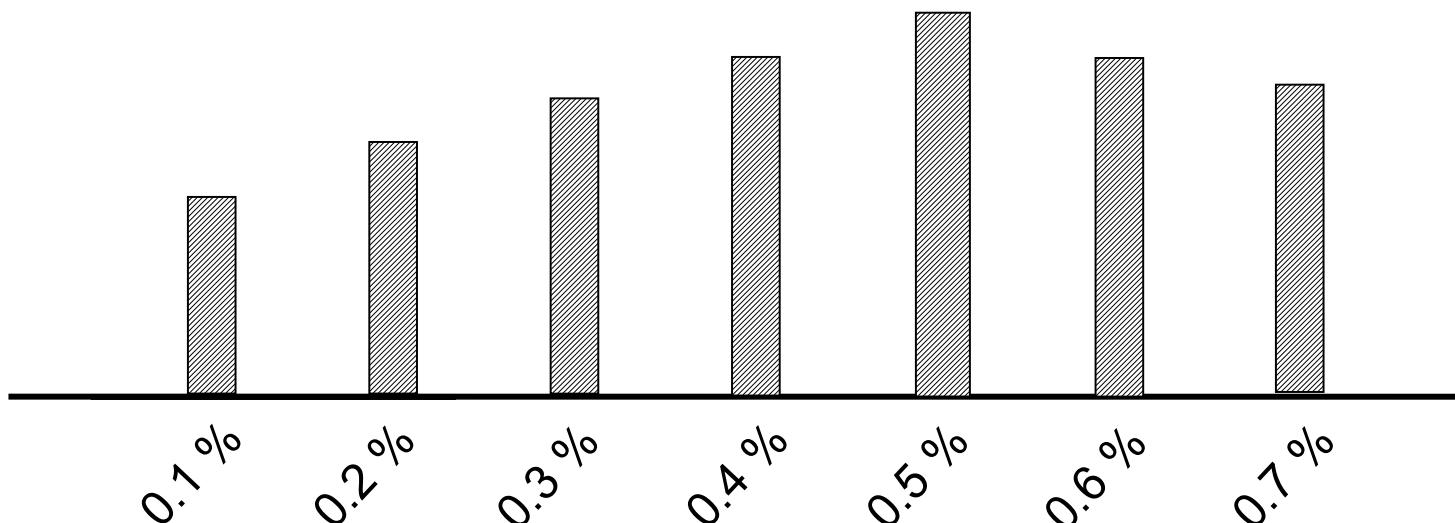
Cash management problem

- To find we have to search over a two-dimensional, discretized space.
- A form of direct lookahead policy, called the *knowledge gradient*, estimates the value of information from an experiment.



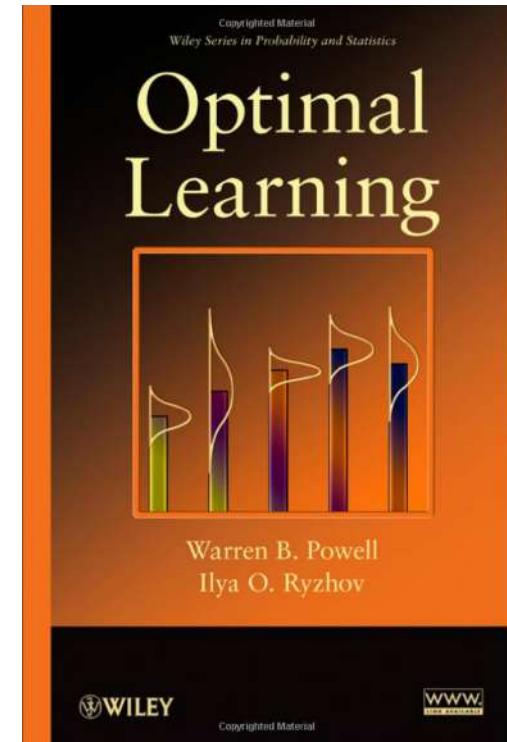
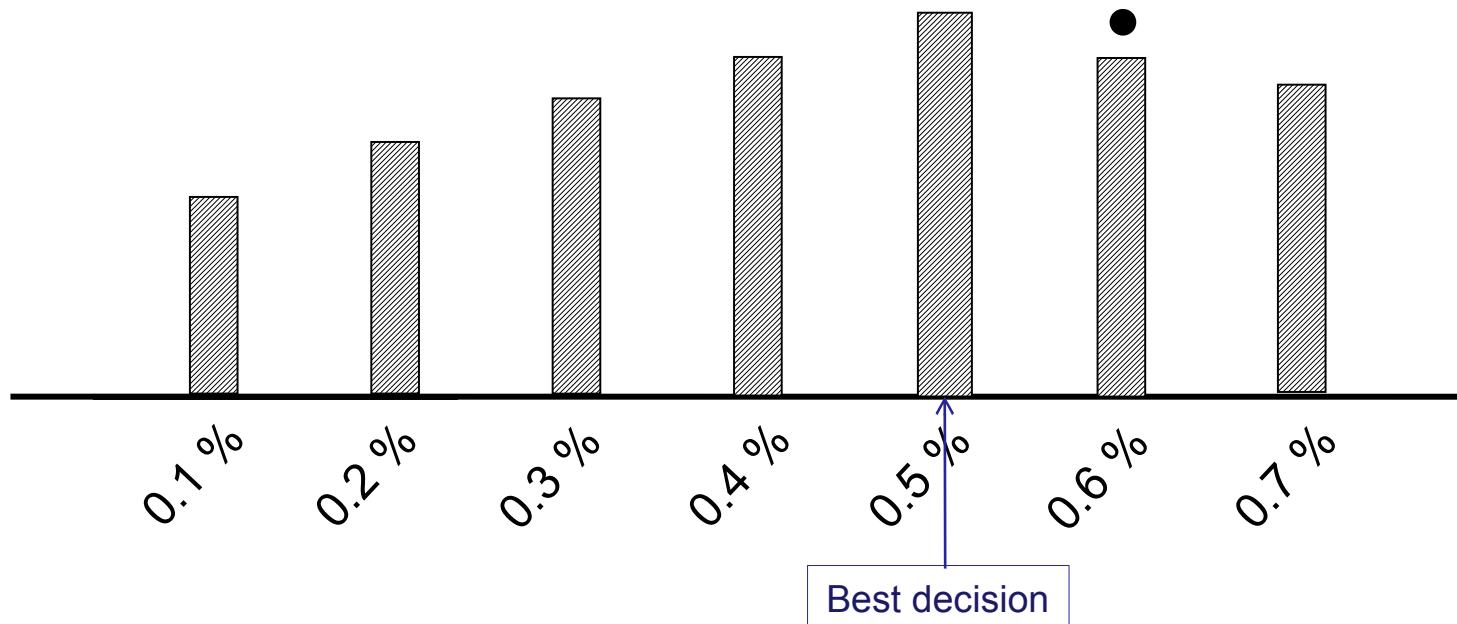
Cash management problem

We have some (possibly large) set of choices, with uncertain beliefs about how well each performs.



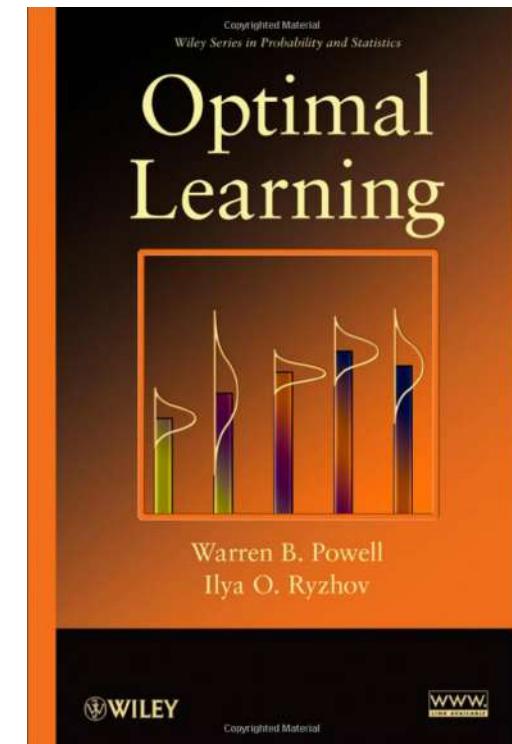
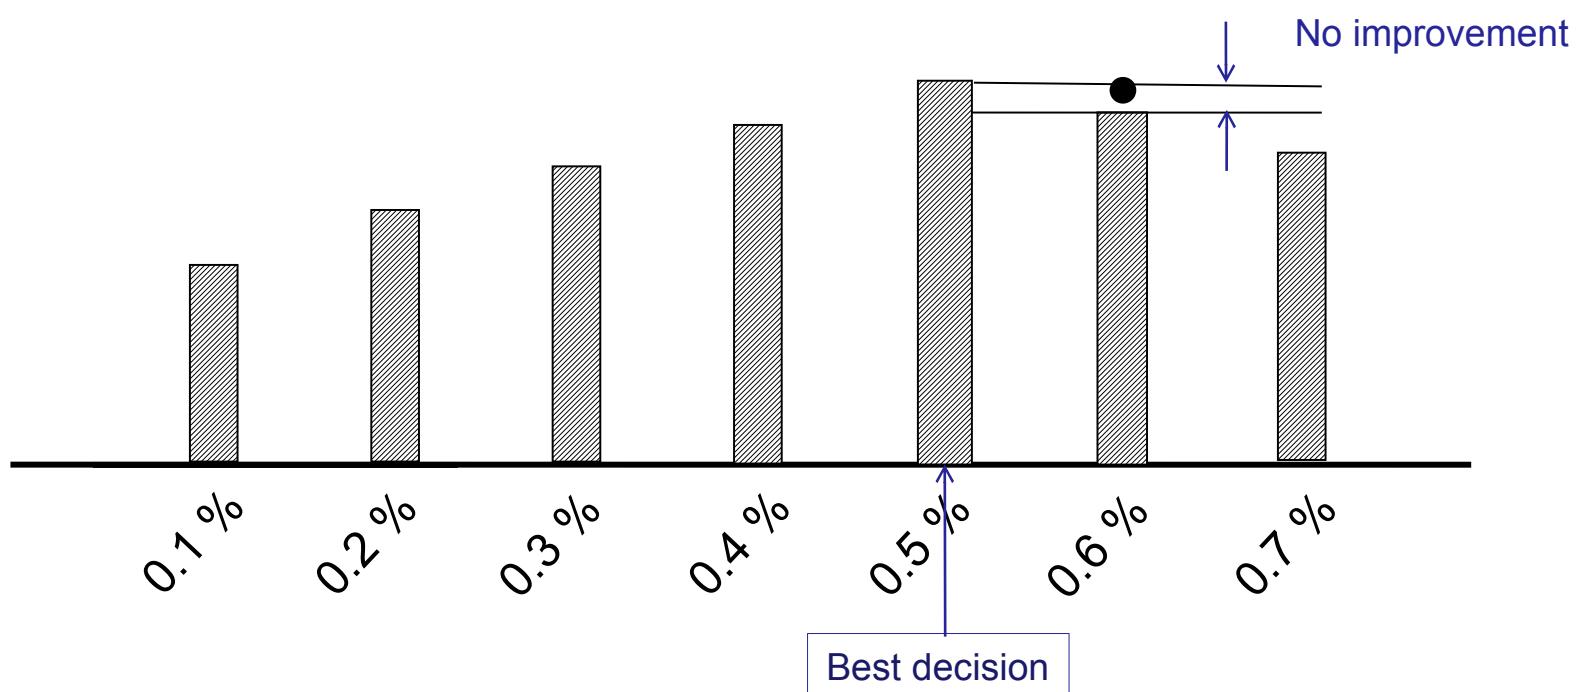
Cash management problem

We have some (possibly large) set of choices, with uncertain beliefs about how well each performs.



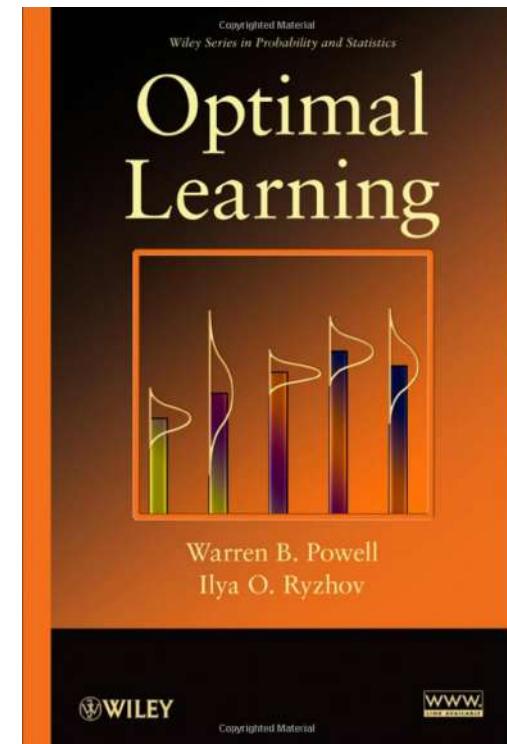
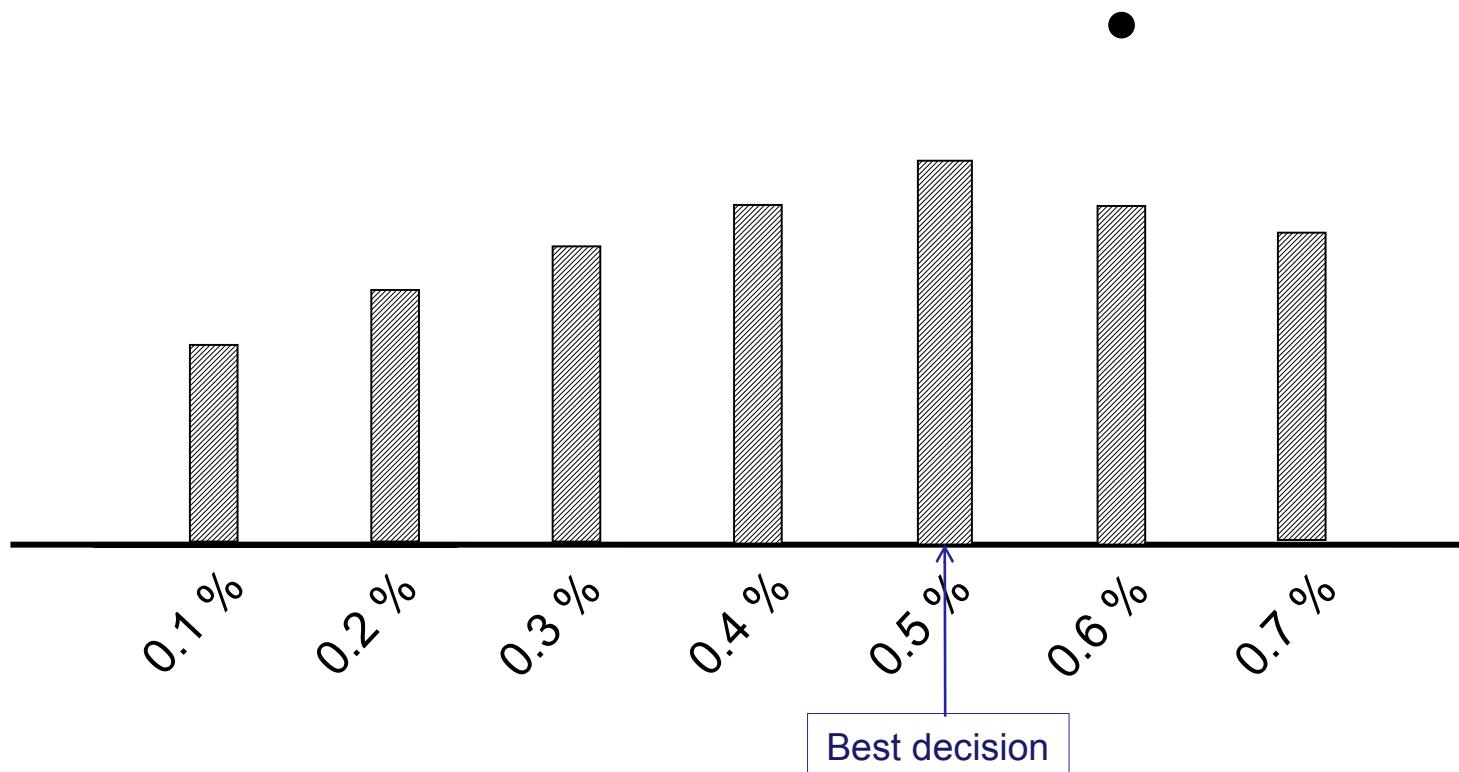
Cash management problem

We have some (possibly large) set of choices, with uncertain beliefs about how well each performs.



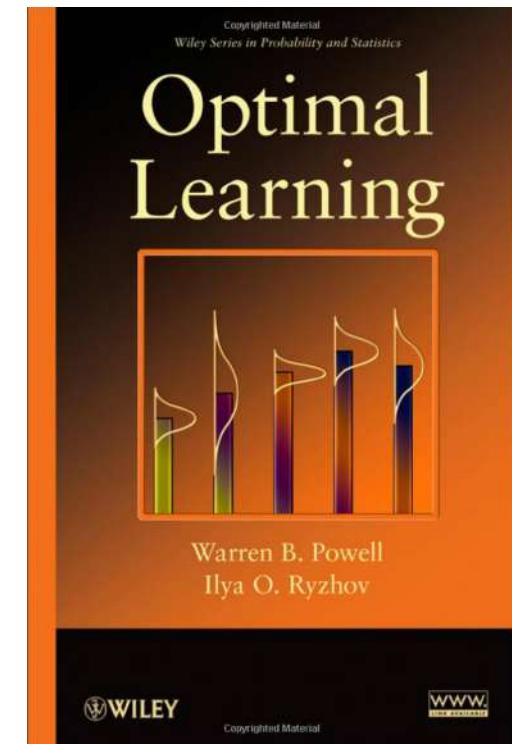
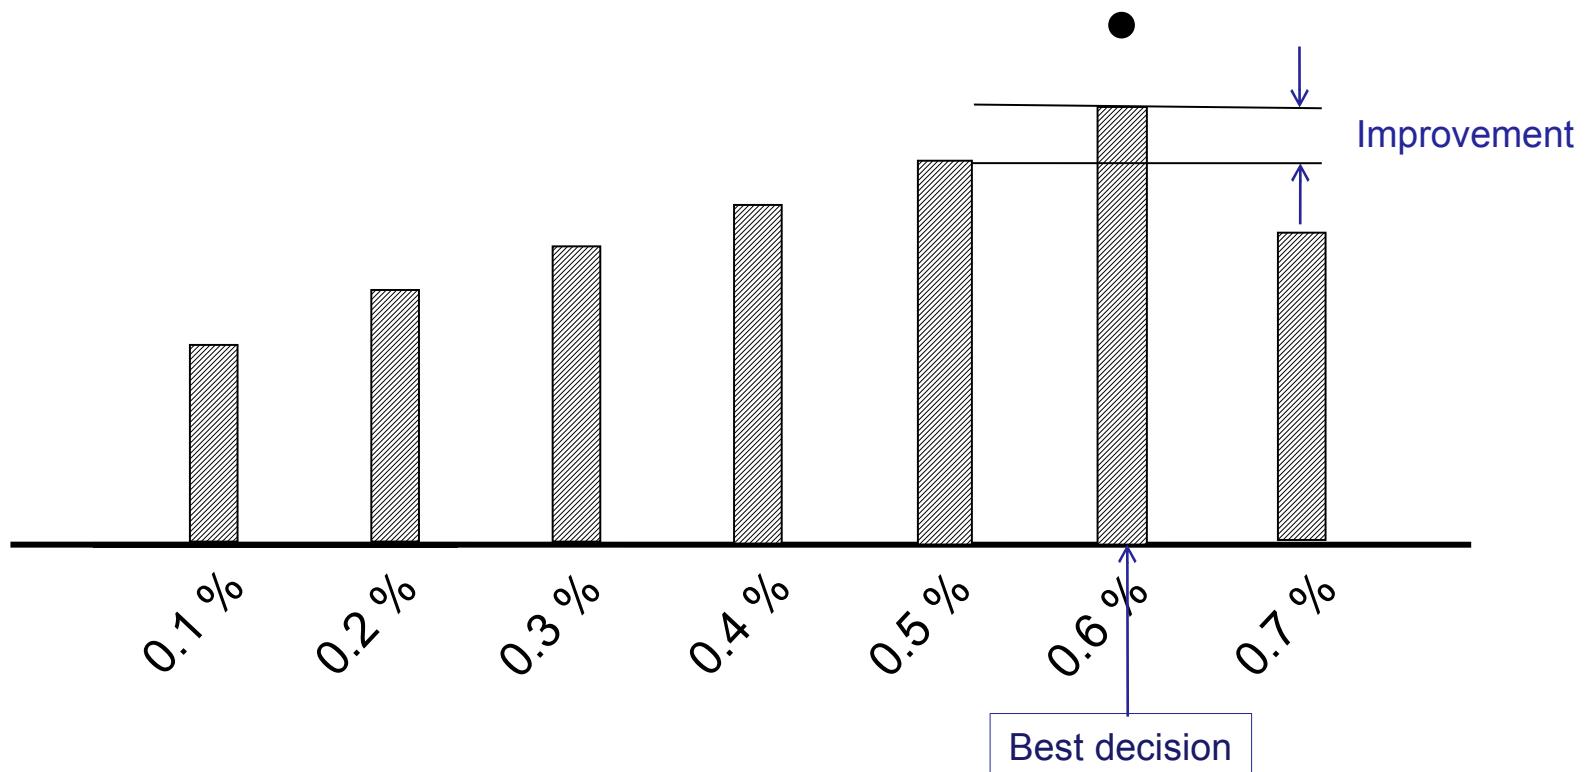
Cash management problem

We have some (possibly large) set of choices, with uncertain beliefs about how well each performs.



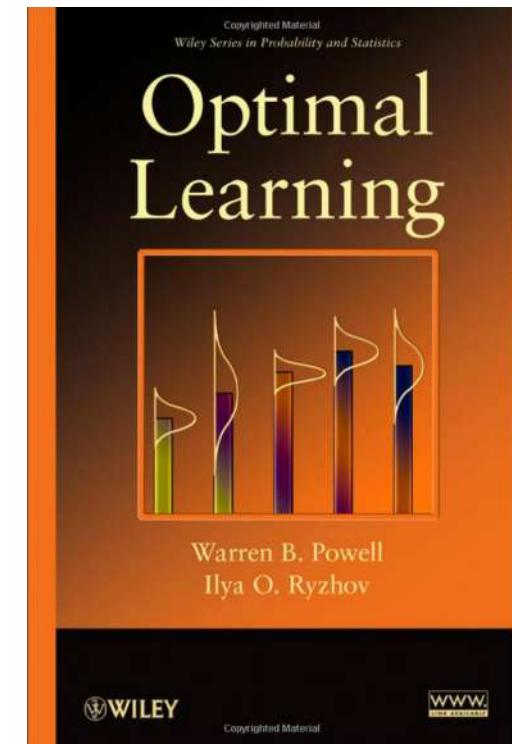
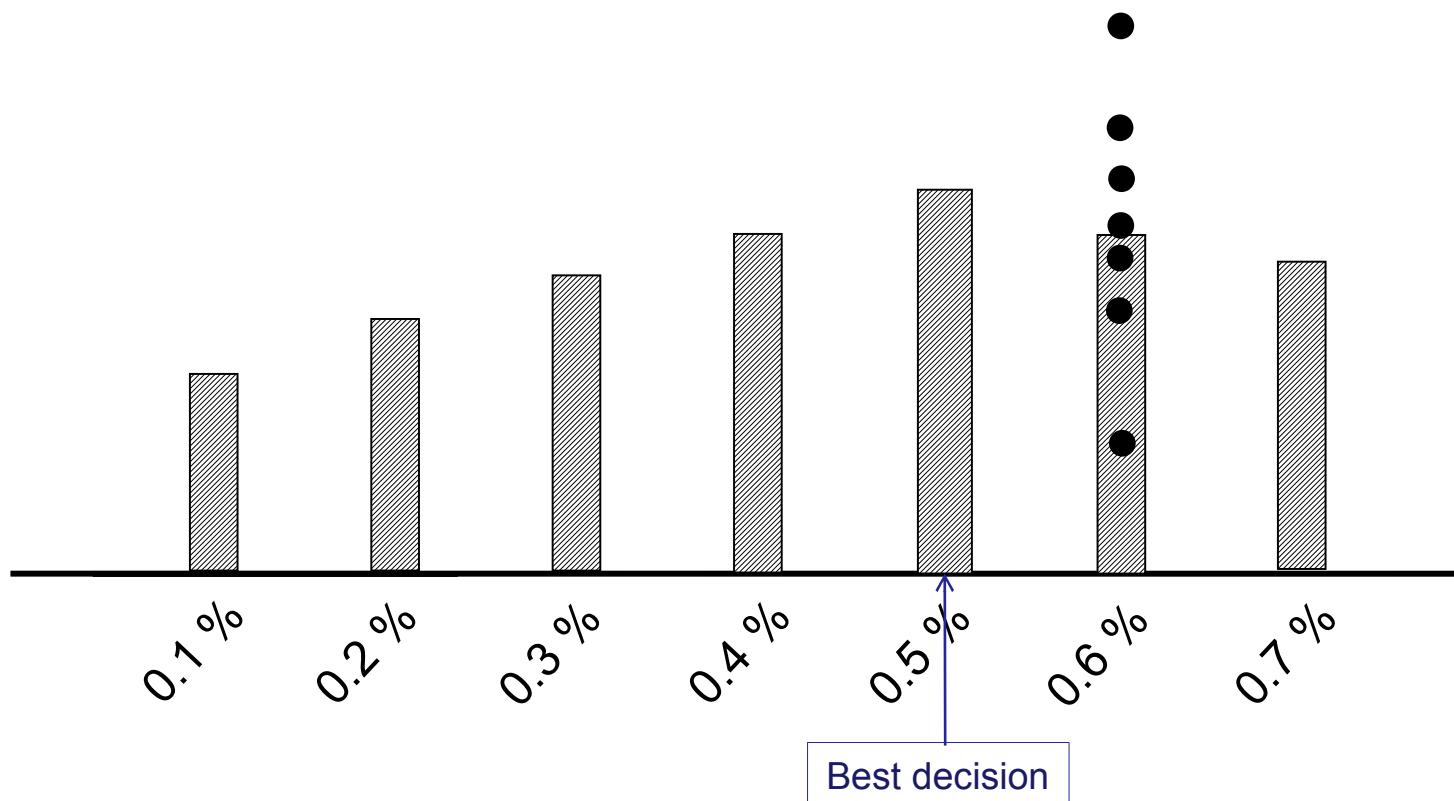
Cash management problem

We have some (possibly large) set of choices, with uncertain beliefs about how well each performs.



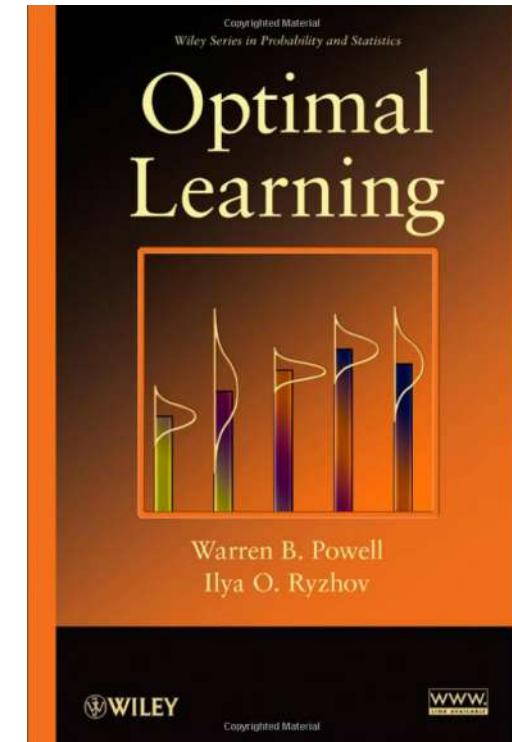
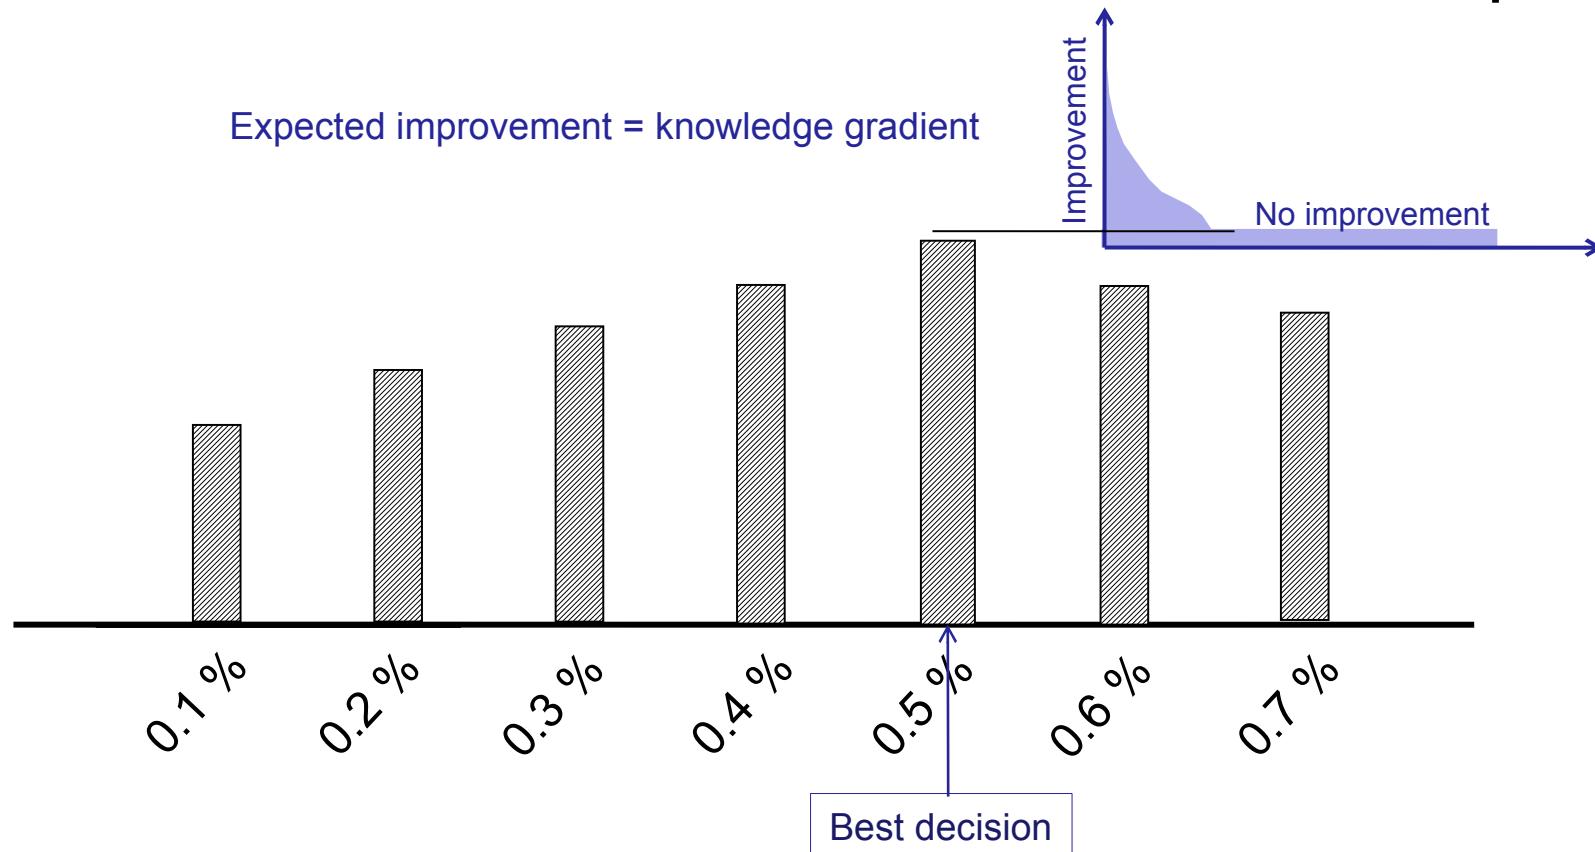
Cash management problem

We have some (possibly large) set of choices, with uncertain beliefs about how well each performs.



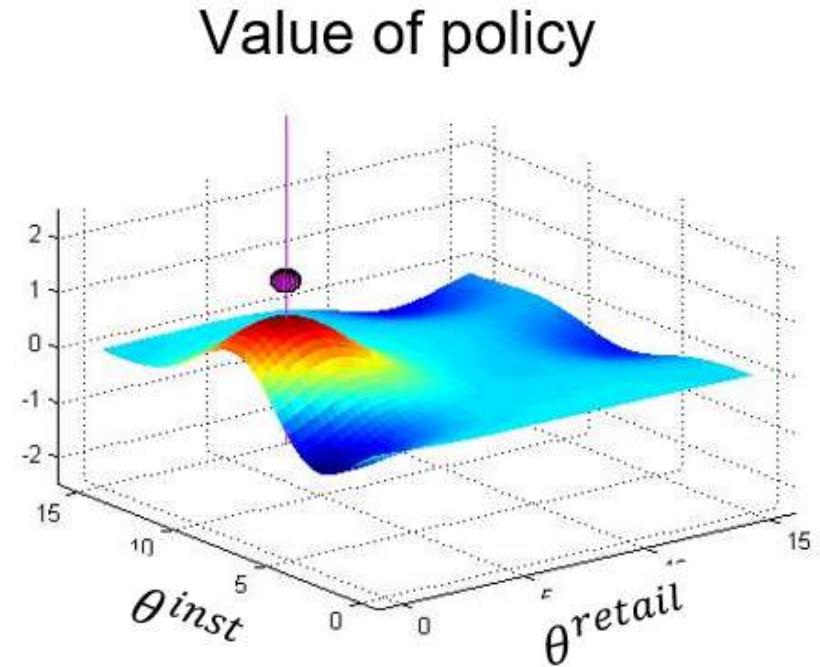
Cash management problem

We have some (possibly large) set of choices, with uncertain beliefs about how well each performs.



Cash management problem

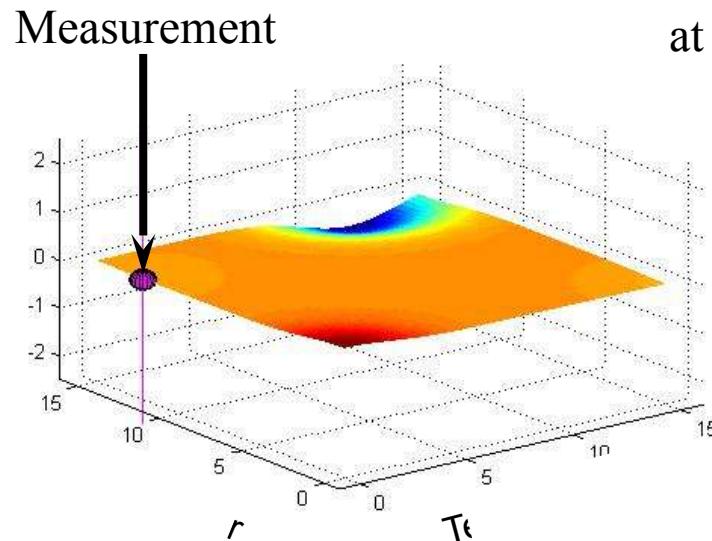
- Assume we discretize θ_{inst} and θ_{retail} into 50 intervals, we would have 2500 points to search over.
- We can easily calculate the knowledge gradient for each of the 2500 points, and then choose the point to observe that has the highest KG value.



Cash management problem

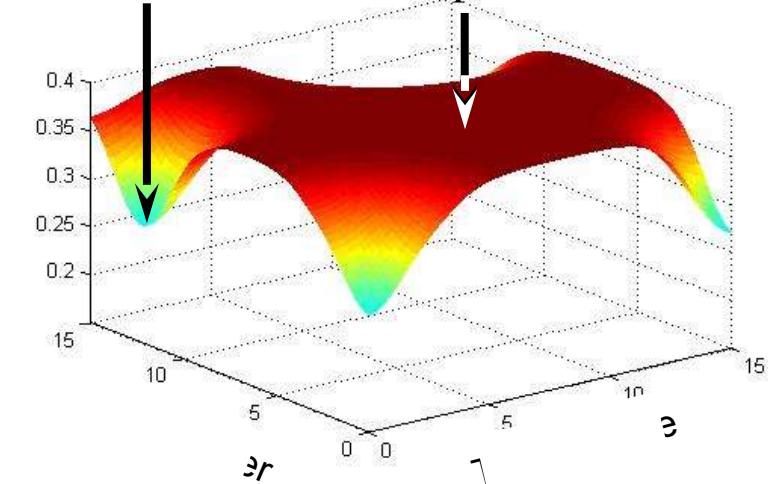
- After four experiments:

Estimated surface



Knowledge gradient

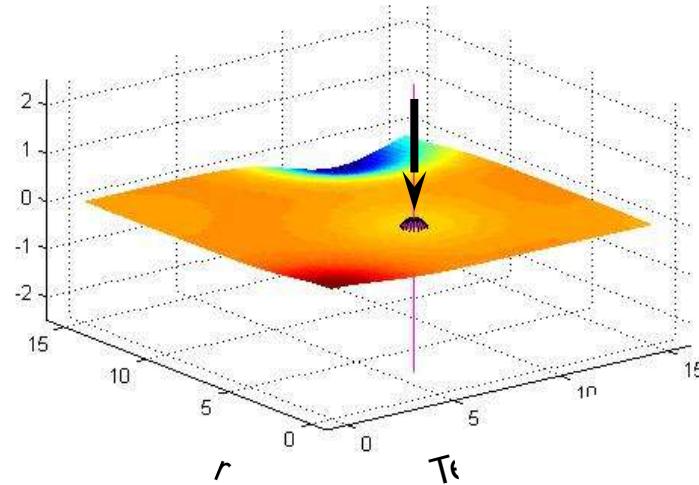
Value of another measurement
at same location.



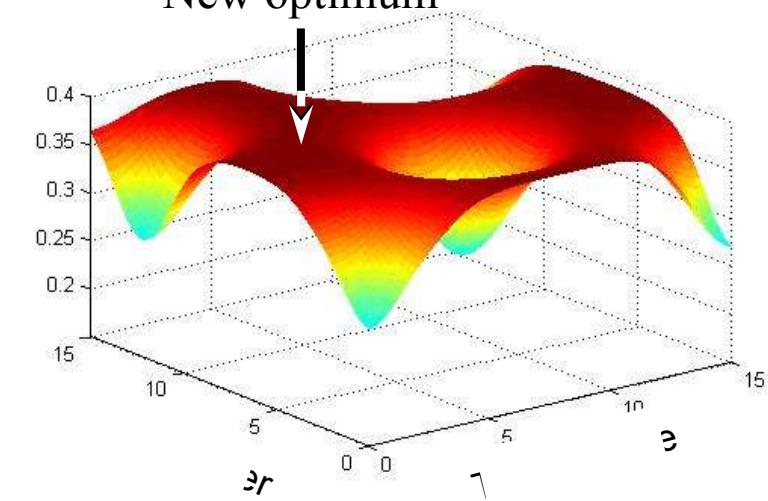
Cash management problem

- After five experiments:

Estimated surface



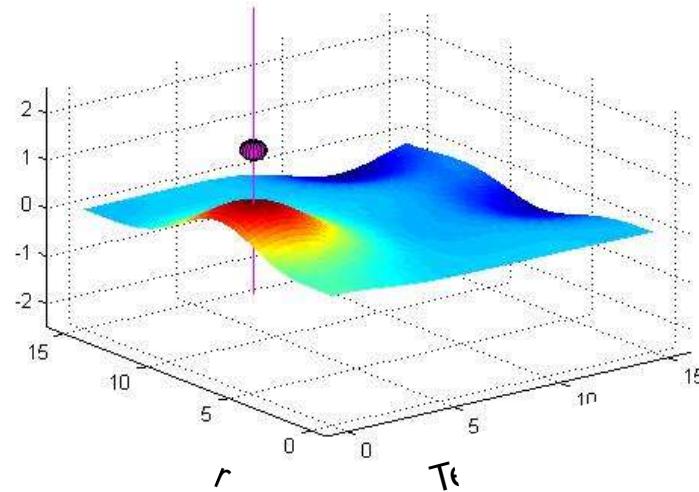
Knowledge gradient
New optimum



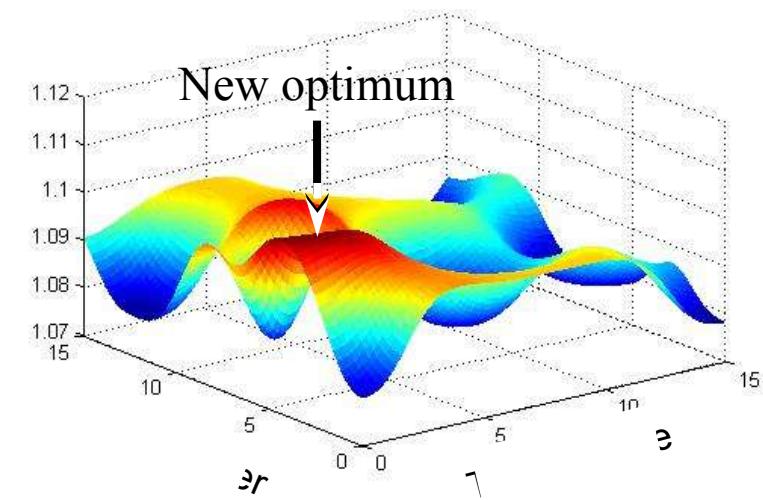
Cash management problem

- After seven experiments:

Estimated surface



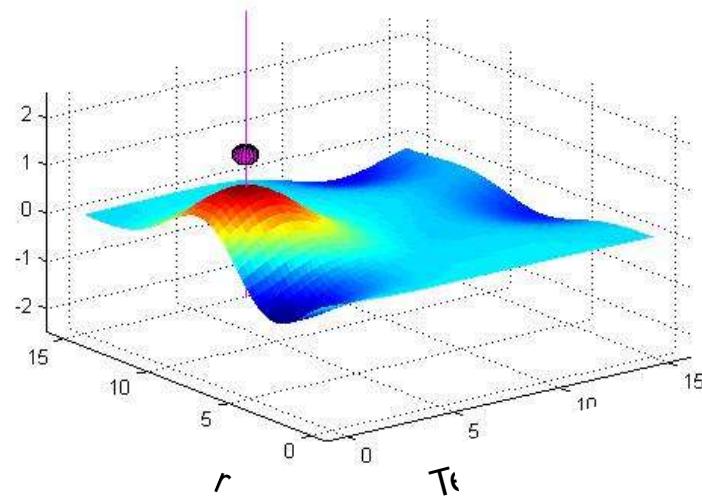
Knowledge gradient



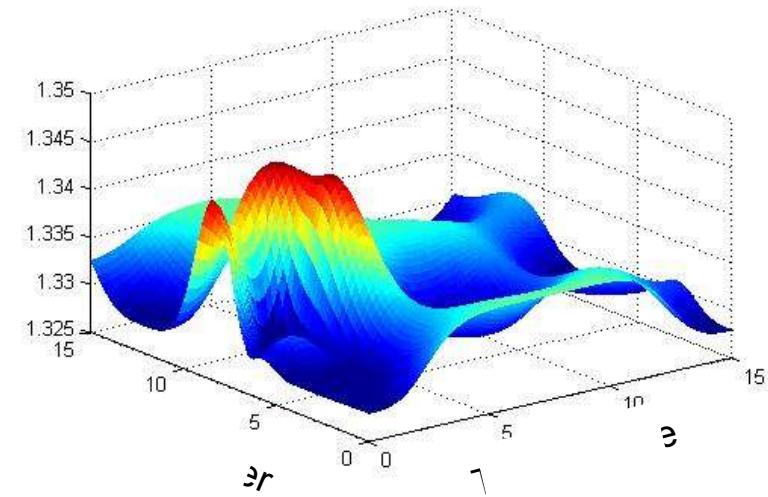
Cash management problem

- After nine experiments:

Estimated surface



Knowledge gradient



After nine experiments, we can identify the best out of 2500 combinations of wafer diameter and the temperature and the heating process.

Cash management problem

- Optimizing the IE policy
 - » We optimize to maximize:

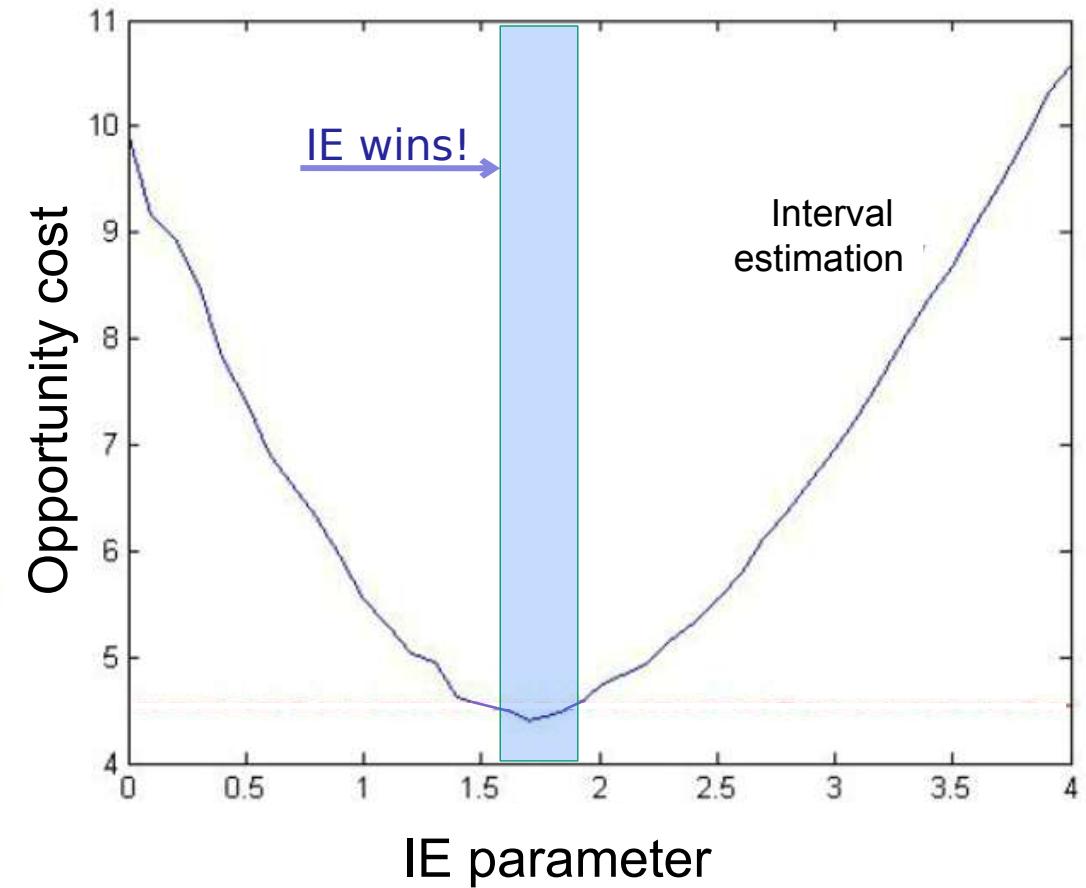
$$\max_{\theta^{IE}} F(\theta^{IE}) = \mathbb{E}F(x^{\pi, N}, W)$$

where

$$x^n = X_{\square}^{IE}(S^n \vee \theta^{IE}) = \operatorname{argmax}_x (\bar{\mu}_x^n + \theta^{IE} \bar{\sigma}_x^n)$$

- Notes:
 - » Interval estimation is a form of cost function approximation (FAA)
 - » Knowledge gradient is a form of direct lookahead approximation (DLA), but requires no tuning.

The price of simplicity is tunable parameters ... *and tuning is hard!*



Designing policies

Optimizing requires solving nested search problems:

- » Searching for the best cash management policy

- PFA? Linear decision rules:

$$X^\pi(S_t | \theta^{\text{asset}}) = \theta^{\text{asset}} \cdot \text{asset } s_t$$

$$X^\pi(S_t | \theta^{\text{asset}}) = \theta^{\text{retail}} \cdot \text{asset } s_t^{\text{retail}} + \theta^{\text{inst}} \cdot \text{asset } s_t^{\text{inst}}$$

- CFA? VFA (Bellman)? DLA (stochastic programming)?

- » Searching for the best parameter . Search policy:

- CFA (Interval estimation?)
 - DLA (knowledge gradient?)

These policies have their own tunable parameters

Designing policies

Optimizing requires solving nested search problems:

- » Searching for the best search policy

- CFA?
 - Interval estimation

- DLA?
 - Knowledge gradient

- » Searching for the best search parameter :

- PFA? Stochastic gradient algorithm
- CFA? (UCB, Interval estimation, Thompson sampling)
- DLA? (Knowledge gradient, expected improvement)

These policies have their own tunable parameters

Cash management problem

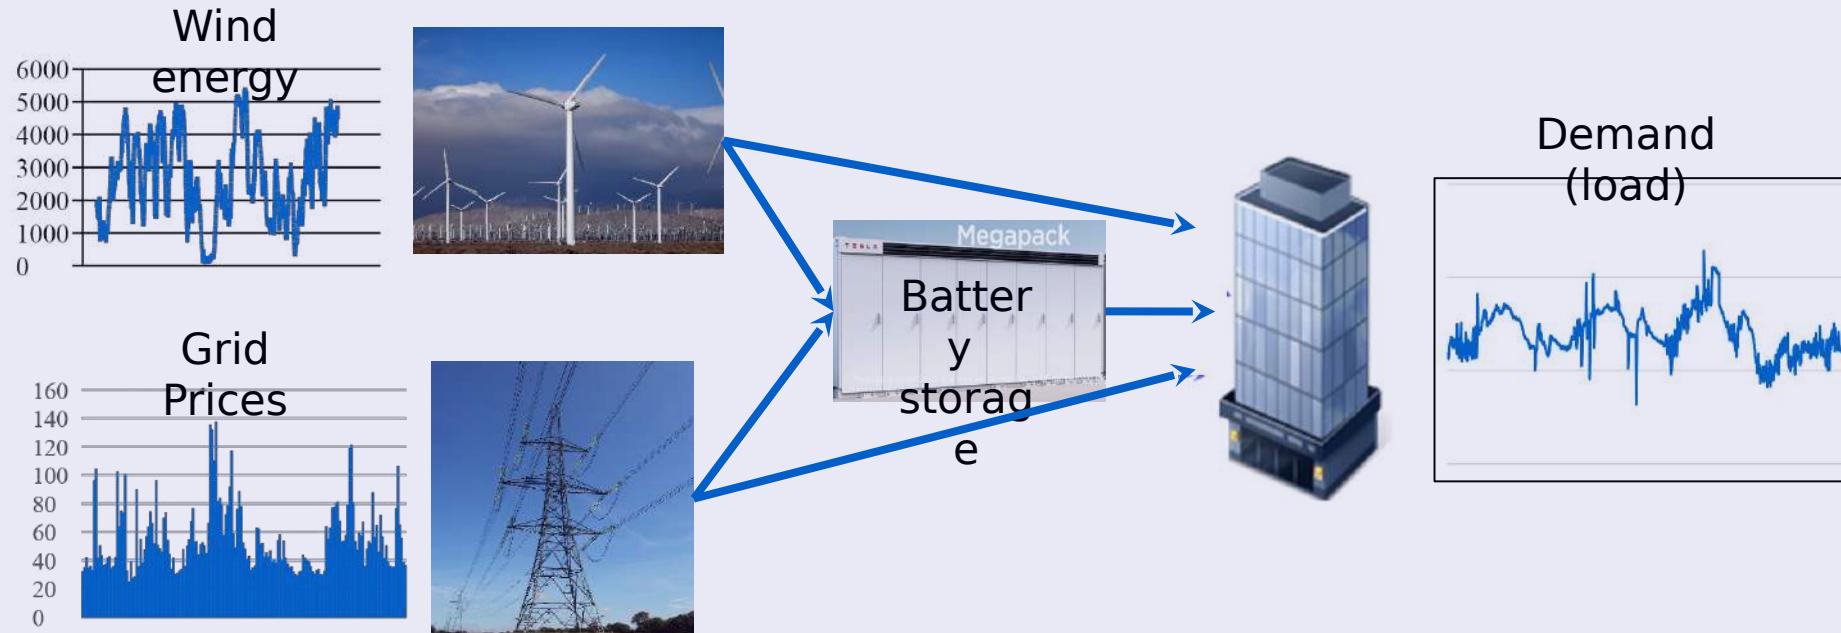
- Notes:
 - » The underlying dynamic model can be completely general.
 - » We can simulate any stochastic process for redemptions, deposits, and the evolution of the market.
- Future research:
 - » We can try any of a range of parametric and nonparametric models for the policy:
 - Tree regression
 - Neural networks
 - » We need additional research to guide the process of learning in the field, where policy search is slow, noisy, and costs are cumulative.

OUTLINE

- The seven levels of artificial intelligence
- The universal modeling framework
- Designing policies
- Mutual fund cash balance optimization
- Choosing the best policy
- A new educational field: sequential decision analytics

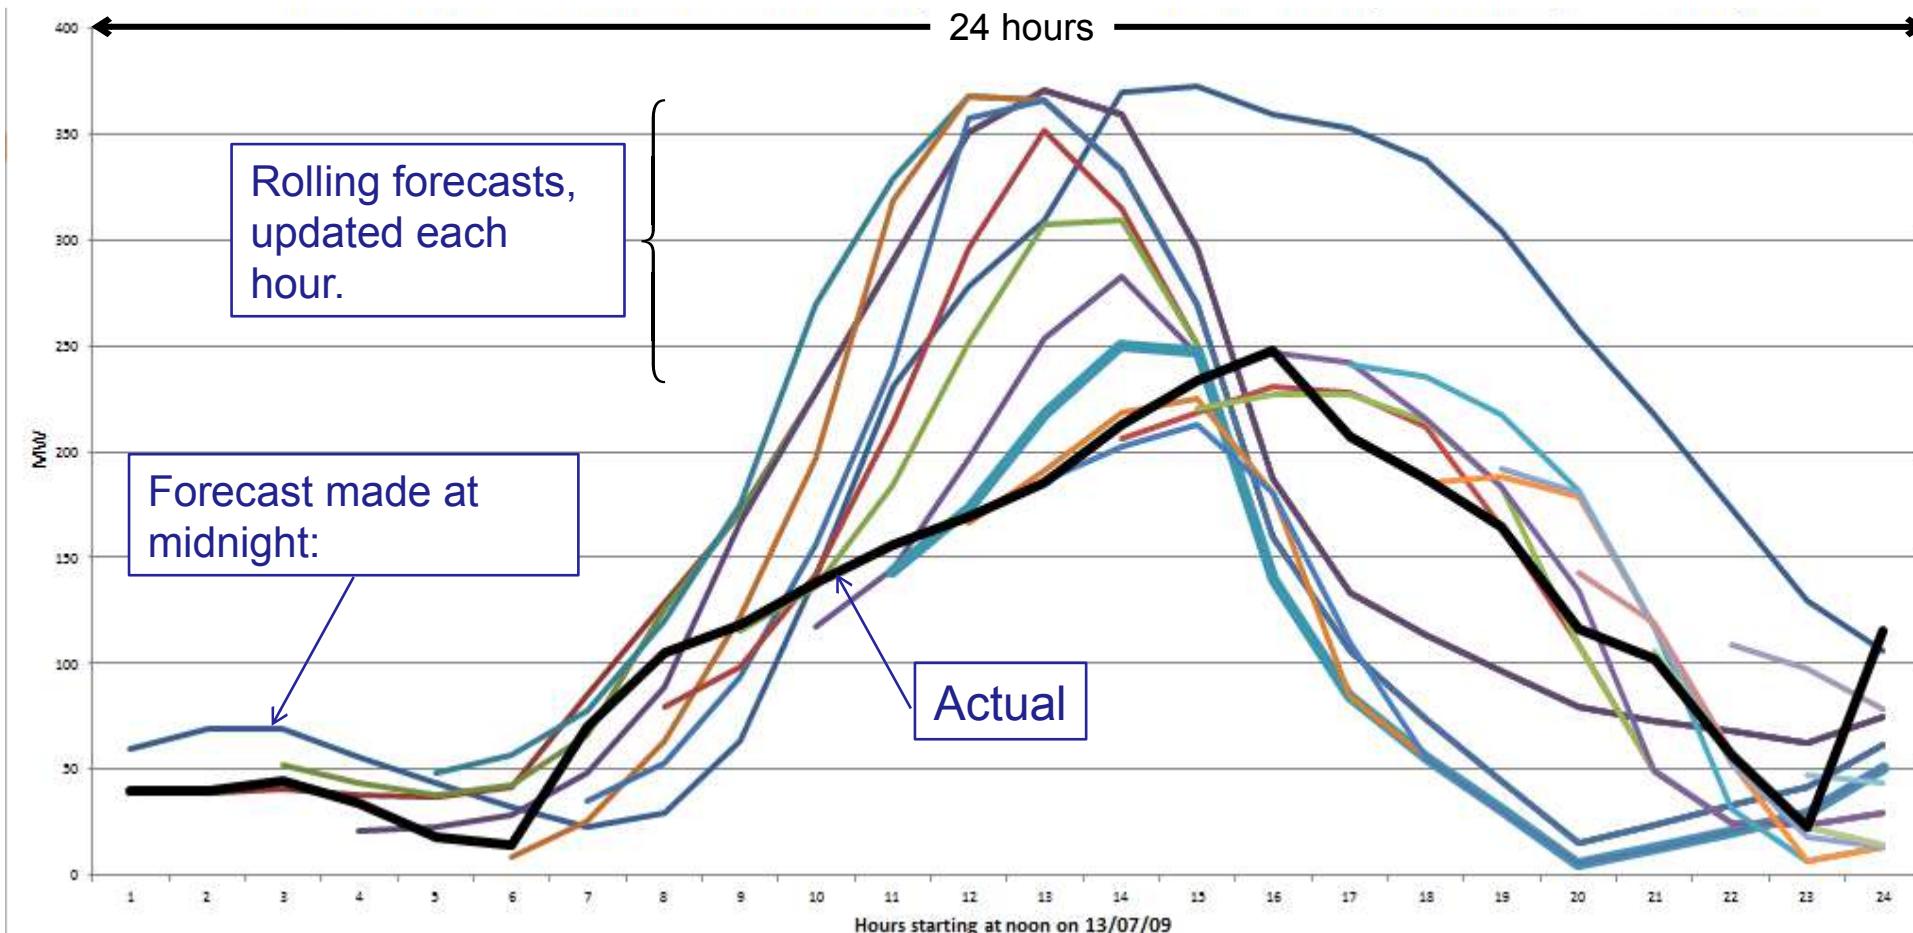
An energy storage application

Consider a basic energy storage problem



An energy storage application

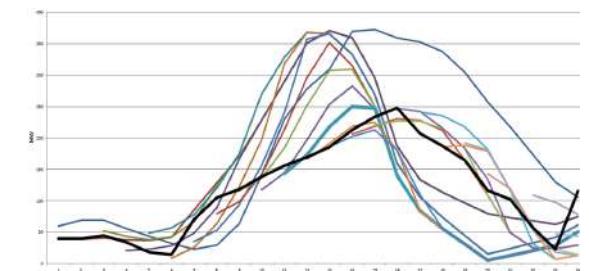
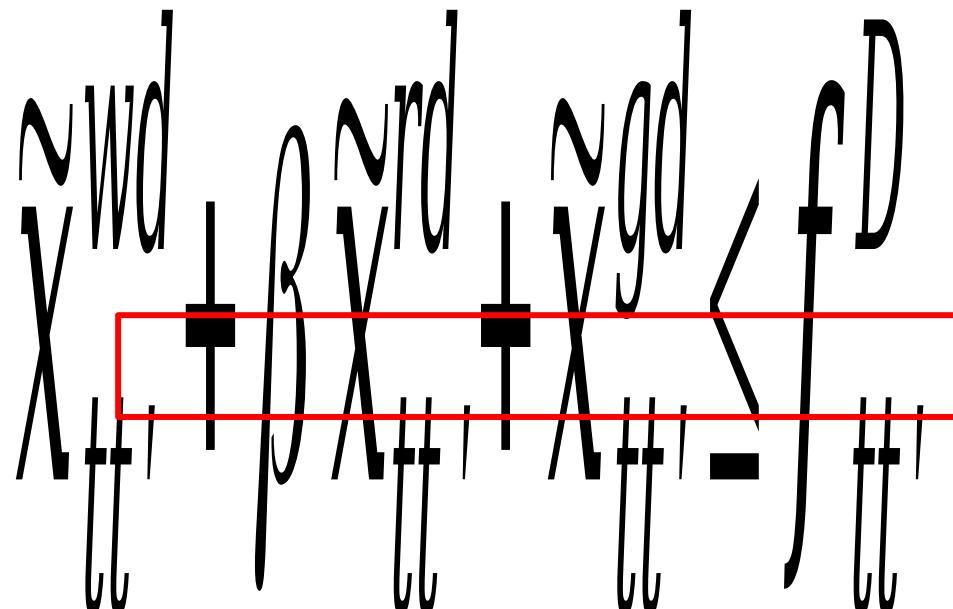
Forecasts evolve over time as new information arrives:



An energy storage application

- Benchmark policy – Deterministic lookahead

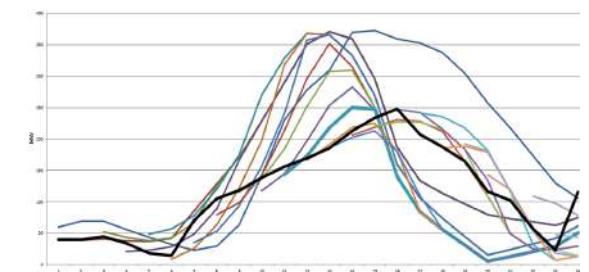
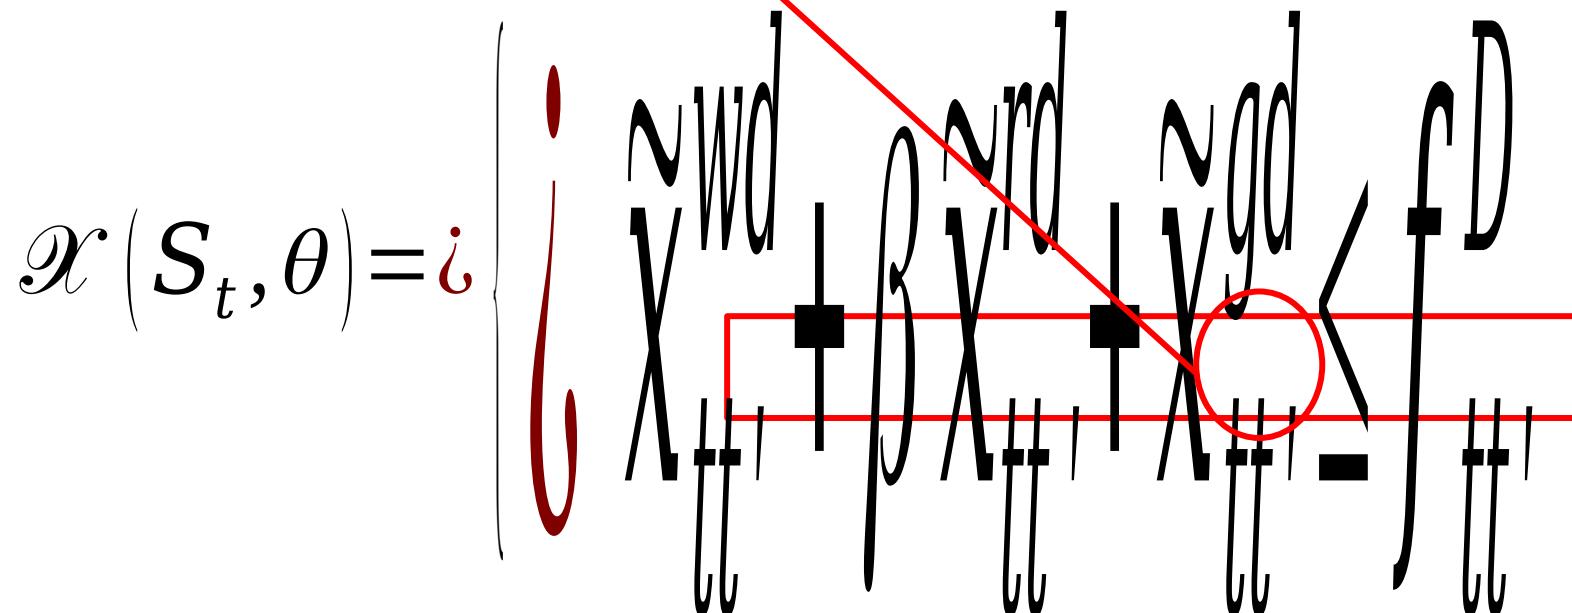
$$X_t^{D\text{-}LA}(S_t) = \arg \min_{x_t, (\tilde{x}_{tt'}, t'=t+1, \dots, t+H)} \left(C(S_t, x_t) + \left[\sum_{t'=t+1}^{t+H} \tilde{c}_{tt'} \tilde{x}_{tt'} \right] \right)$$



An energy storage application

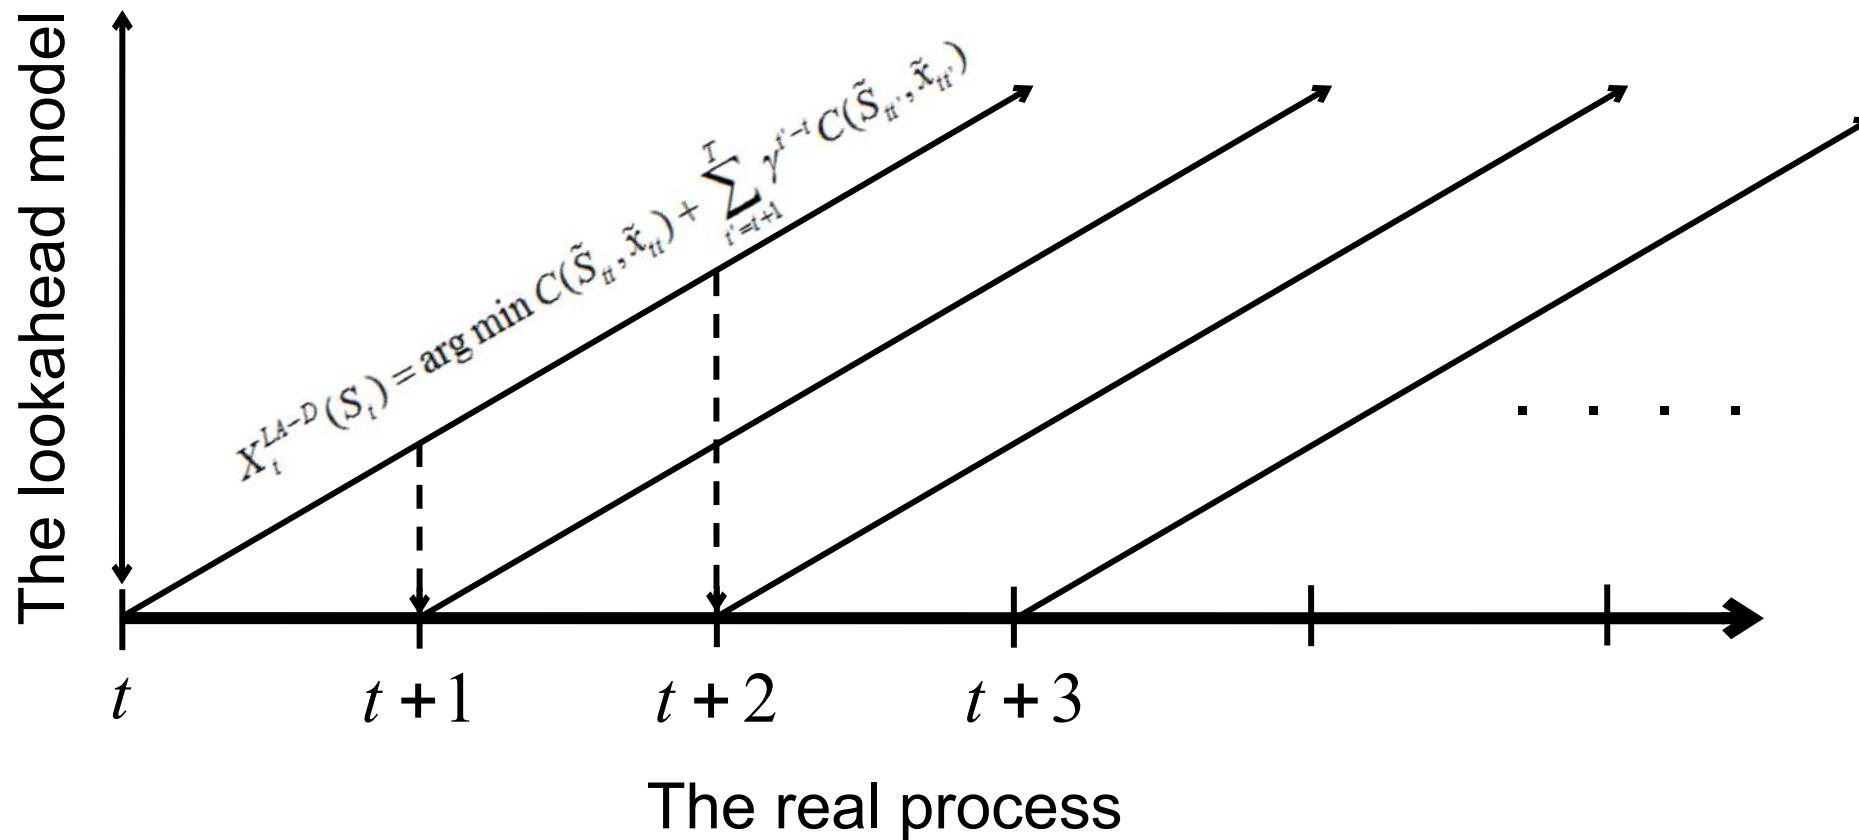
- Benchmark policy – Deterministic lookahead

$$X_t^{D\text{-}LA}(S_t | \theta) = \arg \min_{x_t, (\tilde{x}_{tt'}, t' = t+1, \dots, t+H)} \left(C(S_t, x_t) + \left[\sum_{t'=t+1}^{t+H} \tilde{c}_{tt'} \tilde{x}_{tt'} \right] \right)$$



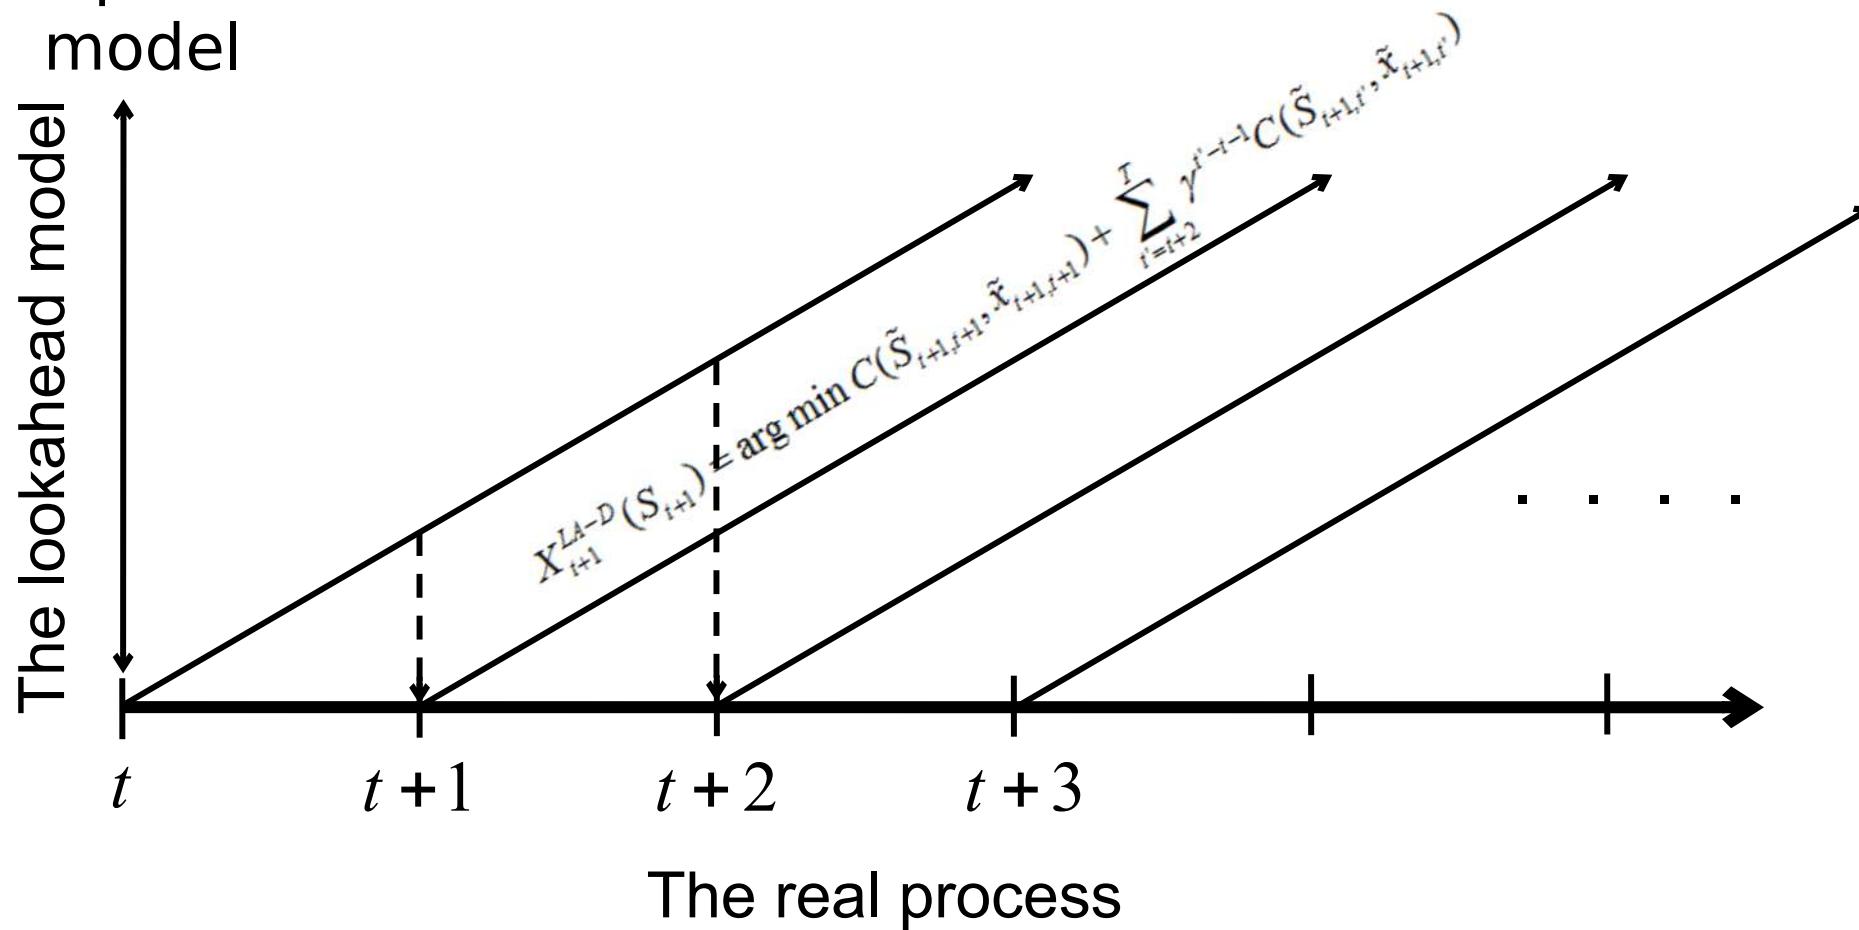
An energy storage application

- Lookahead policies peek into the future
 - » Optimize over deterministic lookahead model



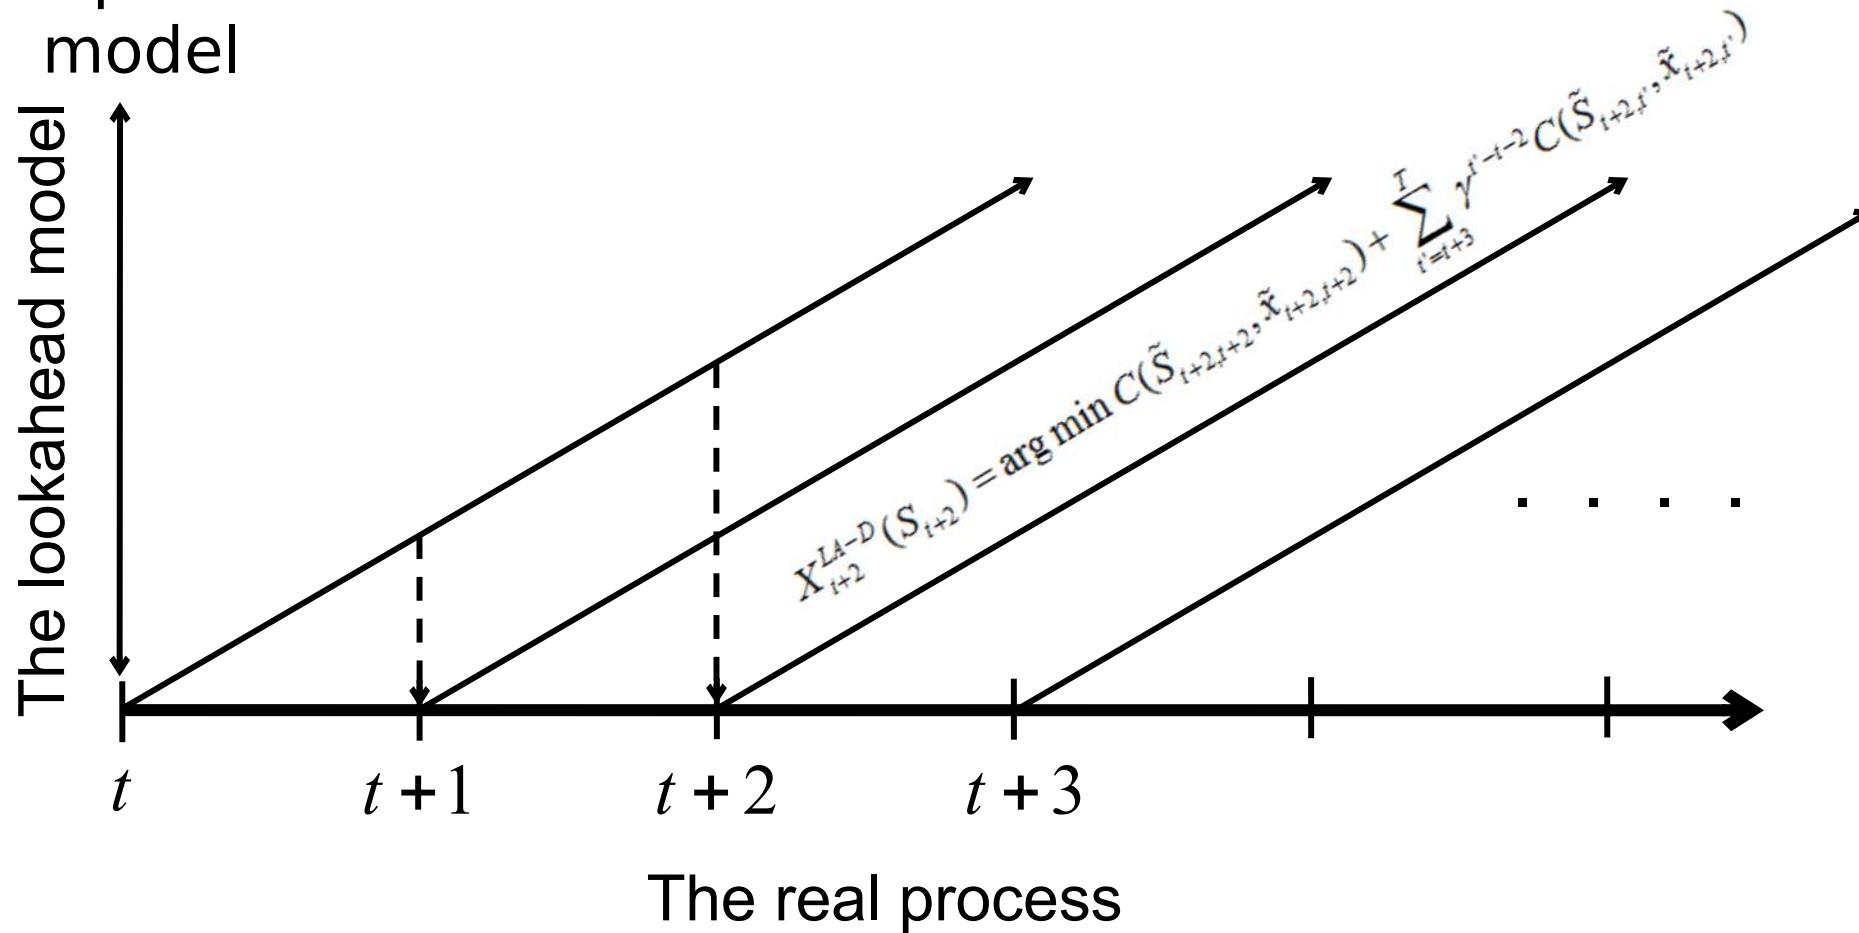
An energy storage application

- Lookahead policies peek into the future
 - » Optimize over deterministic lookahead model



An energy storage application

- Lookahead policies peek into the future
 - » Optimize over deterministic lookahead model



An energy storage application

- The policy (deterministic linear program)

$$x_t^\pi(s_t \vee \theta) = \operatorname{argmax}_{x_t \in \mathcal{X}(s_t, \theta)} \sum_{t'=t}^{t+H} c_{t'} x_{t'}$$

- The objective function (evaluated using simulation)
where

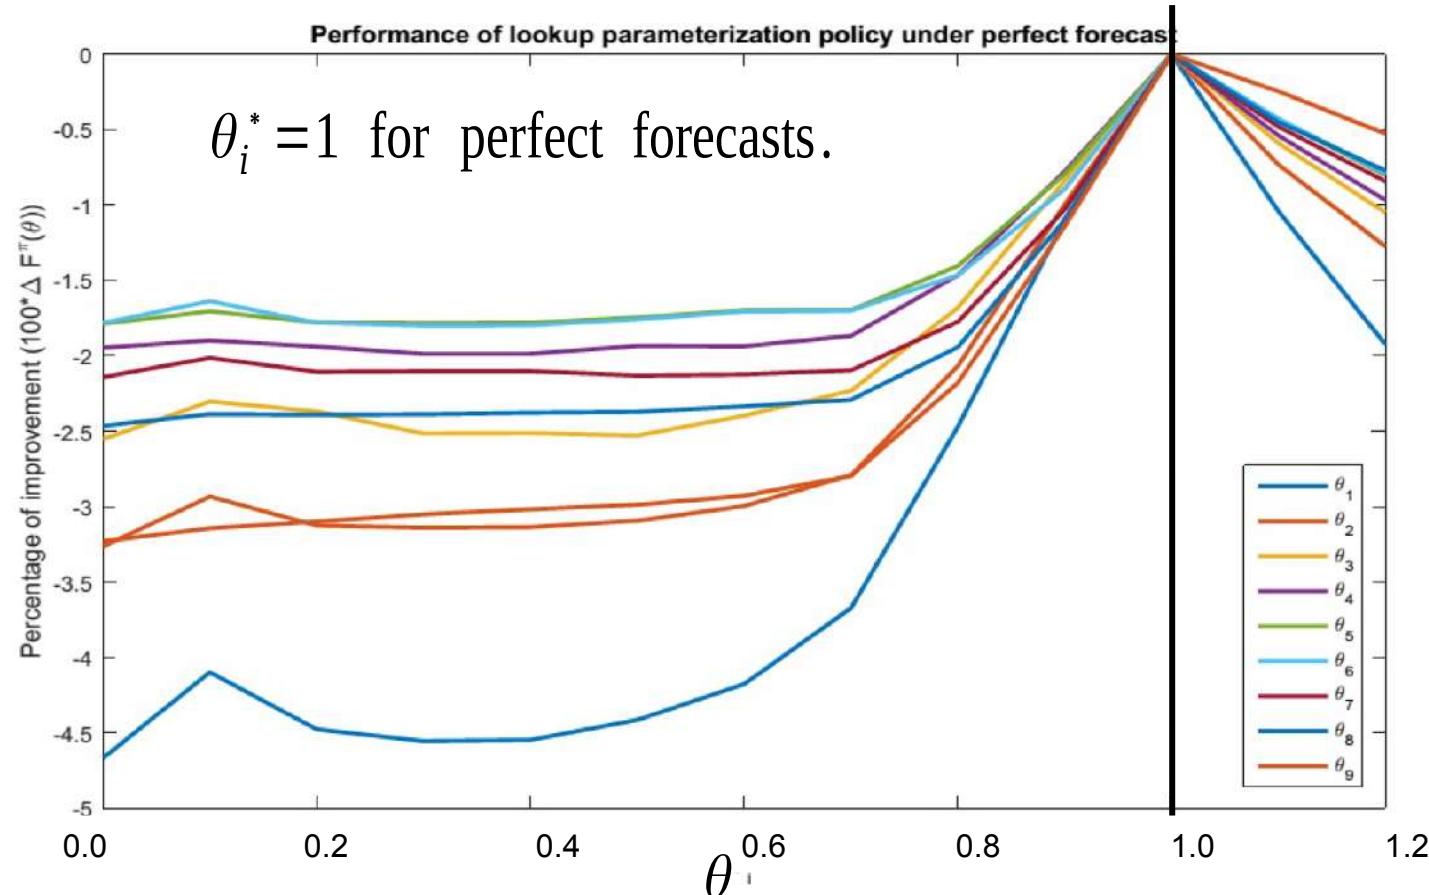
- Stochastic optimization problem

$$\max_{\pi=(f,\theta)} \mathbb{E}_W F^\pi(\theta, W)$$

- Stochastic gradient algorithm (using Spall's SPSA algorithm)

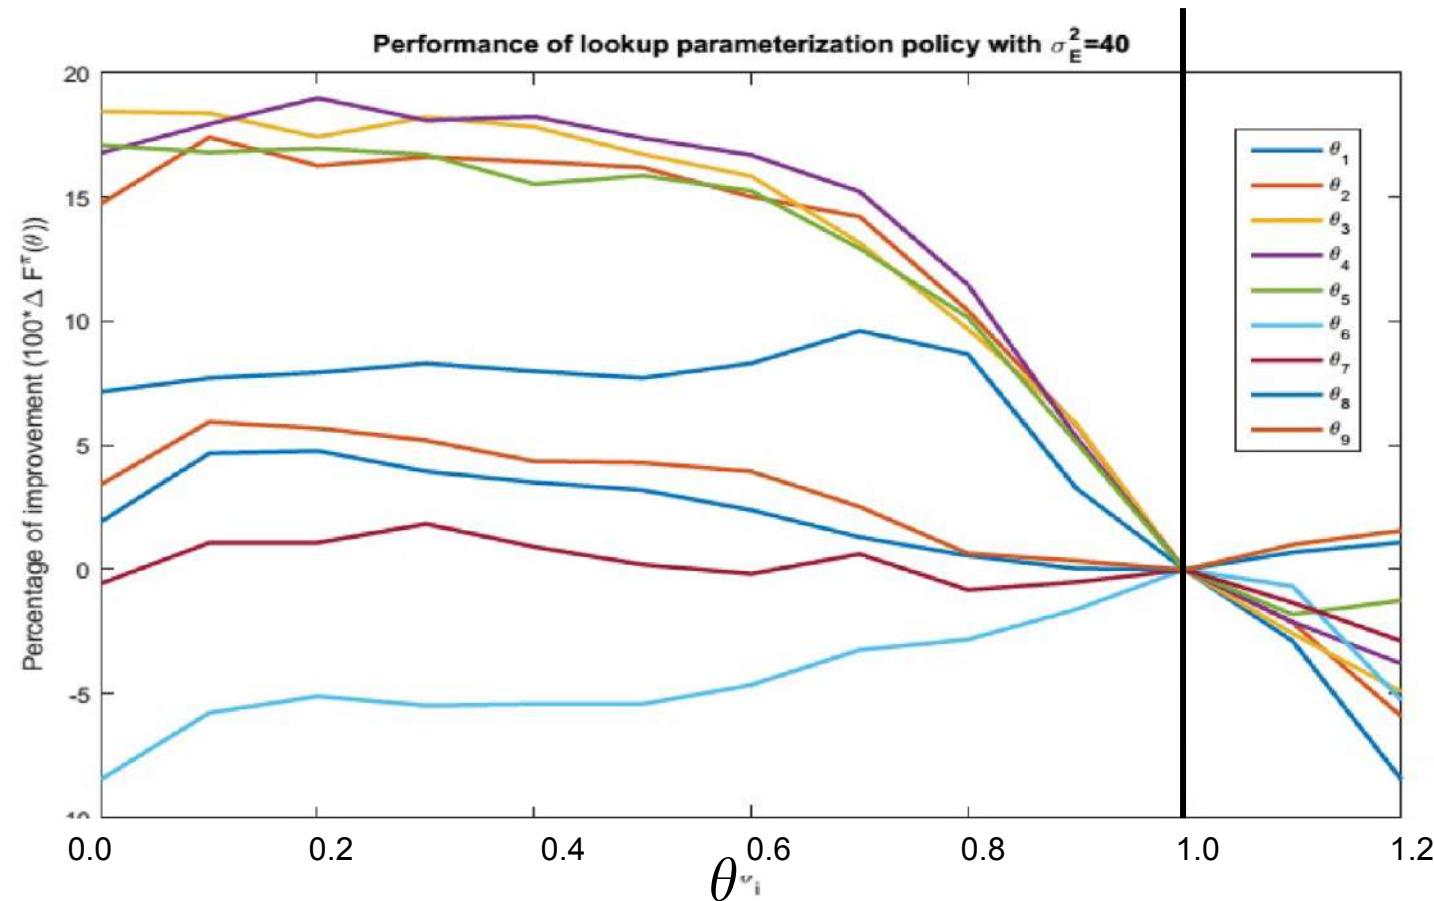
An energy storage application

- One-dimensional contour plots – perfect forecast



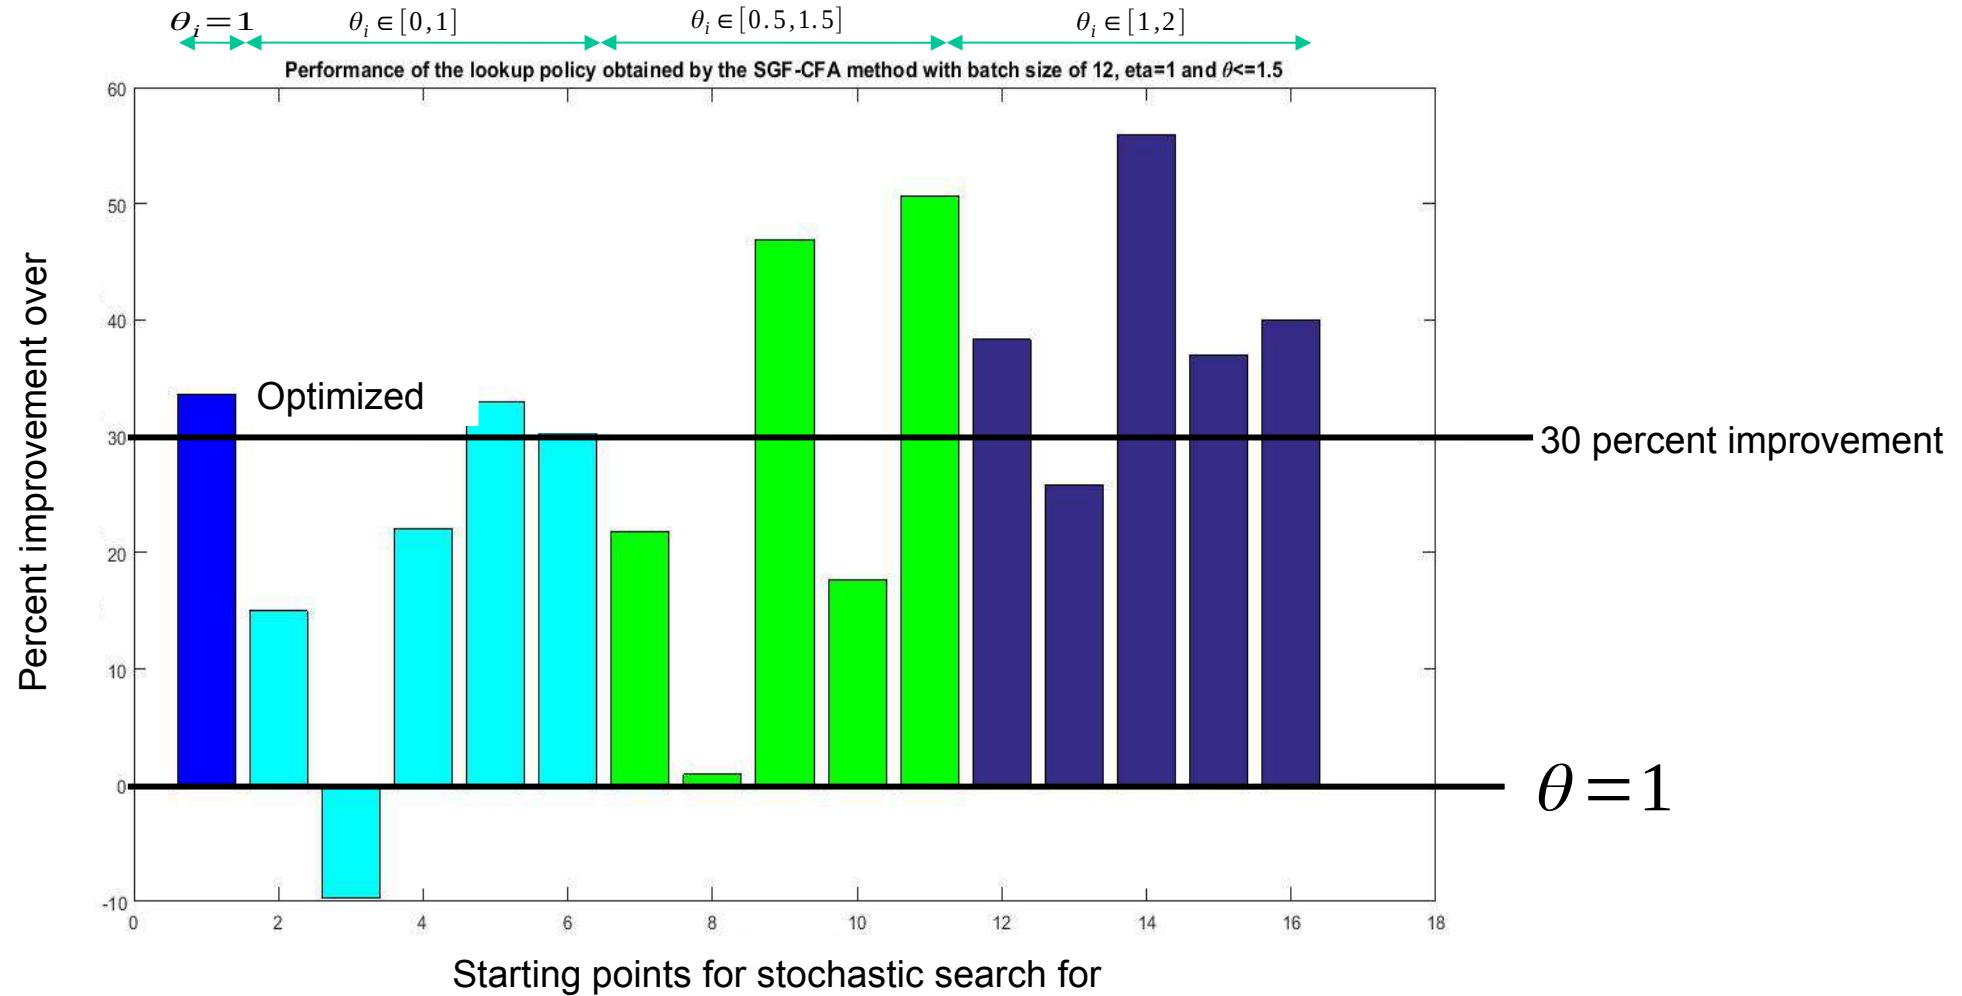
An energy storage application

- One-dimensional contour plots - uncertain forecast



An energy storage application

- Tuning the parameters



An energy storage application

Each policy is best on certain problems

Problem:	Problem description	PFA	CFA	VFA	DLA	DLA/CFA
A	A stationary problem with heavy-tailed prices, relatively low noise, moderately accurate forecasts.	0.959	0.839	0.936	0.887	0.887
B	A time-dependent problem with daily load patterns, no seasonalities in energy and price, relatively low noise, less accurate forecasts.	0.714	0.752	0.712	0.746	0.746
C	A time-dependent problem with daily load, energy and price patterns, relatively high noise, forecast errors increase over horizon.	0.865	0.590	0.914	0.886	0.886
D	A time-dependent problem, relatively low noise, very accurate forecasts.	0.962	0.749	0.971	0.997	0.997
E	Same as (C), but the forecast errors are stationary over the planning horizon.	0.865	0.590	0.914	0.922	0.934

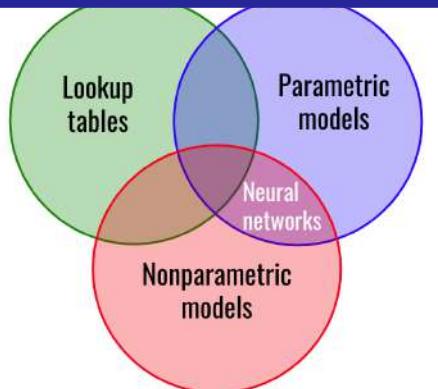
Joint research with Prof. Stephan Meisel, University of Twente, Netherlands.

BRIDGING MACHINE LEARNING & SEQUENTIAL DECISIONS

Machine learning

$$\min_{f \in F, \theta \in \Theta^f} \frac{1}{N} \sum_{n=1}^N (y^n - f(x^n | \theta))^2$$

Searching over functions



"Big dataset"

Sequential decisions

$$\max_{\pi=(f \in F, \theta \in \Theta^f)} \frac{1}{N} \sum_{n=1}^N \sum_{t=0}^T C(S_t^n, X^\pi(S_t^n | \theta))$$

$$S_{t+1} = S^M(S_t, x_t, W_{t+1})$$

Searching over policies

Policy function approximations
Cost function approximations
Value function approximations
Direct lookahead approximations

System model

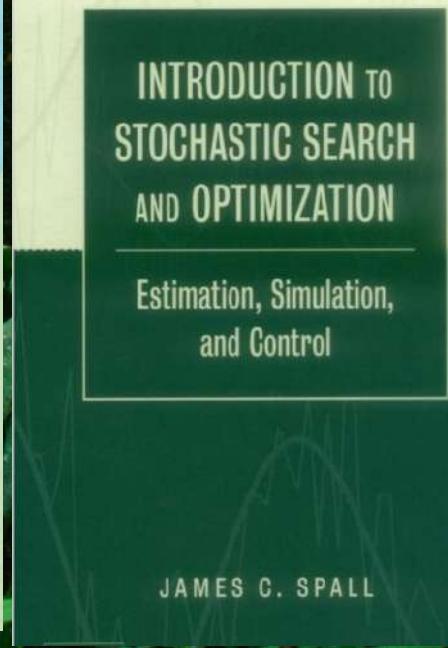
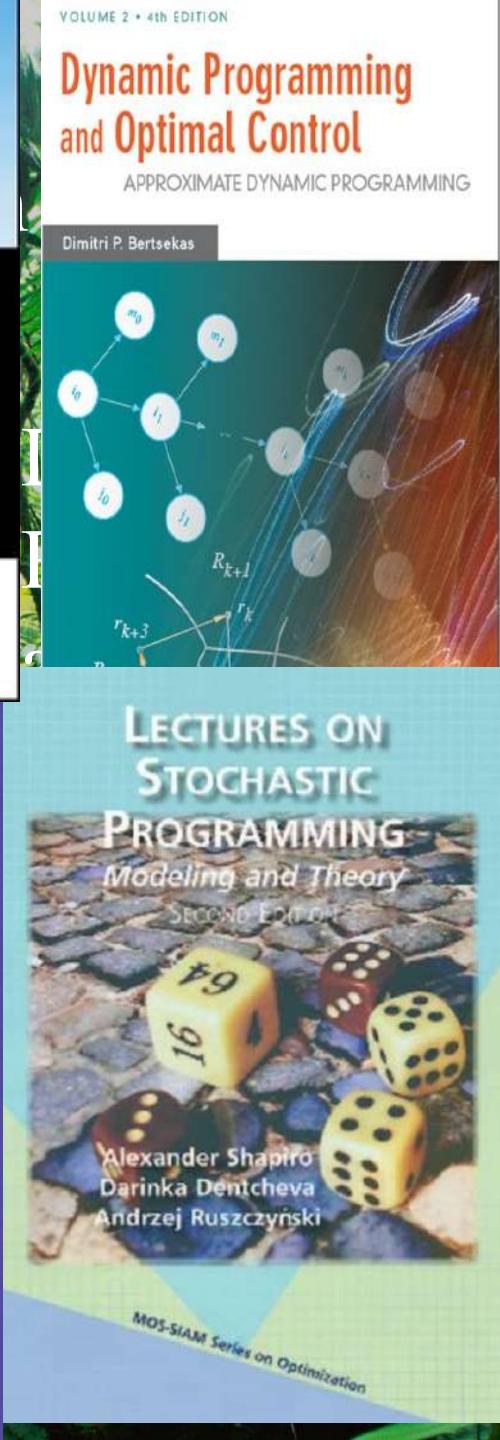
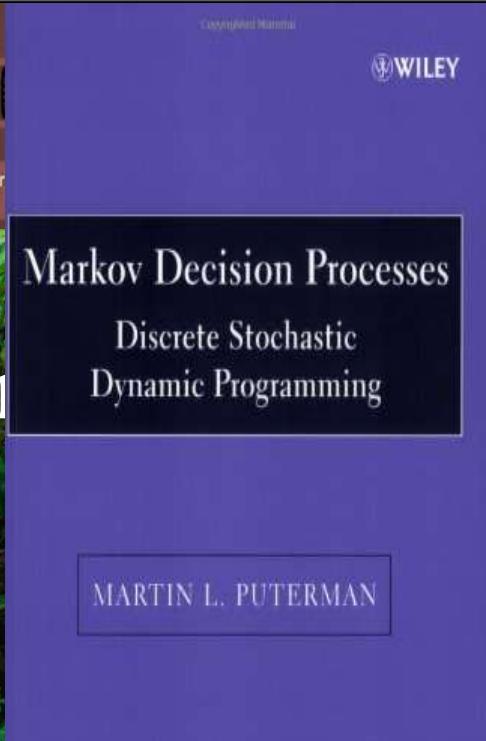
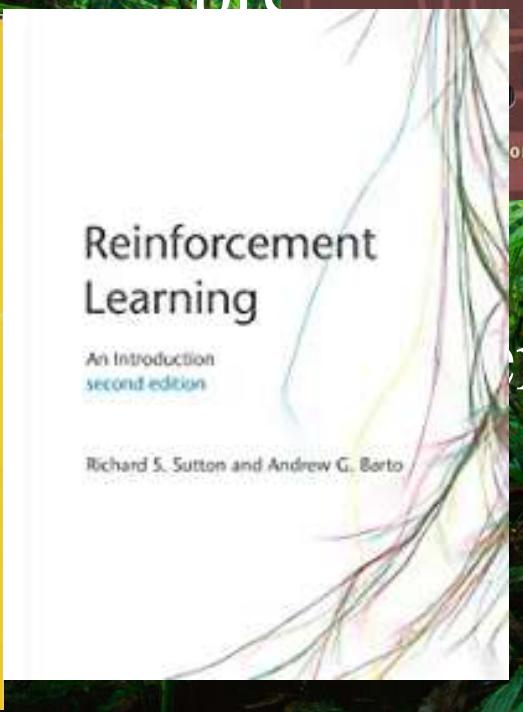
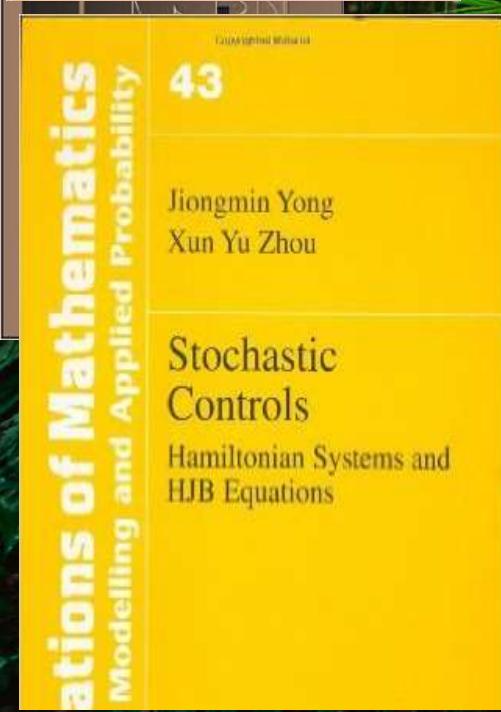
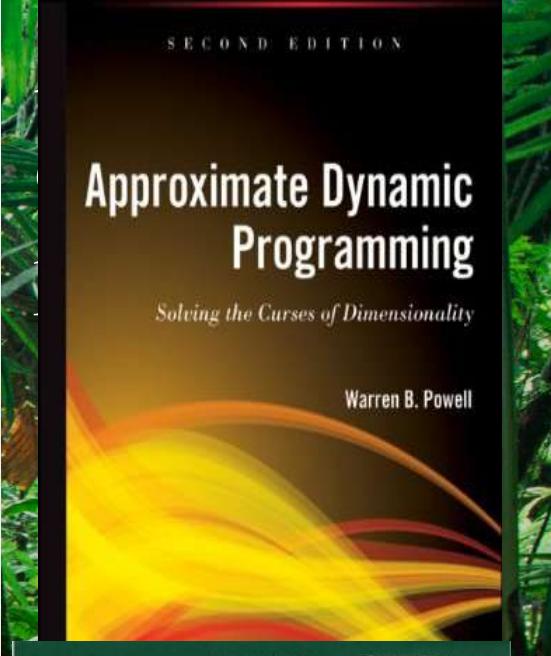
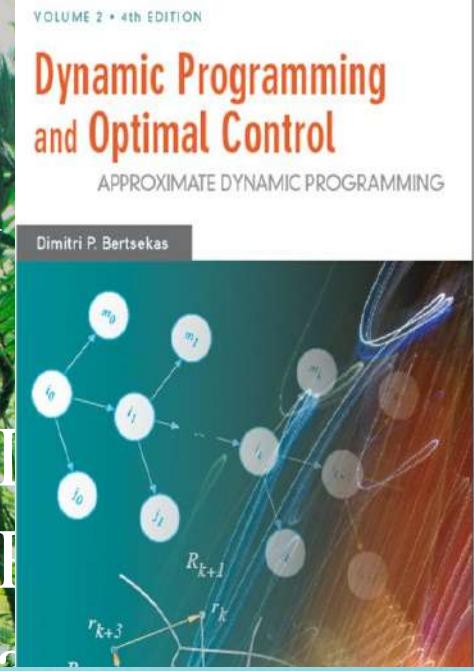
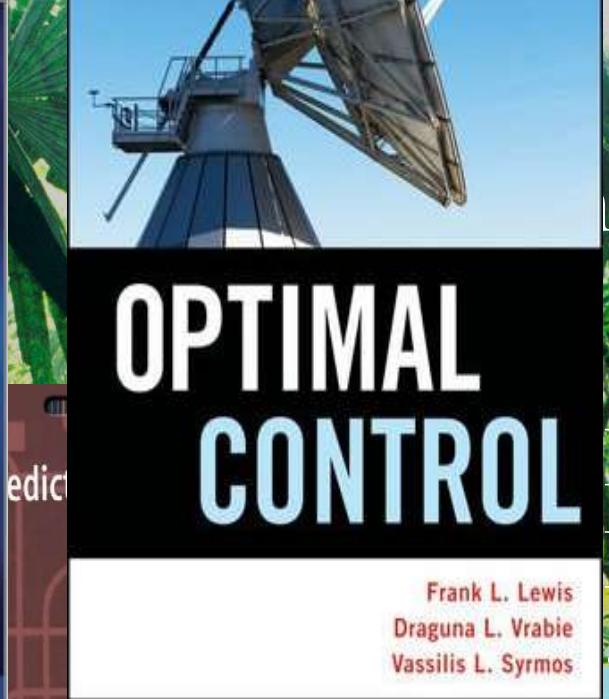
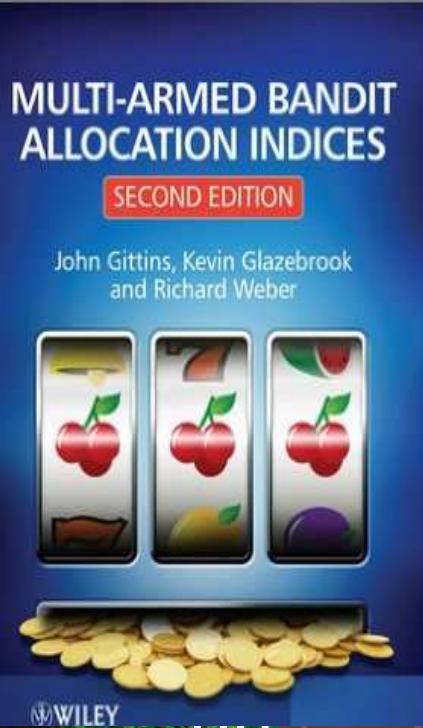
Analytical functions
Optimization problem
Optimization problem
Optimization problem
Optimization problem



Michael C. Fu *Editor*

Handbook of Simulation Optimization

Springer





Choosing a policy class

- It helps to identify five types of policies:
 - » 1) Policy function approximations (PFAs) – Simple rules, analytical functions.
 - » 2) Cost function approximations (CFAs) – Parameterized deterministic optimization models (typically static)
 - » 3) Policies based on value function approximations (VFAs) – Policies that use an approximation of the value of landing in a downstream state
 - » 4) Policies based on direct lookahead approximations (DLAs) – These should be divided into two subclasses:
 - 4a) DLAs using deterministic lookaheads (Det-DLA) – These may be parameterized.
 - 4b) DLAs using stochastic lookaheads (Stoch-DLA)
- So, which are the most useful?

Choosing a policy class

- We can divide the five types of policies into three categories from most to least popular:

- » Category 1 - This category consists of:
 - 1) PFAs - Rules/analytical functions
 - 2) CFAs - Parameterized det. optimization
 - 4a) Det-DLAs - Deterministic lookaheads
- » Category 2 - This category consists of:
 - 4b) Stochastic direct lookaheads
- » Category 3 - This category consists of:
 - 3) Policies based on VFAs.

By far the most widely used in practice. The choice among the three tends to be obvious from the structure of the problem.

Useful for more complex problems where planning into an uncertain future is required, and risk is important.

A very powerful strategy for a very small number of specialized problems.

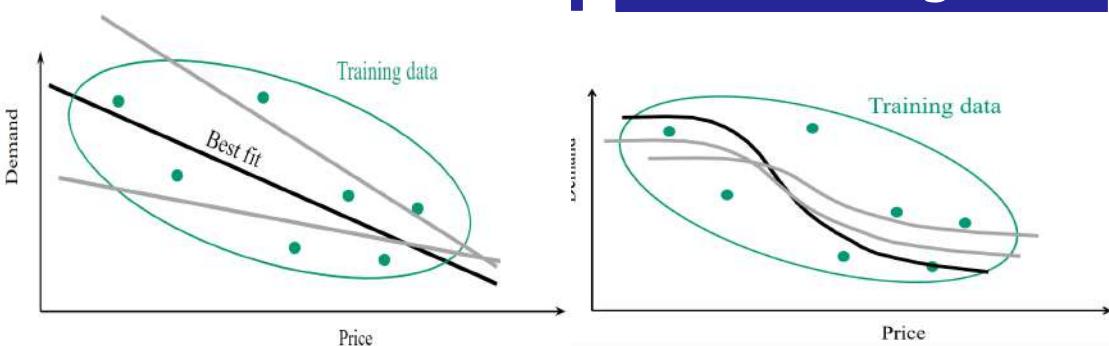
BRIDGING MACHINE LEARNING & SEQUENTIAL DECISIONS

Machine
learning

$$\min_{f \in F} \min_{\theta \in \Theta^f} \frac{1}{N} \sum_{n=1}^N (y^n - f(x^n | \theta))^2$$

Searching over
functions

Parameter
tuning



Parameterized
policies

$$\min_{f \in F} \min_{\theta \in \Theta^f} \frac{1}{N} \sum_{n=1}^N \sum_{t=0}^T C(S_t^n, X^\pi(S_t | \theta))$$

$$S_{t+1} = S^M(S_t, x_t, W_{t+1})$$

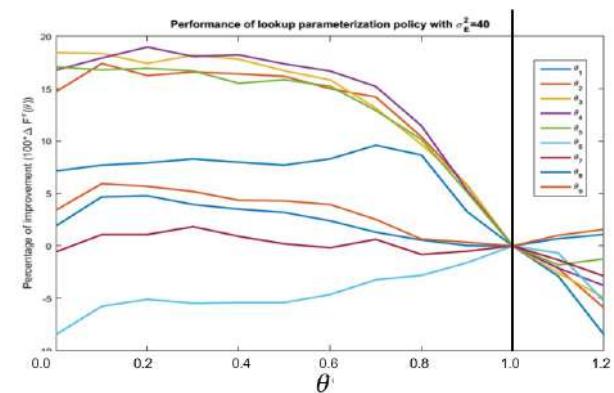
Searching over
functions

Parameter tuning

PFAs

CFAs

Determ.
DLAs



OUTLINE

- The seven levels of artificial intelligence
- The universal modeling framework
- Designing policies
- Mutual fund cash balance optimization
- Choosing the best policy
- A new educational field: sequential decision analytics

Teaching decision analytics

The fields that deal with decisions and uncertainty are completely fragmented.

- » Sequential decision analytics is not a recognized field.
- » There are 15 distinct communities that deal with decisions under uncertainty
- » Each community offers tools that work only for specific problems
- » Real applications require skills that span a wide range of problem settings.



Teaching decision analytics

A graduate-level textbook (2022):

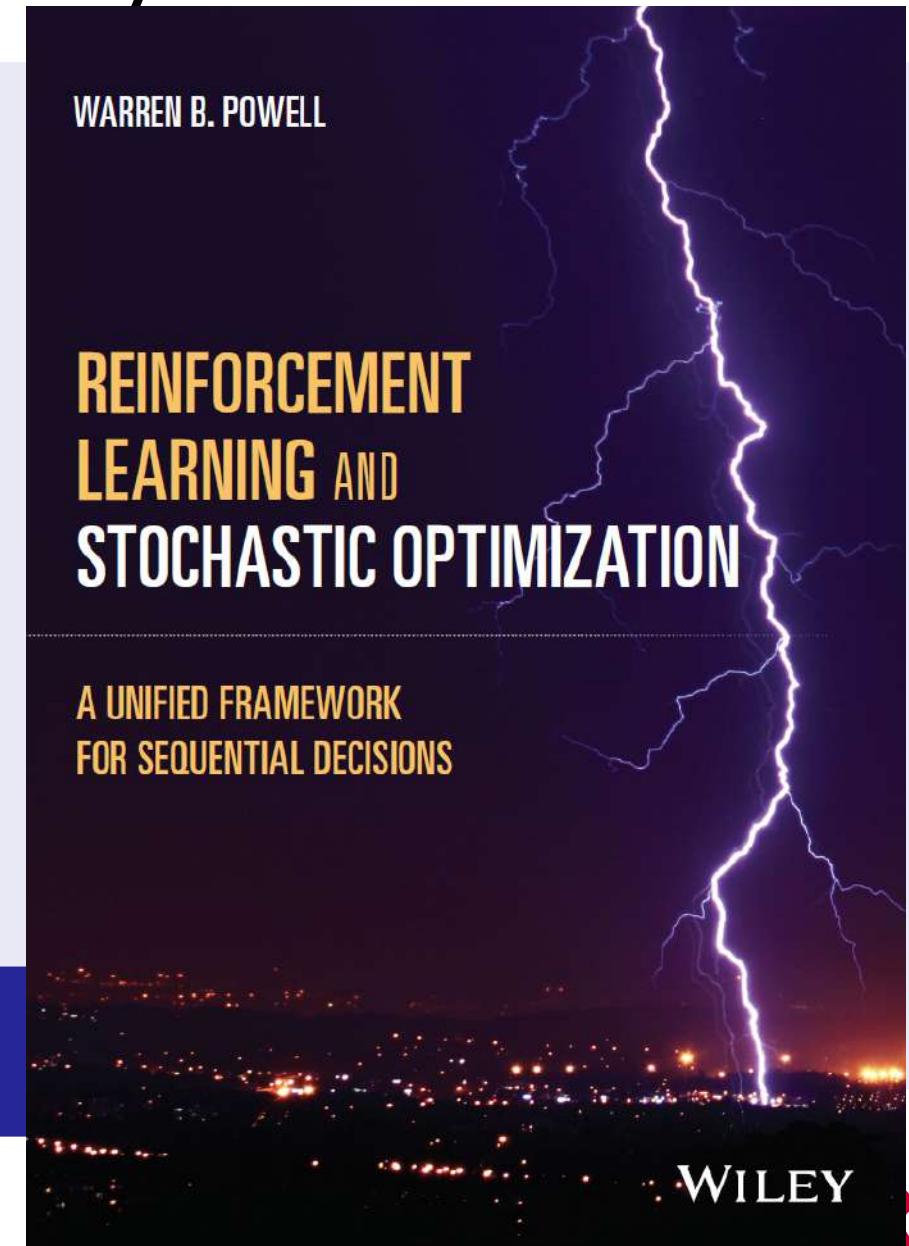
- » First book to introduce a universal modeling framework, covering all four classes of policies.
- » Describes the tools for modeling and solving *any* sequential decision problem, from simple learning problems to truckload fleets to complex supply chains.
- » Aimed at a technical audience interested in writing software to develop models such as those described in this presentation.
- » Provides the foundation for a new field we are calling *sequential decision analytics*.

<http://tinyurl.com/RLandSO/>

WARREN B. POWELL

REINFORCEMENT LEARNING AND STOCHASTIC OPTIMIZATION

A UNIFIED FRAMEWORK
FOR SEQUENTIAL DECISIONS



WILEY

Teaching decision analytics

An introductory book (2022):

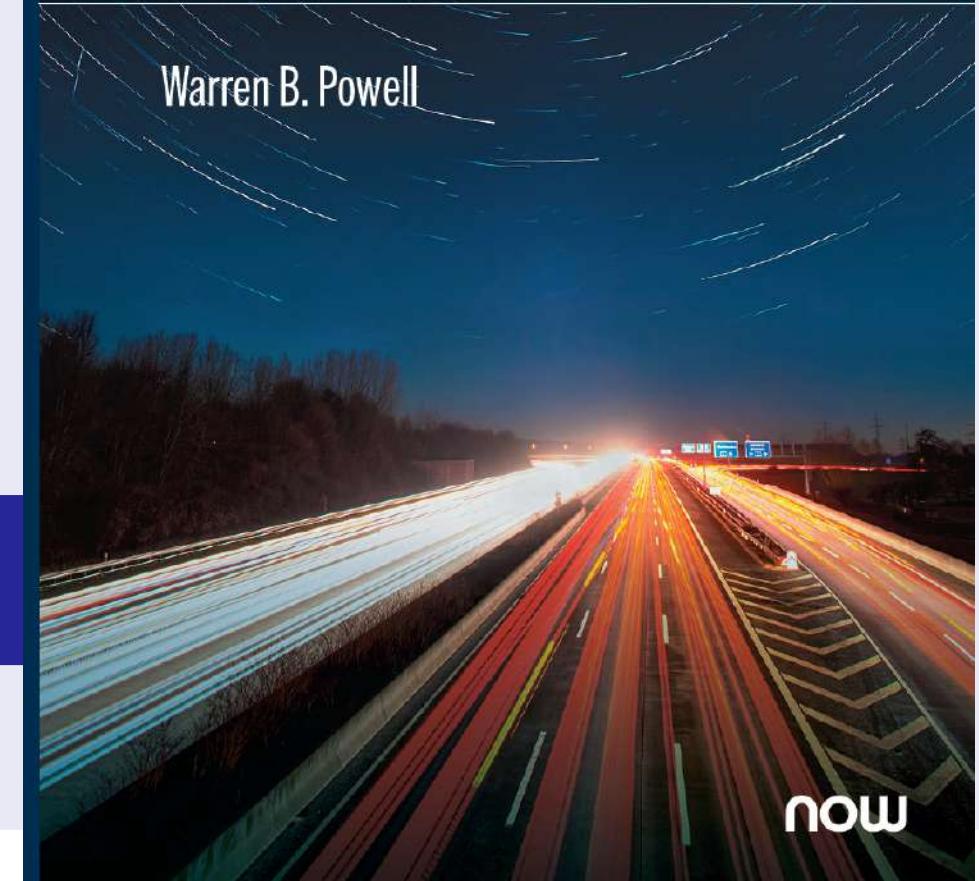
- » Uses a teach-by-example style
- » Illustrates how to model sequential decision problems using a rich set of examples
- » Illustrates all four classes of policies
- » Highlights uncertainty modeling
- » Downloadable from:

<http://tinyurl.com/sdamodeling>

**Sequential Decision Analytics
and Modeling**

Modeling with Python

Warren B. Powell



now

Teaching decision analytics

Chapter 1 - Introduction

Chapter 2 - An asset selling problem

Chapter 3 - Adaptive market planning

Chapter 4 - Learning the best diabetes medication

Chapter 5 - Stochastic shortest path problems - Static

Chapter 6 - Stochastic shortest path problems - Dynamic

Chapter 7 - Applications, revisited

Chapter 8 - Energy storage I

Chapter 9 - Energy storage II

Chapter 10 - Supply chain management I - The two-agent newsvendor problem

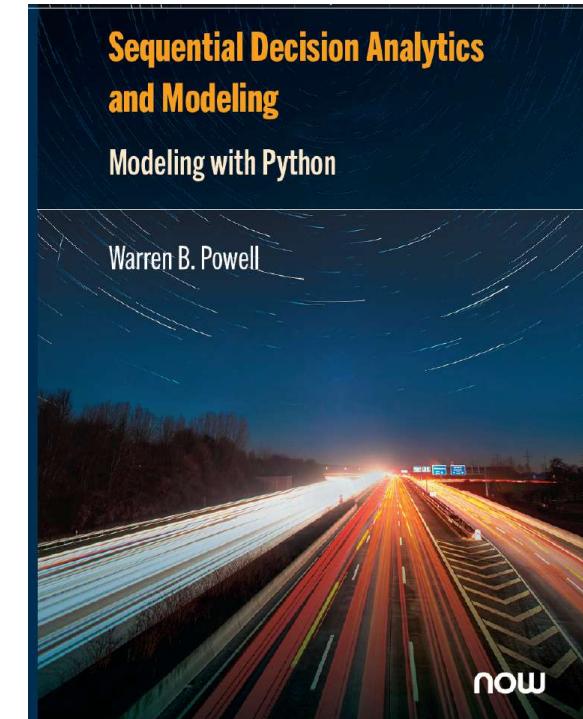
Chapter 11 - Supply chain management II - The beer game

<https://tinyurl.com/sdamodeling/>

Chapter 12 - Ad-click optimization

Chapter 13 - Blood management problem

Chapter 14 - Optimizing clinical trials



Teaching decision analytics

A book for instructors (2024):

- » Fundamentally changes how to teach introduction to optimization
- » Transitions from machine learning, through simpler sequential decision problems, to static and sequential linear, integer and nonlinear problems
- » Illustrates both static and sequential decision problems
- » Available as free download

<http://tinyurl.com/TeachingOpt>

A Modern Approach to Teaching an Introduction to Optimization

Warren B. Powell



Teaching introduction to optimization

The course is divided into 11 topics:

Topic 1 – Machine learning

Topic 2 - Asset selling, inventory problems

Topic 3 - Newsvendor problem

Topic 4 - Optimal learning

Topic 5 - Shortest path problems (static, sequential)

Topic 6 - Overview of sequential decision problems

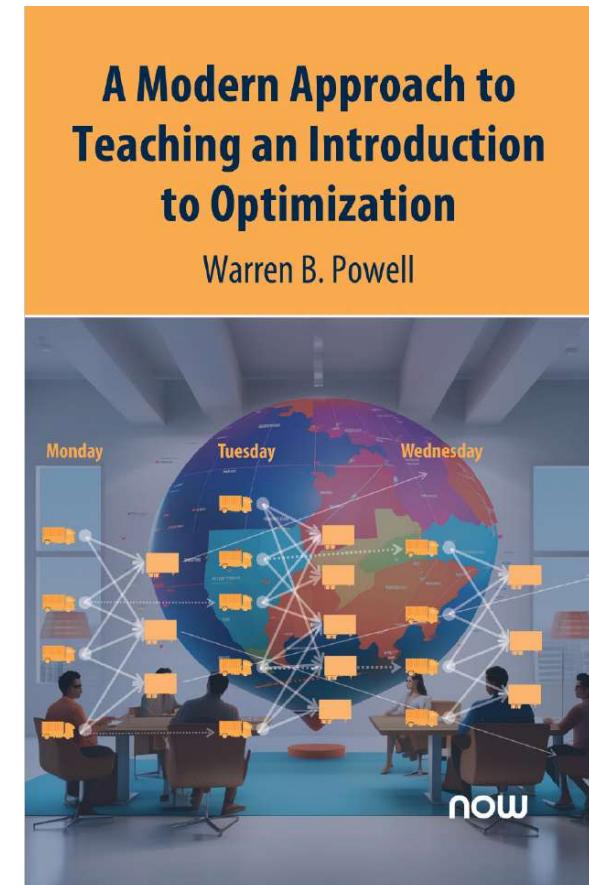
Topic 7 - Linear programming

Topic 8 - Dynamic inventory problem

Topic 9 - Integer programming I (static facility location)

Topic 10 - Integer programming II (dynamic facility location) <https://tinyurl.com/TeachingOpt/>

Topic 11 - Nonlinear programming (static)



Thank you!

An information resource page for sequential decisions:

<http://tinyurl.com/SDAlinks>

This page includes links to:

- » Books (introductory and advanced)
- » Videos
- » Educational webpages
- » Courses and teaching methods
- » A series of educational LinkedIn posts